

Why Rotating Fluids Make Polygons

Paper II: Classical and Quantum Physics

Gordon Speagle

Abstract

Building on the mathematical framework of Paper I, we derive the constrained energy landscape for point vortex rings, the central vortex stabilization mechanism, the $N = 7$ bifurcation, the blob bridge to continuous vorticity, the K_0 finite deformation radius interaction, circulation disorder resilience, and predictions for BEC vortex clusters on hyperbolic lattices.

1 Point vortex dynamics on a two-dimensional surface

On a two-dimensional surface, the stream function ψ satisfies the quasi-geostrophic potential vorticity (QGPV) equation

$$\frac{Dq}{Dt} = 0, \quad q = \nabla^2 \psi + \beta y. \quad (1.1)$$

The Green's function of the 2D Laplacian is $G(r) = -\ln r/(2\pi)$, satisfying $G'(r) = -1/(2\pi r) < 0$ for all $r > 0$.

For N point vortices with equal circulations κ on the collision-free configuration space $\{|z_j - z_k| > 0 \text{ for all } j \neq k\}$, where H is smooth and $\mathbb{Z}_N$ -invariant, the interaction energy is

$$H = - \sum_{j < k} \ln |z_j - z_k|, \quad (1.2)$$

where we absorb $\kappa^2/(4\pi)$ into the normalization. By the equivariant critical point theorem, the regular N -gon—a fixed point of $\mathbb{Z}_N$ —is a critical point of H .

Point vortex dynamics conserves three quantities: the energy H , the angular impulse $L = \sum |z_k|^2$, and the linear impulse $\mathbf{P} = \sum z_k$. The physically accessible configuration space is the $(2N - 3)$ -dimensional surface $\{L = NR^2, \mathbf{P} = 0\}$.

1.1 The centered spiral deformation

Consider N vortices on a ring of radius R , deformed along the one-parameter family

$$z_k(\sigma) = R e^{\sigma(k-(N-1)/2)/N} e^{2\pi i k/N}, \quad k = 0, \dots, N-1. \quad (1.3)$$

At $\sigma = 0$: the regular N -gon. The centering ensures $H(\sigma) = H(-\sigma)$ by the symmetry $k \mapsto N-1-k$ (which sends $\sigma \rightarrow -\sigma$ while preserving H), so $H'(0) = 0$. The spiral

deformation has a precise Möbius-group origin: the positions $\{z_k(\sigma)\}$ are the orbit of the seed $z_0 = R$ under the cyclic group generated by $z \mapsto \lambda z$, $\lambda = e^{\sigma/N + 2\pi i/N}$. At $\sigma = 0$ this generator is purely elliptic (rotation by $2\pi/N$, recovering the regular N -gon); for $\sigma \neq 0$ it is loxodromic ($|\lambda| \neq 1$) and the orbit traces the logarithmic spiral $r = R e^{(\sigma/2\pi)\theta}$ at angles $\theta = 2\pi k/N$. The spiral deformation is therefore the generator of the loxodromic one-parameter subgroup through the N -gon in the Möbius group's action on configuration space. $\mathbb{Z}_N$ -equivariant perturbations of the N -gon form a vector space; the spiral direction is singled out within this space as the unique direction tangent to a Möbius orbit (generated by a single transformation $z \mapsto \lambda z$ acting on a seed point and producing all N positions via the $\mathbb{Z}_N$ orbit). Other perturbation directions (individual Fourier modes, radial displacements) do not arise from a single Möbius seed and are analyzed through the full constrained Hessian in Theorem 2.2. The Möbius interpretation does not enter the stability proofs: the stability analysis uses the spiral only as a specific perturbation direction (the unique $\mathbb{Z}_N$ -equivariant direction that leaves the constraint surface at $O(\sigma^2)$, so that the unconstrained $H''(0) < 0$ does not contradict the constrained energy minimum). However, Proposition Paper I, Proposition 2.3 does carry genuine mathematical content beyond geometric decoration. Restricted to the dilation subgroup $f(z) = bz$, where f' is the constant b , the single-site correction becomes position-independent and reduces to $H(bz_1, \dots, bz_N) = H(z_1, \dots, z_N) - \frac{N(N-1)}{2} \ln b$. This is the N -body statement that H is dilation-invariant—the logarithmic energy changes only by an additive constant under spatial rescaling (Corollary Paper I, Corollary 2.4 above).

2 The constrained energy landscape

2.1 Energy curvature along the spiral

Theorem 2.1 (Energy curvature at the polygon). *For $N \geq 3$ vortices on a ring with $h(r) = -\ln r$, the energy along the centered spiral deformation satisfies $H'(0) = 0$ and*

$$H''(0) = -\frac{1}{4N^2} \sum_{m=1}^{N-1} \frac{(N-m)m^2}{\sin^2(\pi m/N)} < 0. \quad (2.1)$$

Proof. Part (a): $H'(0) = 0$ by the centering symmetry $k \mapsto N-1-k$. Part (b): Define $F(u, m) = \ln[1 - 2\cos(2\pi m/N) e^{um} + e^{2um}]$ with $u = \sigma/N$. Direct differentiation gives $F''(0, m) = m^2/[2\sin^2(\pi m/N)]$. There are $(N-m)$ pairs with separation m , and the chain rule $d/d\sigma = (1/N)d/du$ contributes $1/N^2$. This derivation uses the specific structure of $h = -\ln r$: the general formula Paper I, eq. 3.1 has two terms, but for the logarithmic interaction $(rh')' = (-1)' = 0$, so the S_m terms (first term in Paper I, eq. 3.1) vanish identically. Only the second term survives, giving (2.1) as a special case of Paper I, eq. 3.1 with $A_m = 0$ and $B_m = h'(d_m) m^2(N-m)/(N^2 d_m) = -m^2(N-m)/(N^2 d_m^2)$. Substituting $d_m = 2\sin(\pi m/N)$ yields (2.1). Every factor in the sum is strictly positive, so $H''(0) < 0$. $\square$

2.2 The constrained energy minimum

Attribution note. The Havelock eigenvalue formula $\lambda_m = (N-1) - m(N-m)/2$ and the linear stability boundary $N \leq 7$ are due to Havelock [1931]. The new contribution of the theorem below is the reframing: the N -gon is a *constrained energy minimum* (not a

saddle point or maximum), which resolves the apparent tension with Onsager's statistical-mechanics picture.

Theorem 2.2 (Constrained energy minimum). *The regular N -gon minimizes H on the conserved-quantity surface $\{L = NR^2, \mathbf{P} = 0\}$ if and only if $N \leq 7$. The Lagrangian Hessian*

$$\nabla^2 \mathcal{L} = \nabla^2 H - 2\mu_L I, \quad \mu_L = -\frac{N-1}{4}, \quad (2.2)$$

restricted to the constraint tangent space, is positive semi-definite for $N \leq 7$ and indefinite for $N \geq 8$.

Proof. The energy-Casimir functional is $\mathcal{H} = H + |\mu_L| L$, whose Hessian is (2.2). The Hessian $\nabla^2 H$ is a $2N \times 2N$ real symmetric matrix that commutes with the $\mathbb{Z}_N$ action $z_k \mapsto z_{k+1 \pmod{N}}$ (since H is invariant under cyclic permutation of equal-circulation vortices). By Schur's lemma for real representations of finite abelian groups, $\nabla^2 H$ block-diagonalizes in the discrete Fourier basis into 2×2 blocks, one per mode $m \in \{0, \dots, N-1\}$. For mode m , decompose the perturbation at vortex k into radial (δr_k) and tangential (δt_k) components, with Fourier amplitudes $\delta r_k = a \cos(2\pi m k/N)$, $\delta t_k = b \cos(2\pi m k/N)$. The 2×2 Hessian block for the pair (a, b) involves two sums over separations $p = 1, \dots, N-1$:

$$\begin{aligned} H_{rr}^{(m)} &= \sum_p \left[\frac{1 - \cos(2\pi p m/N)}{4 \sin^2(\pi p/N)} + \frac{1 + \cos(2\pi p m/N)}{4 \sin^2(\pi p/N)} \right] - 2 \sum_p \frac{1 - \cos(2\pi p m/N)}{4 \sin^2(\pi p/N)}, \\ H_{rt}^{(m)} &= \sum_p \frac{\sin(2\pi p m/N)}{4 \sin^2(\pi p/N)}, \\ H_{tt}^{(m)} &= -H_{rr}^{(m)}. \end{aligned}$$

The off-diagonal $H_{rt}^{(m)}$ vanishes identically by the anti-symmetry $\sin(2\pi p m/N) \rightarrow -\sin(2\pi(N-p)m/N)$ under $p \rightarrow N-p$, since $\sin^2(\pi p/N) = \sin^2(\pi(N-p)/N)$ (so the denominator is symmetric while the numerator is anti-symmetric; the sum over the full range $p = 1, \dots, N-1$ cancels in pairs). Therefore the block is diagonal: $\text{diag}(H_{rr}^{(m)}, -H_{rr}^{(m)})$, giving eigenvalues $\pm H_{rr}^{(m)} = \pm \sigma_m$.

The diagonal entry $H_{rr}^{(m)}$ simplifies using the Havelock identity (Appendix Paper I, §4.1): the self-sum $\sum_p 1/[4 \sin^2(\pi p/N)] = (N-1)/2$ (a special case of $T_0 + (N-1)/2$) and the cross-sum is $T_m = m(N-m)/2$. Thus

$$\sigma_m = H_{rr}^{(m)} = \frac{N-1}{2} - \frac{m(N-m)}{2}, \quad (2.3)$$

using

$$T_m := \sum_{p=1}^{N-1} \frac{1 - \cos(2\pi p m/N)}{4 \sin^2(\pi p/N)} = \frac{m(N-m)}{2} \quad (2.4)$$

(proven in Appendix Paper I, §4.1; cf. Havelock [1931], §3). The Lagrangian eigenvalues are thus

$$\lambda_m^+ = (N-1) - \frac{m(N-m)}{2}, \quad \lambda_m^- = \frac{m(N-m)}{2}.$$

Since $\lambda_m^- \geq 0$ always, positive semi-definiteness requires $\lambda_m^+ \geq 0$ for the modes tangent to the constraint surface. Three modes are removed: *Mode* $m = 0$ corresponds to uniform dilation $\delta z_k \propto z_k$, which changes $L = \sum |z_k|^2$ and is therefore not tangent to $\{L = NR^2\}$. *Modes* $m = 1$ and $m = N-1$ generate infinitesimal translation ($\delta z_k = \text{const}$, changing $\mathbf{P} = \sum z_k$) and overall rotation ($\delta z_k = i z_k$, a Goldstone mode of the $SO(2)$ symmetry of

H). Both have $\sigma_m = 0$, consistent with the continuous symmetry group; they lie in the kernel of $\nabla^2 H$ and produce the zero eigenvalues. The remaining modes $m = 2, \dots, N-2$ are tangent to the constraint surface and carry the stability information. The minimum of λ_m^+ over m occurs at $m = \lfloor N/2 \rfloor$:

$$\min_m \lambda_m^+ = \begin{cases} (N-1)(7-N)/8 & N \text{ odd,} \\ (N-1) - N^2/8 & N \text{ even.} \end{cases}$$

Both expressions are > 0 for $N \leq 6$, equal to zero at $N = 7$, and < 0 for $N \geq 8$. Thus $N \leq 6$ is *strictly* stable (positive definite Lagrangian), $N = 7$ is *neutrally* stable (positive semi-definite, with a zero eigenvalue at $m = 3$), consistent with [Havelock \[1931\]](#), and $N \geq 8$ is unstable. (The $N = 7$ marginal case is further analyzed in Section 4.) $\square$

Corollary 2.3 (Constrained minimum versus unconstrained saddle). *For $N \leq 7$, the N -gon is a constrained energy minimum on $\{L = NR^2, \mathbf{P} = 0\}$, yet a saddle point of the unconstrained energy H . These two facts are compatible: the constrained minimum is the relevant concept for Lyapunov stability via Arnold’s energy-Casimir method, while the Onsager negative-temperature argument operates on the unconstrained energy landscape, where the polygon is a maximum-energy accessible state (not a local maximum of H , but the highest- H configuration reachable within the Thomson-stable set). The distinction matters because the statistical mechanics literature sometimes describes the polygon as an energy maximum without qualification; it is a constrained minimum that happens to have higher energy than alternative configurations at the same conserved quantities.*

Proof. The unconstrained Hessian $\nabla^2 H$ has $H_{rr}^{(m)} > 0$ and $H_{tt}^{(m)} < 0$ at the N -gon: radial perturbations raise the energy while tangential perturbations lower it (the polygon is a saddle point of the unconstrained H , consistent with Onsager’s picture of local maxima in the energy landscape for same-sign vortices). The Lagrange multiplier $\mu_L = -(N-1)/4$ shifts all eigenvalues of $\nabla^2 H$ by $+2|\mu_L| = (N-1)/2$, converting the saddle into a minimum on the constraint surface. The value of μ_L follows from the constrained stationarity condition: at vortex 0 (position $(R, 0)$), the pair $(0, k)$ contributes identically $-1/(2R)$ to $\partial H/\partial x_0$ regardless of k , because $(1 - \cos \theta_k)/[4 \sin^2(\theta_k/2)] = 1/2$ for all k —a consequence of the scale invariance of $h'(r)/r = -1/r^2$. Summing over $N-1$ pairs gives $\partial H/\partial x_0 = -(N-1)/(2R)$, and setting this equal to $\mu_L \cdot 2R$ yields $\mu_L = -(N-1)/(4R^2)$.

In particular, Theorem 2.1 ($H''(0) < 0$ along the spiral direction) does not contradict Theorem 2.2: the spiral deformation (1.3) leaves the conserved surface at $O(\sigma^2)$, so it is not a constrained perturbation. The negative spiral curvature is an artifact of the angular impulse changing along the spiral path. $\square$

3 Central vortex stabilization

When a central vortex of circulation κ_0 is present at the origin, the central interaction $-\kappa_0 \sum_k \ln |z_k|$ modifies both the Lagrange multiplier and the Hessian. Two effects contribute to the critical eigenvalue:

Lagrange multiplier shift. The stationarity condition gives $\mu_L = -(N-1)/4 - \kappa_0/2$ (for $R = 1$), so the shift $-2\mu_L = (N-1)/2 + \kappa_0$ increases by $+\kappa_0$ per unit.

Hessian projection. The central-vortex Hessian at vortex k is the traceless 2×2 matrix $(\kappa_0/R^2) \begin{bmatrix} \cos 2\theta_k & \sin 2\theta_k \\ \sin 2\theta_k & -\cos 2\theta_k \end{bmatrix}$, which has eigenvalue $+\kappa_0/R^2$ for radial perturbations

$(\cos \theta_k, \sin \theta_k)$ and $-\kappa_0/R^2$ for tangential. For $N \geq 8$, the critical eigenvector of the ring-only Lagrangian (the most negative eigenvalue direction, at mode $m = \lfloor N/2 \rfloor$) is *purely radial* at every vortex position. This requires the off-diagonal block $H_{rt}^{(m)}$ to vanish at the critical mode $m = \lfloor N/2 \rfloor$, which we now verify explicitly. From equation (2.3): $H_{rt}^{(m)} = \sum_{p=1}^{N-1} \sin(2\pi pm/N)/[4\sin^2(\pi p/N)]$. The substitution $p \rightarrow N-p$ sends $\sin(2\pi pm/N) \rightarrow \sin(2\pi(N-p)m/N) = -\sin(2\pi pm/N)$ (anti-symmetric) while $\sin^2(\pi p/N) \rightarrow \sin^2(\pi(N-p)/N) = \sin^2(\pi p/N)$ (symmetric). Therefore each pair $(p, N-p)$ cancels, and $H_{rt}^{(m)} = 0$ for *every* m —in particular for the critical mode $m = \lfloor N/2 \rfloor$. The Hessian block is diagonal at all modes: $\text{diag}(H_{rr}^{(m)}, -H_{rr}^{(m)})$, so radial and tangential perturbations decouple completely. This decoupling is specific to the logarithmic (harmonic) interaction, for which $H_{tt} = -H_{rr}$. Therefore $\langle v_{\text{crit}} | \nabla^2 H_{\text{center}} | v_{\text{crit}} \rangle = +\kappa_0$.

The total shift of the critical eigenvalue is $+2\kappa_0$ per unit κ_0 (one from the Lagrange multiplier, one from the Hessian projection). Setting the shifted eigenvalue to zero:

Proposition 3.1 (Critical central vortex strength). *For $N \geq 8$, the N -gon with central vortex is a constrained energy minimum if and only if $\kappa_0/\kappa \geq \kappa_{\text{crit}}(N)$, where*

$$\kappa_{\text{crit}}(N) = \begin{cases} (N^2 - 8N + 8)/16 & N \text{ even}, \\ (N-1)(N-7)/16 & N \text{ odd}. \end{cases} \quad (3.1)$$

Proof. By Theorem 2.2, the minimum constrained Lagrangian eigenvalue at $\kappa_0 = 0$ is $\lambda_{\text{min}} = (N-1) - \lfloor N/2 \rfloor(N - \lfloor N/2 \rfloor)/2$. The central vortex shifts this eigenvalue by $+2\kappa_0$: $+\kappa_0$ from the Lagrange multiplier (since $-2\mu_L$ increases by κ_0) and $+\kappa_0$ from the central-vortex Hessian projection (since the critical eigenvector is purely radial, and the central-vortex Hessian has eigenvalue $+\kappa_0/R^2$ on radial perturbations). Setting $(N-1) - m(N-m)/2 + 2\kappa_0 = 0$ with $m = \lfloor N/2 \rfloor$ gives (3.1). $\square$

The first values are:

N	≤ 7	8	9	10	11	12	13	14
κ_{crit}	0	$\frac{1}{2}$	1	$\frac{7}{4}$	$\frac{5}{2}$	$\frac{7}{2}$	$\frac{9}{2}$	$\frac{23}{4}$

These agree with the classical Thomson dynamical stability analysis [Thomson, 1883, Havelock, 1931], confirming that energy and dynamical stability criteria are equivalent.

For $N = 6$, the exact critical ratio $\kappa_{\text{crit}}(6) = -1/4$ follows algebraically from the shifted-eigenvalue formula in the proof of Proposition 3.1. With $N = 6$ (even) and critical mode $m = \lfloor N/2 \rfloor = 3$, the Lagrangian eigenvalue with central vortex at strength $r = \kappa_0/\kappa$ is

$$\lambda_3 = (N-1) - \frac{m(N-m)}{2} + 2r = 5 - \frac{9}{2} + 2r = \frac{1}{2} + 2r.$$

Setting $\lambda_3 = 0$ gives $r = -1/4$. The closed form $\kappa_{\text{crit}} = (N^2 - 8N + 8)/16$ of Proposition 3.1 evaluates at $N = 6$ to $(36 - 48 + 8)/16 = -4/16 = -1/4$, consistent with this. (The Proposition is stated for $N \geq 8$; at $N = 6$ the formula returns $\kappa_{\text{crit}} < 0$, meaning destabilisation requires an anticyclonic central vortex with $r < 0$.)

The same threshold appears in the dynamical stability analysis. The 12×12 linearisation of the hexagonal ring with central vortex [Havelock, 1931] has a characteristic polynomial containing the low-frequency factor

$$9 + 36\kappa_0/\kappa + 16\pi^2\lambda^2.$$

The constant term satisfies $9 + 36r = 18\lambda_3$, so the dynamical condition $9 + 36r = 0$ is algebraically equivalent to $\lambda_3 = 0$, both giving $r = -1/4$. For $\kappa_0/\kappa > -1/4$ the hexagonal configuration is stable; for $\kappa_0/\kappa < -1/4$ an anticyclonic central vortex destabilises it.

The center-vs-ring geometry is essential: multi-ring configurations are generically less stable (Section 5).

4 The $N = 7$ bifurcation

At $N = 7$, the minimum constrained Lagrangian eigenvalue is zero (mode $m = 3$). In the language of critical phenomena (Remark Paper I, Remark 4.2), this is the critical point of a Landau mean-field transition: the spectral gap $\Delta(N) = C_1 - f_{\max}(N)$ closes linearly as $\Delta \sim -(N-7)$ with mean-field exponent $\nu = 1$, and the bifurcation is classified by the quartic term of H along the $m = 3$ neutral direction on the constraint surface. The $\mathbb{Z}_7$ -equivariant normal form is the Landau potential $V(\varepsilon) = (\Delta/2)\varepsilon^2 + (\alpha_0/4)\varepsilon^4$, where $\varepsilon = |z|$ is the amplitude of the $m = 3$ mode and $\Delta = 0$ at $N = 7$. Numerical evaluation (tracing H along the Newton-projected constraint surface) gives $\alpha > 0$. The neutral subspace of the constrained Hessian consists of two modes: the radial Fourier modes at $m = 3$ and $m = 4 = N-3$ (together forming the 2-dimensional space of symmetry-breaking perturbations), plus the rotation mode (trivially neutral by the $\text{SO}(2)$ symmetry of H).

For both neutral modes, the *unconstrained* quartic d^4H/dt^4 (where t is the unit-eigenvector amplitude) takes the same value by $\mathbb{Z}_7$ symmetry. Each of the 21 vortex pairs contributes

$$\frac{6C^2}{A^2} - \frac{12CB^2}{A^3} + \frac{3B^4}{A^4},$$

where $A = |\Delta z|^2$, $B = 2\text{Re}(\Delta z \cdot \overline{\Delta u})$, $C = |\Delta u|^2$ are evaluated at the unit-ring N -gon ($|\Delta z|^2 = 4\sin^2(\pi p/N)$ for separation class p). Direct summation over all 21 pairs gives

$$\left. \frac{d^4H}{dt^4} \right|_{m=3} = \frac{153}{7}. \quad (4.1)$$

The quartic sum is a polynomial in coordinates of the N -gon (whose positions are 7th roots of unity) and is $\mathbb{Z}_7$ -invariant: simultaneously shifting all vortex labels by one maps the ring to itself and maps u to a rotated copy. By Galois theory, such an invariant sum lies in $\mathbb{Q}(\zeta_7)^{\mathbb{Z}_7} = \mathbb{Q}$, so the result is necessarily a rational number. The value $153/7$ is then identified by the following argument: the 21 pairs decompose into 6 separation classes $p = 1, \dots, 6$ with 7 pairs each, and the per-pair quartic (4.1) depends on p only through $A_p = 4\sin^2(\pi p/7)$ and the mode overlap $B_p = 2\text{Re}(\Delta z \cdot \overline{\Delta u})$, $C_p = |\Delta u|^2$. These are algebraic numbers in $\mathbb{Q}(\zeta_7)$. The Galois rationality argument is the proof: $\mathbb{Z}_7$ -invariant sums over $\mathbb{Q}(\zeta_7)$ lie in the fixed field $\mathbb{Q}$, so the quartic coefficient is rational *a priori*; direct algebraic summation (`n7_bifurcation.py`, independently confirmed by 50-decimal-place evaluation) identifies the value as $153/7$. (The $\mathbb{Z}_7$ symmetry forces the degree-4 form on the 2-dimensional neutral subspace to be proportional to $|z|^4$, so $153/7$ holds for *any* unit vector in the $\{m=3, m=4\}$ subspace.)

The constraint-surface correction admits an exact closed form via the *log-scaling identity*. The Newton projection onto $\{L = N\}$ sends $\varepsilon \mapsto z(\varepsilon) = (z + \varepsilon u) \sqrt{N/L(\varepsilon)}$, where $L(\varepsilon) = N + \varepsilon^2$ exactly (since $\text{Re} \sum_k z_k \overline{u_k} = 0$ for the radial $m=3$ mode and $\|u\|^2 = 1$).

Because the logarithmic energy is a sum of $-\ln|\cdot|$ over $N(N-1)/2$ pairs, the overall scale factor $s(\varepsilon) = \sqrt{N/(N+\varepsilon^2)}$ contributes an additive term:

$$H_{\text{constrained}}(\varepsilon) = H_{\text{unconstrained}}(\varepsilon) + \frac{N(N-1)}{4} \ln\left(1 + \frac{\varepsilon^2}{N}\right). \quad (4.2)$$

Expanding the logarithm: $\ln(1 + \varepsilon^2/N) = \varepsilon^2/N - \varepsilon^4/(2N^2) + O(\varepsilon^6)$, so the 4th derivative of the correction term at $\varepsilon = 0$ is

$$\frac{N(N-1)}{4} \cdot \left(-\frac{12}{N^2}\right) = -\frac{3(N-1)}{N}.$$

For $N = 7$ this is $-18/7$. Combined with the unconstrained value $153/7$, the constrained quartic is therefore

$$\left.\frac{d^4 H}{dt^4}\right|_{\text{constrained}} = \frac{153}{7} - \frac{18}{7} = \frac{135}{7}, \quad (4.3)$$

an exact rational result. Richardson extrapolation (three two-point pairs at $h \in \{0.05, 0.04, 0.03, 0.02\}$, symmetric fourth-difference formula $[2\delta H(2h) - 8\delta H(h)]/h^4 \rightarrow 24c_4$ with log-ratio cancellation) gives 19.2857 ± 0.0005 , numerically consistent with the exact value $135/7 = 19.285714\dots$

[Kurakin and Yudovich \[2002\]](#) established nonlinear stability of the heptagonal ring using Arnold's energy-Casimir method, proving positivity of the quartic form by Dirichlet's criterion without computing its value. [Cabral and Schmidt \[1999\]](#) and [Schmidt \[2004\]](#) independently proved stability via Birkhoff normal form and computed a quartic coefficient of $9/28$ in the *Sylvester-normalized* coordinates (where all quadratic terms are rescaled to $1/2$). The present value $\alpha_0 = 45/14$ is expressed in the physical (un-normalized) radial-mode coordinates on the constraint surface; both values confirm $\alpha_0 > 0$ but measure different quantities. The present derivation proceeds via the constrained energy-minimum framework (Theorem 2.2): the log-scaling identity (4.2) yields the exact correction $-3(N-1)/N$, and the unconstrained value $153/7$ is established by the Galois rationality argument combined with direct algebraic summation. In the normal form $\mathcal{H} = (\alpha_0/4)|z|^4$, the exact constrained quartic gives

$$\alpha_0 = \frac{135/7}{6} = \frac{45}{14} \approx 3.214 > 0: \quad (4.4)$$

the energy-Casimir functional $\mathcal{H} = H + |\mu_L|L$ has a strict local minimum along the $m = 3$ direction at quartic order, $\mathcal{H} \geq \mathcal{H}_0 + (\alpha_0/4)\varepsilon^4$. By the Dirichlet stability criterion [[Boatto and Cabral, 2003](#)], the heptagonal configuration is nonlinearly stable. The quartic (rather than quadratic) confinement explains why heptagonal rings are stable but more easily perturbed than hexagonal ones.

5 Multi-ring configurations and the center-ring dichotomy

Proposition 5.1 (Center-ring dichotomy). *Let N outer vortices sit at the regular N -gon on a ring of radius R_2 , and let $M \geq 2$ inner vortices sit at the regular M -gon on a ring of radius $R_1 < R_2$.*

(i) If $N \leq 7$ (outer ring Thomson-stable), the two-ring configuration is unstable: the Lagrangian Hessian has at least one negative eigenvalue.

(ii) If $N \geq 8$ (outer ring Thomson-unstable), the inner ring can reduce but not eliminate the instability: the two-ring Hessian retains at least one negative eigenvalue.

In contrast, a single central vortex ($M = 0$) at the $\mathbb{Z}_N$ fixed point produces a uniform eigenvalue shift $+\kappa_0/R^2$ across all Fourier modes and can achieve full stabilization (Proposition 3.1).

Proof. The proof splits into two cases.

Case 1: $N \geq 8$ (outer ring Thomson-unstable). The Havelock eigenvalue $\lambda_{\lfloor N/2 \rfloor} = (N-1) - \lfloor N/2 \rfloor(N - \lfloor N/2 \rfloor)/2 < 0$, so the outer N -ring alone has a negative constrained eigenvalue. Adding $M \geq 2$ inner vortices introduces new degrees of freedom but does not remove the existing instability: the negative outer mode persists (its eigenvalue is a continuous function of the coupling strength, starting negative at zero coupling). This proves (ii).

Case 2: $N \leq 7$ (outer ring Thomson-stable). The two concentric regular polygons form an equilibrium at the co-rotation ratio

$$R_2/R_1 = \sqrt{(N+M-1)/(M-1)}, \quad (5.1)$$

determined by the Gauss mean-value theorem: the outer ring at $R_2 > R_1$ encloses zero inner-ring circulation, so the inner ring's angular velocity is $(M-1)\kappa/(4\pi R_1^2)$ (self-interaction only), while the outer ring sees the enclosed inner circulation and rotates at $(N+M-1)\kappa/(4\pi R_2^2)$. Equating gives (5.1).

At this equilibrium, the constrained Lagrangian Hessian $\mathcal{H}$ is computed for all 12 cases ($N \in \{5, 6, 7\}$, $M \in \{2, \dots, N-1\}$). Non-positive-definiteness is certified by two algebraic tests:

- **Determinant test** (8 cases with odd number of negative eigenvalues): $\det(\mathcal{H}) < 0$ implies an odd number of negative eigenvalues, hence ≥ 1 .
- **Sylvester minor test** (4 remaining cases with even number of negative eigenvalues): the k -th leading principal minor of $\mathcal{H}$ is negative for some $k \leq \dim(\mathcal{H})$, so $\mathcal{H}$ is not positive semi-definite.

N	M	$\lambda_{\min}$	n_{neg}	Certificate
5	2	-0.375	3	$\det < 0$
5	3	-1.410	3	$\det < 0$
5	4	-3.085	4	$\text{minor}(7) < 0$
6	2	-0.204	2	$\text{minor}(8) < 0$
6	3	-0.795	3	$\det < 0$
6	4	-1.688	4	$\text{minor}(9) < 0$
6	5	-3.136	5	$\det < 0$
7	2	-0.088	1	$\det < 0$
7	3	-0.316	3	$\det < 0$
7	4	-0.803	3	$\det < 0$
7	5	-1.691	5	$\det < 0$
7	6	-3.012	6	$\text{minor}(9) < 0$

Both tests are verified in *exact rational arithmetic* (sympy 1.14, all matrix entries converted to $\mathbb{Q}$ via `nsimplify`; determinants computed symbolically with no floating-point rounding). The signs are algebraically exact: there is no numerical error to bound. Code: `two_ring.py` (numerical), `symbolic_two_ring_proof.py` (exact).

Extension to all $M \geq 2$. The 12 exact cases cover $M \in \{2, \dots, N-1\}$. The result extends to all $M \geq N$ by the following argument. The 2×2 radial Lagrangian Hessian (the breathing modes of the two rings) has off-diagonal entry $H_{12} = \partial^2 H / (\partial R_1 \partial R_2)$ from the attractive inter-ring interaction. At the co-rotation equilibrium, each ring individually has positive radial stiffness, but the shared Lagrange multiplier μ_L satisfies $\mu_{\text{outer}} < \mu_L < \mu_{\text{inner}}$ (the inner ring's natural frequency exceeds the outer's by the ratio $(N-1)/(N+M-1) < 1$). The resulting determinant of the 2×2 radial block satisfies $\det = (H_{11} - 2M\mu)(H_{22} - 2N\mu) - H_{12}^2$. Since H_{12} is proportional to MN (each inner-outer pair contributes) and the diagonal terms scale as M^2 and N^2 , the off-diagonal dominates for $M \geq 2$: $\det < 0$, giving at least one negative eigenvalue. This is confirmed by the 12 exact cases ($\det < 0$ in all 12) and extends to $M \geq N$ by continuity in M (the determinant is a rational function of M with no sign change for $M \geq 2$).

Mechanism. A central vortex ($M = 0$, the $\mathbb{Z}_N$ fixed point) contributes a uniform eigenvalue shift $+\kappa_0/R^2$ across all Fourier modes, compensating the negative Havelock eigenvalues. An inner ring of $M \geq 2$ cannot achieve this: the inter-ring coupling H_{12} breaks the uniformity, and the Lagrange-multiplier frustration ($\mu_{\text{outer}} \neq \mu_{\text{inner}}$) forces a negative eigenvalue. $\square$

Jupiter's observed center-plus-octagon geometry is therefore *uniquely stabilized* among all configurations with a secondary vortex structure: only the single central cyclone achieves full stabilization of the $N = 8$ ring.

6 From point vortices to continuous vorticity: the blob bridge

The results of Sections 1.1–Paper III, §21 are stated for point vortices. The RG classification (Proposition Paper I, Proposition 2.1) guarantees that compact vorticity distributions interact logarithmically at large separation, with irrelevant corrections $O(\varepsilon^2/d^2)$. Two distinct bridges make this quantitative.

6.0.1 The vortex blob bridge (Jupiter regime)

For discrete cyclone configurations ($\beta R^3/\kappa \ll 1$), replace each point vortex with a smooth blob of width ε . The regularized interaction $h_\varepsilon(d) = -\frac{1}{2} \ln(d^2 + 2\varepsilon^2)$ recovers $-\ln d$ as $\varepsilon \rightarrow 0$.

Proposition 6.1 (Blob stability correction). *The constrained Lagrangian eigenvalues for h_ε satisfy*

$$\lambda_m(\varepsilon) = \lambda_m(0) + \left[-\frac{N^2-1}{12} + P_m\right] \varepsilon^2 + O(\varepsilon^4), \quad (6.1)$$

where $\lambda_m(0)$ is the Havelock eigenvalue (Theorem 2.2).

Proof. Lagrange multiplier shift. The stationarity condition $\partial H_\varepsilon / \partial x_0 = \mu_L \cdot 2x_0$ gives $\mu_L(\varepsilon) = \mu_L(0) + \delta\mu_L \cdot \varepsilon^2$ where $\delta\mu_L = (N^2-1)/24$, using $\sum_{k=1}^{N-1} \csc^2(\pi k/N) = (N^2-1)/3$. The shift $-2\delta\mu_L = -(N^2-1)/12$ is *universal* (the same for all Fourier modes).

Hessian correction. The $O(\varepsilon^2)$ corrections to h'' and h'/d are: $\delta(h''(d)) = -6\varepsilon^2/d^4 + O(\varepsilon^4)$, $\delta(h'(d)/d) = 2\varepsilon^2/d^4 + O(\varepsilon^4)$. For a pair at distance $d_p = 2\sin(\pi p/N)$, the Hessian correction per pair is $O(\varepsilon^2/d_p^4)$. The projection $P_m = v_m^\top \delta \nabla^2 H v_m$ of this correction onto the critical eigenvector v_m is computed by summing over all $N - 1$ pairwise distances with the Fourier mode weights.

Operator norm bound. Since $d_{\min} = 2\sin(\pi/N) > 0$:

$$\|\nabla^2 H_\varepsilon - \nabla^2 H_0\|_{\text{op}} \leq 6(2N-3)\varepsilon^2/d_{\min}^4 = O(\varepsilon^2). \quad (6.2)$$

By Weyl's perturbation theorem, each eigenvalue shifts by at most this amount, giving uniform convergence at $O(\varepsilon^2)$.

Profile independence. For a blob with finite second moment σ^2 (Gaussian, compact support), the Fourier transform $|\hat{\phi}(\varepsilon k)|^2 = 1 - \sigma^2 \varepsilon^2 |k|^2 + O(\varepsilon^4 |k|^4)$ shows the $O(\varepsilon^2)$ correction is a constant in k -space (= self-energy, not affecting pairwise eigenvalues). The pairwise correction is $O(\varepsilon^4)$ —*faster* than the Cauchy blob. Since the Cauchy blob is the slowest-converging regularization, all profiles with finite σ^2 inherit its stability conclusions. $\square$

The Hessian projection P_m and the total correction $c_m = -(N^2-1)/12 + P_m$ (specific to the Cauchy blob $h_\varepsilon = -\frac{1}{2}\ln(d^2+2\varepsilon^2)$) are exact rational numbers, computed as follows. Differentiating the Cauchy blob pair Hessian with respect to ε^2 at $\varepsilon = 0$ gives the per-pair correction:

$$\delta H_{jj,xx} = \frac{2}{d_{jk}^6}((\Delta y)^2 - 3(\Delta x)^2), \quad \delta H_{jk,xx} = -\delta H_{jj,xx}, \quad (6.3)$$

with analogous expressions for yy and xy (see proof of Proposition 6.1). Here $(\Delta x, \Delta y, d_{jk})$ are the pair separation components. The critical radial Fourier mode is $v_k \propto z_k \cos(2\pi m k/N)$ with $m = \lfloor N/2 \rfloor$. The projection

$$P_m = v_m^\top \delta \nabla^2 H v_m \quad (6.4)$$

is a finite quadratic form: a sum over $N(N-1)$ pair contributions, each involving trigonometric values at N -th roots of unity. Since $\sin(\pi p/N)$ and $\cos(2\pi p/N)$ are algebraic for integer p and N , every term is rational and so is the sum. Explicit values (verified to machine precision in `blob_correction.compute_Pm(N)` and notebook `05.BlobCorrection.ipynb`):

N	3	4	5	6	7	8
$-(N^2-1)/12$	$-\frac{2}{3}$	$-\frac{5}{4}$	-2	$-\frac{35}{12}$	-4	$-\frac{21}{4}$
P_m	$-\frac{1}{3}$	$+\frac{1}{4}$	$+1$	$+4$	$+8$	$+\frac{65}{4}$
c_m	-1	-1	-1	$+\frac{13}{12}$	$+4$	$+11$

The Hessian projection P_m changes sign between $N = 3$ ($P_3 = -1/3$) and $N = 4$ ($P_4 = +1/4$): for the triangle, the blob softening reduces the nearest-neighbor curvature (all pairs are at the same short distance $d = \sqrt{3}$), while for the square and larger polygons, the long-range pairs increasingly dominate the projection and $P_m > 0$. Despite this sign change, $c_m = -1$ universally for $N \leq 5$ because the Lagrange multiplier shift $-(N^2-1)/12$ dominates P_m in this range. For $N \geq 6$, P_m grows faster and overwhelms the shift.

For blob profiles with finite second moment σ^2 (Gaussian, compact support), the $O(\varepsilon^2)$ correction is confined to the self-energy and does not affect pairwise eigenvalues (Section 6, profile independence argument). The pairwise correction is $O(\varepsilon^4)$ with a

profile-dependent coefficient. Since the $O(\varepsilon^2)$ corrections from the Cauchy blob table above preserve stability for $N \leq 7$, and all finite- σ^2 blobs have *smaller* corrections at any given ε , the stability conclusions hold *a fortiori* for all smooth profiles.

Three consequences follow:

(i) *Stability preserved.* For $N \leq 5$: $\lambda_0 \geq 1$ and $c = -1$, so $\lambda(\varepsilon) = \lambda_0 - \varepsilon^2 > 0$ for $\varepsilon < 1$. For $N = 6$: $c > 0$, so $\lambda(\varepsilon) > \lambda_0$ —the blob *enhances* stability.

(ii) *Marginal mode resolved.* At $N = 7$, $\lambda_0 = 0$ and $c = +4$, so $\lambda(\varepsilon) = 4\varepsilon^2 > 0$ for any $\varepsilon > 0$. The heptagonal marginal stability is an artifact of the singular point vortex limit.

(iii) *Instability partially compensated.* For $N = 8$: $\lambda_0 = -1$ and $c = 11$, so the octagon is stabilized for $\varepsilon > 1/\sqrt{11} \approx 0.30$.

6.0.2 Full infinite-dimensional convergence

Proposition 6.1 addresses center-position perturbations of N smooth blobs. The full vorticity perturbation space is infinite-dimensional, including blob *shape* deformations. A complete stability analysis requires the Hessian of the 2D Euler energy $E[\omega] = \frac{1}{2} \iint \omega(x) G(x, y) \omega(y) dx dy$ at the blob configuration $\omega_\varepsilon = \sum_k \kappa \phi_\varepsilon(\cdot - z_k)$, decomposed into three sectors:

(a) *Center translations.* The Hessian restricted to the $2N$ -dimensional subspace of center displacements is *exactly* the blob Hessian of Proposition 6.1. Convergence to the point vortex Hessian at $O(\varepsilon^\alpha)$ follows from the convolution estimates.

(b) *Shape perturbations.* Mean-zero, center-preserving deformations of each blob's profile enter through the self-energy $E_{\text{self}} = \iint \phi_\varepsilon(x) G(x, y) \phi_\varepsilon(y) dx dy$. The second variation at mean-zero $\delta\omega$ is, by Plancherel and the Fourier representation $\hat{G}(\xi) = 1/(4\pi^2|\xi|^2)$ of the 2D log-Green's function:

$$\delta^2 E_{\text{self}} = \iint \delta\omega(x) G(x, y) \delta\omega(y) dx dy = \int_{\mathbb{R}^2} \frac{|\hat{\delta\omega}(\xi)|^2}{4\pi^2|\xi|^2} d\xi = \frac{1}{4\pi^2} \|\delta\omega\|_{H^{-1}}^2. \quad (6.5)$$

Since $\hat{\delta\omega}(0) = \int \delta\omega = 0$ (mean-zero), the integrand is zero only on a set of measure zero, so $\delta^2 E_{\text{self}} > 0$ for all non-zero mean-zero $\delta\omega$: the self-energy Hessian is positive definite in H^{-1} . The rescaling $\phi_\varepsilon(x) = \varepsilon^{-2} \phi(x/\varepsilon)$ gives $\hat{\phi}_\varepsilon(\xi) = \hat{\phi}(\varepsilon\xi)$, so

$$E_{\text{self}}[\phi_\varepsilon] = \int \frac{|\hat{\phi}(\varepsilon\xi)|^2}{4\pi^2|\xi|^2} d\xi = \varepsilon^{-2} \int \frac{|\hat{\phi}(\eta)|^2}{4\pi^2|\eta|^2} d\eta = \frac{c}{\varepsilon^2},$$

and the smallest shape eigenvalue scales as c/ε^2 . Shape modes are *unconditionally stable*.

(c) *Center-shape coupling.* The cross terms involve overlap integrals

$$C_{jk} = \iint \nabla \phi_\varepsilon(x - z_j) \cdot \nabla G(x, y) \delta\psi(y - z_k) dx dy$$

between non-adjacent blobs $j \neq k$ ($|z_j - z_k| \geq d_{\min} \gg \varepsilon$), where $\delta\psi$ is a mean-zero, center-preserving shape perturbation. Since both the monopole ($\int \delta\psi = 0$) and dipole ($\int y \delta\psi = 0$) moments of $\delta\psi$ vanish, the leading term in the multipole expansion of $G(x, y)$ around z_k is the quadrupole:

$$G(x, y) = G(x, z_k) + \nabla_y G(x, z_k) \cdot (y - z_k) + O(|y - z_k|^2/d^2), \quad d = |x - z_k|.$$

The first two terms contribute zero (by the vanishing moments), so the coupling is determined by the quadrupole interaction: $C_{jk} = O(\varepsilon^2/d_{\min}^3)$. (Numerically verified:

$C_{jk} \approx 4\varepsilon^2/d_{\min}^3$ for the Cauchy blob at $d_{\min} \geq 3\varepsilon$.) This is *algebraically small*—the correct bound for the Cauchy blob, which has algebraic not exponential spatial decay. The ratio of coupling to shape eigenvalue is $O(\varepsilon^4/d_{\min}^3) \rightarrow 0$ as $\varepsilon \rightarrow 0$ with $d_{\min}$ fixed, which is the only bound needed below.

Proposition 6.2 (Inertia preservation). *For sufficiently small $\varepsilon > 0$, the full Hessian of $E[\omega_\varepsilon] + \mu_L L[\omega_\varepsilon]$ on the space of $\mathbb{Z}_N$ -symmetric vorticity perturbations has the same inertia (number of positive, zero, and negative eigenvalues along the constrained directions) as the point vortex Lagrangian Hessian of Theorem 2.2. The point vortex model is therefore the exact limit of smooth vorticity dynamics for stability analysis: no eigenvalue is lost, no sign is changed, and no new instability arises from the infinite-dimensional shape degrees of freedom. The stability boundary $N \leq 7$ is not an approximation of the smooth result—it is the smooth result in the $\varepsilon \rightarrow 0$ limit.*

Proof. Write the constrained full Hessian in block form:

$$\mathcal{H} = \begin{pmatrix} H_{cc} & H_{cs} \\ H_{cs}^\top & H_{ss} \end{pmatrix},$$

where the subscripts c and s denote the center and shape sectors.

Shape sector. By (6.5), $H_{ss} \geq (c/\varepsilon^2) \text{Id}_{H^{-1}}$ with $c > 0$. All shape eigenvalues are bounded below by $c/\varepsilon^2 > 0$.

Center sector. By (a) and the convolution estimates of Proposition 6.1, the center-sector eigenvalues satisfy $\lambda_c^{(k)}(\varepsilon) = \lambda_k^{\text{Hav}} + O(\varepsilon^2)$, where λ_k^{Hav} are the Havelock eigenvalues.

Coupling bound. From (c), the off-diagonal block satisfies $\|H_{cs}\|_{\text{op}} = O(\varepsilon^2/d_{\min}^3)$. By the Schur complement, the full Hessian has the same inertia as $H_{cc} - H_{cs}H_{ss}^{-1}H_{cs}^\top$. The correction term is

$$\|H_{cs}H_{ss}^{-1}H_{cs}^\top\|_{\text{op}} \leq \|H_{cs}\|_{\text{op}}^2 \|H_{ss}^{-1}\|_{\text{op}} = O\left(\frac{\varepsilon^4}{d_{\min}^6}\right) \cdot O(\varepsilon^2) = O\left(\frac{\varepsilon^6}{d_{\min}^6}\right),$$

which goes to zero as $\varepsilon \rightarrow 0$ with $d_{\min}$ fixed. Hence, for sufficiently small ε , the Schur complement eigenvalues are within $O(\varepsilon^6/d_{\min}^6)$ of the center-sector eigenvalues H_{cc} , and in particular preserve their signs. Since $\|H_{cc} - H_{cc}^{\text{Hav}}\|_{\text{op}} = O(\varepsilon^2)$ (from the blob correction), for small enough ε the Schur-complement eigenvalues have the same signs as the Havelock eigenvalues. Therefore the full Hessian inertia equals the Havelock inertia—which is the point vortex Lagrangian inertia of Theorem 2.2. $\square$

Proposition 6.2 extends to *general* $\mathbb{Z}_N$ -symmetric vorticity ω concentrated in N patches (not necessarily a prescribed blob profile). Decompose $\omega = \sum_k \omega_k$ via Voronoi cells, split each perturbation into center-translation and shape components $\delta\omega_k = \delta\omega_k^{(c)} + \delta\omega_k^{(s)}$ with $\delta\omega_k^{(c)} = -\delta\bar{z}_k \cdot \nabla\omega_k$. Integration by parts gives $\int x \delta\omega_k^{(s)} = 0$ (dipole cancellation), so the center-shape coupling is quadrupole: $|C_{jk}| = O(\varepsilon^2/d_{\min}^2)$ for compactly supported profiles, $O(\varepsilon^2/d_{\min}^3)$ for the Cauchy blob. The self-energy Hessian is positive definite for any mean-zero profile, and $\mathbb{Z}_N$ symmetry forces congruent patches. The phrase *sufficiently concentrated* is made quantitative as follows. The operator norm bound (6.2) is derived for the Cauchy blob, which has algebraic ($|x|^{-3}$) rather than exponential spatial decay and is therefore the slowest-converging regularization among profiles with finite second moment; profiles with faster decay (Gaussian, compact support) satisfy the same bound *a fortiori*, with a smaller implicit constant. The Cauchy-blob bound gives

$\|\nabla^2 H_\varepsilon - \nabla^2 H_0\|_{\text{op}} \leq 6(2N-3)\varepsilon^2/d_{\min}^4$. Sign preservation of the Havelock eigenvalues requires this to be less than $\lambda_{\min}(N)$ (the smallest positive Havelock eigenvalue at the configuration), yielding the explicit concentration condition

$$\frac{\varepsilon}{d_{\min}} < \delta(N) := \sqrt{\frac{\lambda_{\min}(N)}{6(2N-3)}}, \quad (6.6)$$

where $d_{\min} = 2\sin(\pi/N)$ is the nearest-neighbor separation (unit ring). Numerically:

N	4	5	6	7
$\lambda_{\min}(N)$	1	1	$\frac{1}{2}$	0
$\delta(N)$	0.18	0.15	0.096	—

The $\delta(N)$ values are set by the center-sector bound, not the coupling. At $\varepsilon/d_{\min} = \delta(N)/\sqrt{2}$, the center-sector eigenvalue shift is at most $\lambda_{\min}/2$ (from (6.6)), leaving each center eigenvalue $\geq \lambda_{\min}/2$. The Schur correction at this ε is $O(\delta^6/8) = O(10^{-6})$, which is smaller than $\lambda_{\min}/2$ by a factor of at least 10^{-5} —so the Schur term cannot reverse any eigenvalue sign, and the full inertia preservation holds with the explicit bound $\varepsilon/d_{\min} < \delta(N)/\sqrt{2}$.

Corollary 6.3 (Inertia preservation for $N \leq 6$). *For $N \leq 6$ and any $\mathbb{Z}_N$ -symmetric vorticity distribution with patch radius $\varepsilon/d_{\min} < \delta(N)$ (equation (6.6)), the full Hessian inertia equals the Havelock inertia. The argument uses only the spectral-gap condition $\lambda_{\min}(N) > 0$ and the Weyl perturbation bound; no information about the $O(\varepsilon^2)$ blob correction is needed.*

Corollary 6.4 (Marginal-mode resolution for $N = 7$). *For $N = 7$, the concentration bound (6.6) does not apply ($\lambda_{\min} = 0$, so $\delta(7) = 0$). A separate, logically distinct argument shows $\lambda_{\min}(\varepsilon) > 0$ for all $\varepsilon \in (0, d_{\min})$ via two regions:*

- Region 1 ($\varepsilon < \varepsilon_0 = 2/\sqrt{|c'_7|}$, where $|c'_7| \leq 17$ is the bounded $O(\varepsilon^4)$ coefficient): $\lambda \geq \varepsilon^2(4 - |c'_7|\varepsilon^2) > 0$, since the blob correction $c_7 = +4 > 0$ (Proposition 6.1) dominates.
- Region 2 ($\varepsilon \in [\varepsilon_0, d_{\min})$): interval arithmetic on a 2,000-point grid with mean-value-theorem bounds gives $\lambda \geq 0.34 > 0$ on every subinterval.

The two corollaries rest on different mechanisms: Corollary 6.3 on a spectral gap, Corollary 6.4 on a blob correction. Both extend beyond prescribed-profile blobs to general $\mathbb{Z}_N$ -symmetric vorticity patches satisfying their respective conditions.

6.0.3 The Rossby bridge (Saturn regime)

For continuous jet meanders ($\beta R^3/\kappa \gg 1$), the bridge is not vortex concentration but *Rossby wave quantization*. The QGPV linearized about a zonal jet $U(y)$ gives the Rayleigh–Kuo equation:

$$(U - c)(\hat{\psi}'' - k^2\hat{\psi}) + (\beta - U'')\hat{\psi} = 0.$$

For a stationary wave ($c = 0$) at the jet center ($y = 0$) of a Gaussian jet with $U''(0) = -U_{\max}/\sigma^2$:

$$\hat{\psi}'' + \left[\frac{\beta_{\text{eff}}}{U_{\max}} - \frac{n^2}{R^2} \right] \hat{\psi} = 0, \quad \beta_{\text{eff}} = \beta + \frac{U_{\max}}{\sigma^2},$$

where the effective PV gradient β_{eff} includes the jet curvature (which dominates β by a factor ~ 13 for Saturn). The stationary azimuthal wavenumber is

$$n^* = R_S \cos \varphi \sqrt{\beta_{\text{eff}}/U_{\max}}, \quad (6.7)$$

where $R_S \cos \varphi$ is the ring radius at latitude φ ($\approx 1.31 \times 10^7$ m for Saturn at 76° N). The nearest integer $n = \text{round}(n^*)$ is the selected polygon wavenumber. For Saturn: $\beta_{\text{eff}}/U_{\max} \approx 1.72 \times 10^{-13} \text{ m}^{-2}$, giving $n^* \approx 5.4$ from the leading-order dispersion relation. The observed $n = 6$ is obtained in numerical eigenvalue computations of the full Rayleigh–Kuo equation [Sánchez-Lavega et al., 2014]; the discrepancy between the leading-order $n^* \approx 5.4$ and the observed $n = 6$ reflects corrections from the finite jet width (the meridional eigenfunction structure contributes an $O(1/\sigma^2)$ correction to the effective wavenumber that is not captured by the dispersion relation alone).

The discretization mechanism is *azimuthal periodicity*: the wave must fit an integer number of wavelengths around the ring.

The connection to the point vortex energy landscape proceeds through the *critical layer* of the saturated wave. At finite amplitude, the stationary Rossby wave develops a cat’s-eye structure at the critical line $U = c = 0$: vorticity is trapped into N lobes of width $\sim \sigma$ centered at the meander peaks. As the jet sharpens ($\sigma/R \rightarrow 0$ at fixed wave amplitude), these lobes concentrate into N isolated patches near the N -gon vertices.

Remark 6.1 (Cat’s-eye Hessian convergence). *The cat’s-eye vorticity decomposes as $\omega = \omega_{\text{lobe}} + \omega_{\text{braid}} + \omega_{\text{bg}}$, where the lobes are the trapped fluid inside the separatrices, the braids are thin filaments along the separatrix itself, and the background is the passing fluid outside the cat’s-eye.*

Two conditions are required for Proposition 6.2 to apply:

(i) **Lobe concentration.** *The lobes satisfy the concentration condition (6.6) with $\varepsilon \sim \sigma$ and the centroids converge to the N -gon vertices. This is guaranteed by $\mathbb{Z}_N$ symmetry.*

(ii) **Braid smallness.** *The absolute braid circulation Γ_{braid} is negligible compared to the total lobe circulation Γ_{lobe} .*

Both conditions hold for the Stuart vortex $\Psi = -\sigma^2 \ln(\cosh(y/\sigma) - \eta \cos(n\theta))$ ($\eta \in (0, 1)$; $\eta \rightarrow 1$ is point-vortex concentration).

Proof of lobe dominance. The Stuart vortex satisfies $\omega = C/F^2$ with $F = \cosh(y/\sigma) - \eta \cos(n\theta)$. Near the elliptic center ($\theta = 0$, $y = 0$): $F \approx (1-\eta) + (y/\sigma)^2/2 + n^2\eta\theta^2/2$, so

$$\Gamma_{\text{lobe}} = C \iint_{\text{lobe}} \frac{dA}{F^2} \geq C \int_0^\infty \frac{2\pi r dr}{[(1-\eta) + r^2/2]^2} = \frac{2\pi C}{1-\eta}.$$

Outside the separatrix, $F \geq 1+\eta \geq 1$, so $\Gamma_{\text{non-lobe}} = C \iint_{\text{outside}} F^{-2} dA \leq C A_{\text{period}}$, which is bounded independently of η . Therefore

$$\frac{\Gamma_{\text{non-lobe}}}{\Gamma_{\text{total}}} \leq \frac{A_{\text{period}}(1-\eta)}{2\pi + A_{\text{period}}(1-\eta)} = O(1-\eta). \quad (6.8)$$

The lobe fraction $\Gamma_{\text{lobe}}/\Gamma_{\text{total}} = 1 - O(1-\eta)$ as $\eta \rightarrow 1$. Numerically: at $\eta = 0.99$, lobes capture $> 99\%$ of the total circulation.

Combining with Proposition 6.2: the cat’s-eye satisfies the blob bridge conditions for $\sigma/d_{\min} < \delta(N)/\sqrt{2}$ and $1 - \eta < c\delta(N)^2$ (the braid perturbation is $O(1-\eta)$ of the lobe system, negligible at the same threshold). The constrained Hessian inertia of the cat’s-eye equals the Havelock inertia, verified numerically for all $\eta \in [0.3, 0.99]$ and confirmed analytically by (6.8) in the limit $\eta \rightarrow 1$. The same argument applies to the Prandtl–Batchelor cat’s-eye (homogenised PV inside the lobes), the physically relevant model for Saturn’s hexagon. Code: `catseye_decomposition.py`, `catseye_hessian.py`.

The Saturn bridge thus arrives at the *same* energy landscape as the Jupiter bridge, through a different limiting procedure: Saturn concentrates a continuous jet meander into N patches (jet width $\sigma \rightarrow 0$), while Jupiter concentrates N discrete cyclones into

7 Finite Rossby deformation radius: the K_0 interaction

The K_0 interaction provides a second example of the RG flow toward the logarithmic fixed point (Proposition Paper I, Proposition 2.1): the exponential tail $K_0(r/R_d) + \ln r = O(e^{-r/R_d})$ is an irrelevant perturbation (faster than any power law), and the limit $R_d \rightarrow \infty$ is the flow to the fixed point.

In quasi-geostrophic dynamics with finite Rossby deformation radius R_d , the Green’s function is $K_0(r/R_d)$ (modified Bessel function), which recovers $-\ln r$ for $r \ll R_d$. The interaction $h(r) = K_0(r/R_d)$ satisfies $h'(r) = -K_1(r/R_d)/R_d < 0$ and $(rh'(r))' = rK_0(r/R_d)/R_d^2 > 0$ (class (iii)). Numerical computation of the constrained Lagrangian Hessian as a function of $\eta = R/R_d$ (polygon ring radius over deformation radius) reveals an N -dependent *polygon-suppression transition*:

N	Barotropic $\lambda_{\min}$	Critical R/R_d	Status
≤ 5	≥ 1	none	Stable for all R/R_d
6	$1/2$	≈ 1.29	Suppressed for $R/R_d > 1.29$
7	0	≈ 0	Unstable for any finite R/R_d
≥ 8	< 0	N/A	Unstable (barotropic)

Polygons with $N \leq 5$ are robust against stratification effects. The hexagon ($N = 6$) is the first polygon sensitive to R_d : it requires $R/R_d < 1.29$. The heptagon ($N = 7$), neutrally stable in the barotropic limit, is destabilized by any finite R/R_d —providing an additional mechanism for its rarity beyond the quartic confinement.

For Saturn, the barotropic deformation radius (~ 2000 – 3000 km) gives $R/R_d \approx 8$ – 12 at the hexagonal jet, placing it above the critical ratio. The observed hexagonal persistence therefore requires the dynamics to be effectively barotropic over a deep layer—a requirement now directly supported by observational data. Cassini Grand Finale gravity measurements constrain Saturn’s differential rotation to a depth of $8,000$ – $9,000$ km [Iess et al., 2019], implying a coherent barotropic column of order 10^4 km, far exceeding the hexagon’s vertical scale. Saturn is therefore the best-supported case in our sample for the applicability of the 2D logarithmic Green’s function: the columnar depth is observationally established, not merely assumed.

Observable prediction: hexagon suppression signatures. The critical ratio $R/R_d \approx 1.29$ for $N = 6$ yields a testable prediction. If the effective barotropic depth were reduced (by seasonal radiative forcing of the upper troposphere, or by a hypothetical change in deep convective coupling), R_d would decrease and R/R_d would increase toward the critical value. Approaching the threshold, the framework predicts three observable signatures: (1) the minimum Havelock eigenvalue $\lambda_{\min}(\eta) \rightarrow 0$, reducing the hexagonal mode’s restoring frequency and producing a *longer oscillation period* for meridional meander perturbations; (2) the hexagonal amplitude ε_{obs} would *grow*, since the restoring force weakens while convective forcing remains constant; and (3) at $R/R_d > 1.29$, the hexagonal mode becomes linearly unstable and the pattern would transition to a lower- N polygon ($N = 5$, the next Thomson-stable configuration) or break up entirely. Saturn’s hexagon has been observed continuously since Voyager 1 (1980–81); Cassini monitored it through northern spring equinox (2009) to summer solstice (2017) with no sign of weakening [Antuñano et al., 2015, Sánchez-Lavega et al., 2014]. This four-decade persistence across half a Saturnian year places an empirical lower bound on the effective barotropic depth: seasonal radiative forcing does not reduce R_d enough to approach the critical ratio, consistent with the deep ($\sim 10^4$ km) barotropic column implied by Cassini gravity.

Why the leading K_0 correction does not shift stability. The small- r expansion $K_0(r/R_d) = -\ln(r/R_d) - \gamma_E + r^2/(4R_d^2) + O(r^4)$ shows the leading correction to the logarithm is $g(r) = r^2/(4R_d^2)$ —a perturbation of RG exponent $\alpha = 2$. By the isotropic-marginality result (Paper I, Proposition 2.2), the pair Hessian anisotropy $g'' - g'/r = \alpha(\alpha-2)r^{\alpha-2}$ vanishes at $\alpha = 2$: the Hessian perturbation is purely isotropic ($\propto I$) and is exactly compensated by the rotation-frequency shift $\delta\Omega$. Consequently $\delta\lambda_m = 0$ for all binding modes—the $r^2/(4R_d^2)$ term shifts the precession rate of the polygon but does *not* shift its stability eigenvalues. The first correction to stability from K_0 is $O(r^4/R_d^4)$ ($\alpha = 4$, which is *not* isotropic-marginal), and this is numerically small for $R/R_d \ll 1$. This explains why the barotropic ($R_d \rightarrow \infty$) stability predictions remain accurate even at moderate R/R_d : the leading deformation-radius correction is stability-neutral.

For Jupiter the situation is different, and the 2D justification requires separate examination. Juno gravity measurements [Kaspi et al., 2018] constrain Jupiter’s differential rotation to $\sim 3,000$ km depth ($\approx 4\%$ of R_J), far shallower than Saturn’s. Jupiter’s polar cyclones therefore do *not* rest on deep barotropic Taylor columns; the relevant 2D dynamics is that of a shallow weather layer above a near-solid-body rotating deep interior. Two limiting cases bracket the problem:

Deep weather-layer limit ($H_{\text{eff}} \sim 3,000$ km). Taking the full differential-rotation depth as the effective layer and using the equivalent-barotropic deformation radius $R_d = \sqrt{g H_{\text{eff}}}/f \approx \sqrt{25 \times 3 \times 10^6} / (2\Omega_J \sin 84^\circ) \approx 25,000$ km, we obtain $R/R_d \approx 0.3 \ll 1$. In this limit $K_0(r/R_d) \rightarrow -\ln r + \text{const}$, so the logarithmic Green’s function applies—though the physical justification is a shallow equivalent-barotropic layer rather than a deep Taylor column.

Tropospheric limit ($H_{\text{eff}} \sim 30\text{--}100$ km). If the active polar-cyclone layer is shallower, reduced gravity $g' \sim 0.1\text{--}1$ m s $^{-2}$ gives $R_d \sim 2,000\text{--}6,000$ km and $R/R_d \sim 1\text{--}4$. The $N = 8$ ring lies in the finite- R_d regime. Computing the κ_0 -shifted constrained Hessian eigenvalue for the K_0 interaction with $\kappa_0/\kappa = 0.7$ (the Juno central estimate):

N	κ_0/κ	$(R/R_d)_{\text{crit}}$	Tropospheric stable?
6	0.7	17.2	yes ($\gg 4$)
7	0.7	4.2	marginal
8	0.7	0.79	no (< 1)
9	0.7	0.10	no

Here $(R/R_d)_{\text{crit}}$ is the ratio at which the minimum constrained eigenvalue crosses zero: the ring is stable for $R/R_d < (R/R_d)_{\text{crit}}$ and unstable above. For $N = 8$ the critical ratio is $(R/R_d)_{\text{crit}} \approx 0.79$, which lies *below* the tropospheric range $R/R_d \in [1, 4]$. **The $N = 8$ prediction therefore requires the deep weather-layer limit** ($R/R_d \lesssim 0.8$, i.e. $R_d \gtrsim 1.3 R \approx 9,700$ km); it does not survive into the tropospheric regime.

The effective layer depth is not directly constrained by current Juno data, making Jupiter the *least* observationally supported case for the 2D assumption, in contrast to Saturn. A direct measurement of the depth of the polar cyclone roots—analogueous to the gravity constraints on the jets—would resolve this ambiguity. The computation above sets a quantitative target: $R_d > 9,700$ km is necessary for $N = 8$ stability under the K_0 interaction.

The critical ratio is defined implicitly by the equation $\bar{\lambda}(\eta_{\text{crit}}, N) = 0$, where the dimensionless eigenvalue $\bar{\lambda}(\eta, N) = R^2 \cdot \lambda_{\min}(\eta)$ interpolates smoothly from the Havelock value $\bar{\lambda}(0, N) = (N-1) - \lfloor N/2 \rfloor (N - \lfloor N/2 \rfloor)/2$ to zero.

The Lagrange multiplier is determined by the stationarity condition $\partial H/\partial R + \mu_L \partial L/\partial R = 0$. For $H = -\sum_{j < k} K_0(d_{jk}/R_d)$ with $d_{jk} = 2R \sin(\pi p/N)$ and $L = NR^2$, the radial derivative of each pair contributes $-K'_0(d_{jk}/R_d) \cdot 2 \sin(\pi p/N)/R_d = K_1(d_{jk}/R_d) \cdot 2 \sin(\pi p/N)/R_d$ (using $K'_0 = -K_1$). Projecting onto the constraint gradient $\partial L/\partial R = 2NR$ and solving gives

$$-2\mu_L = \frac{\eta}{R} \sum_{p=1}^{N-1} \frac{(1 - \cos \theta_p) K_1(2\eta \sin(\pi p/N))}{2 \sin(\pi p/N)}, \quad \theta_p = 2\pi p/N, \quad (7.1)$$

where $\eta = R/R_d$ and the factor $(1 - \cos \theta_p)$ arises from the Fourier decomposition of the pairwise Hessian along the radial mode. As $\eta \rightarrow 0$: $K_1(x) \sim 1/x$ for small x , so each summand reduces to $(1 - \cos \theta_p)/(4\eta \sin^2(\pi p/N)) \cdot (\eta/R) = (1 - \cos \theta_p)/(4R \sin^2(\pi p/N))$. Summing over p using the Havelock identity $\sum_p (1 - \cos \theta_p)/(4 \sin^2(\pi p/N)) = (N-1)/2$ recovers the flat-plane value $-2\mu_L = (N-1)/(2R^2)$. For $N = 6$, the critical mode $m = 3$ has zero Havelock weight at $p = 2, 4$ (since $\cos(2\pi \cdot 2 \cdot 3/6) = 1$), so the critical eigenvalue depends on K_0 and K_1 at only *two* distinct arguments, η and 2η , rather than three.

No closed-form analog of the Havelock identity is known for K_0 . The orthogonality mechanism that collapses the logarithmic sum to $m(N-m)/2$ relies on the identity $\sum_p e^{2\pi i m p/N} \ln |2 \sin(\pi p/N)| = 0$, which has no counterpart for K_0 : the Bessel values $K_0(2\eta \sin(\pi p/N))$ are transcendental at nonzero algebraic arguments by Siegel's theorem on E -functions, but this does not preclude the full weighted sum from having a closed form. Whether $\sum_p (\cos \theta_{pm}) K_0(2\eta \sin(\pi p/N))$ admits a closed form for any N and m is an open question; no such form has been found, and all stability computations for K_0 interactions are carried out numerically. A perturbative expansion $K_0(x) \approx -\ln x + O(x^2)$ gives $\eta_{\text{crit}} \approx 0.9$ for $N = 6$ —30% below the exact value—because the critical ratio lies in the non-perturbative regime where K_0 departs significantly from $-\ln$. The transition occurs precisely where the exponential decay of K_0 (destroying long-range correlations between non-adjacent vortices) overtakes the logarithmic regime, making the suppression

inherently non-perturbative. The implicit equation is nevertheless well-posed, and the critical η is uniquely determined by bisection for each N .

The two phase boundaries. For K_0 , two distinct boundaries exist in the (N, η) plane: (1) the *spiral sign change* $\bar{H}''(0) = 0$, where the positive A_m term from $(rh')' > 0$ overwhelms the negative B_m term; and (2) the *constrained stability loss* $\bar{\lambda}_{\min} = 0$, where the Lagrange multiplier shift consumes the stability margin. The constrained boundary ($\eta \approx 1.29$ for $N = 6$) is always *inside* the spiral boundary ($\eta \approx 3.2$ for $N = 6$) because the Lagrange shift reduces the margin before H'' changes sign. Physically, the polygon requires long-range correlations between non-adjacent vortices. When $\eta \gg 1$ and the K_0 interaction decays exponentially, only nearest-neighbor correlations survive and the polygon loses its energetic advantage.

For comparison, the Cauchy-blob regularization $-\frac{1}{2} \ln(d^2 + 2\varepsilon^2)$ has $H'' < 0$ for *all* ε —the soft-core never flips the spiral sign. The Gaussian-screened interaction $-\ln(d) \cdot e^{-d^2/L^2}$ shows an earlier transition ($\eta \approx 0.9$ for $N = 6$) due to its faster-than-exponential decay.

The constrained stability boundary. The interaction class determines N_{crit} , defined here as the largest N for which the *constrained* Lagrangian Hessian $\bar{\lambda}_{\min} \geq 0$ (i.e. the polygon is a constrained energy minimum). This is distinct from the spiral-direction boundary $H''(0) = 0$, which occurs at larger η (e.g. $\eta \approx 3.2$ vs. 1.29 for $N = 6$ with K_0 ; see the two phase boundaries above):

Class	Condition	N_{crit} (constrained)
(i)	$(rh')' < 0$	≥ 7 (increases with $ (rh')' $)
(ii)	$(rh')' = 0$ (only $-\ln r$)	$= 7$ exactly
(iii)	$(rh')' > 0$ (e.g. K_0)	≤ 7 (decreases with screening)

Numerical computation confirms: for $h = -r^\alpha$ (class (i)), $N_{\text{crit}} = 7$ at $\alpha \rightarrow 0$ and increases with α ($N_{\text{crit}} = 8$ at $\alpha = 1$, $N_{\text{crit}} = 14$ at $\alpha = 2$). At $N = 7$, $\lambda_{\min}(\alpha) > 0$ for all $\alpha > 0$ tested, with $d^2 \lambda_{\min} / d\alpha^2 = +2.5 > 0$ at $\alpha \rightarrow 0$: the marginal stability at $N = 7$ is *unique to the logarithmic interaction*, and any class (i) interaction makes the heptagon strictly stable. For $K_0(r/R_d)$ (class (iii)), the dimensionless eigenvalue $\bar{\lambda}(\eta)$ decreases monotonically in η for each Fourier mode (verified for $N = 5, 6$), consistent with the physical expectation that screening weakens the stabilizing long-range correlations. For $N_{\text{crit}} = 6$ at $R/R_d = 1$ and decreases to 5 at $R/R_d = 3$. The logarithmic interaction is the *watershed*: the critical case between interactions that stabilize polygons (class (i)) and interactions that suppress them (class (iii)).

Proposition 7.1 (Min/max duality). *For any pairwise interaction $h(r)$, the constrained energy minimum of $H = \sum h(d_{jk})$ and the constrained energy maximum of $-H$ have the same stability boundary N_{crit} .*

Proof. The Lagrange multiplier for $-H$ is $\mu_L^{-H} = -\mu_L^H$ (from $-\nabla H = (-\mu_L)\nabla L$). The Lagrangian Hessian $\mathcal{L}_{-H} = -\nabla^2 H + 2\mu_L I = -\mathcal{L}_H$. The eigenvalues are negated; constrained minimum (all ≥ 0) of H is equivalent to constrained maximum (all ≤ 0) of $-H$. $\square$

For gravity ($h = -1/r$, $h' > 0$), the N -gon is a constrained energy maximum. The duality gives $N_{\text{crit}}^{\text{gravity}} = N_{\text{crit}}^{\text{Coulomb}}$. Numerical computation of the Coulomb ($h = +1/r$) constrained Hessian gives $N_{\text{crit}} = 5$: the minimum eigenvalue is $+0.73$ at $N = 5$ and

-0.02 at $N = 6$. This is consistent with Maxwell’s (1859) result that gravitating rings require a central mass for $N \geq 7$; the one-unit difference ($N_{\text{crit}} = 5$ vs Maxwell’s 6) arises because our analysis does not include the central mass (Maxwell’s stability includes the stabilizing effect of the central body).

The Coulomb/gravity $N_{\text{crit}} = 5$ is less than the logarithmic $N_{\text{crit}} = 7$, confirming the class (iii) reduction: the $1/r$ interaction has $(rh')' = 1/r^2 > 0$, placing it in the destabilizing regime relative to $-\ln r$.

8 Circulation disorder and stability domain robustness

8.0.1 Resilience to circulation disorder

In both planetary regimes the vortex circulations κ_k are not perfectly equal; convection injects or removes circulation from individual vortices continuously. We model this as $\kappa_k = \kappa(1 + \eta_k)$ with i.i.d. fractional disorders $\eta_k \sim \mathcal{N}(0, \sigma_\eta^2)$.

We work on $\mathbf{H}^2$ rather than the flat plane because the hyperbolic setting simultaneously tests two effects: circulation disorder and curvature stabilization. In the flat limit $\rho/a \rightarrow 0$ the $\mathbf{H}^2$ eigenvalues recover the Havelock values exactly, so the flat-plane result is contained as a special case; conversely, the curvature-stabilized configurations ($N \geq 8$ on $\mathbf{H}^2$) have no flat-plane analog and can only be probed on the curved surface.

On $\mathbf{H}^2$, the disordered Hamiltonian is

$$H_\kappa = - \sum_{j < k} \kappa_j \kappa_k \ln |z_j - z_k| + \frac{1}{2} \sum_j \kappa_j (K_{\text{tot}} - \kappa_j) \ln(a^2 - |z_j|^2), \quad (8.1)$$

where $K_{\text{tot}} = \sum_k \kappa_k$. At $\eta_k \equiv 0$ this reduces exactly to the uniform Hamiltonian H_{hyp} .

Holding the ring positions at the uniform equilibrium (exact to $O(\eta)$; the $O(\eta)$ shift of equilibrium produces only an $O(\eta^2)$ eigenvalue correction), we compute the minimum Fourier-mode eigenvalue $\lambda_{\min}(\eta)$ via finite differences. Monte Carlo over $n = 1000$ realizations estimates $P(\lambda_{\min} < 0)$ as a function of (σ_η, N, ρ) .

Physical range of σ_η . Juno imagery [Adriani et al., 2018, Gavriel and Kaspi, 2021] shows Jupiter’s eight north-polar cyclones have peak wind speeds ranging from ~ 60 to $\sim 100 \text{ m s}^{-1}$ and diameters from $\sim 4,000$ to $\sim 7,000 \text{ km}$. Taking circulation $\kappa \propto V_{\text{max}} r_c$, the coefficient of variation across the eight cyclones is $\sigma_\eta \approx 0.15\text{--}0.25$. Saturn’s hexagonal jet is more uniform (Cassini wind profiles give $\sigma_\eta \lesssim 0.05$ for the azimuthal circulation variation). The physically relevant range is therefore $\sigma_\eta \in [0.05, 0.25]$.

Key findings (all uncertainties are $\pm 1\sigma$ binomial standard errors $\sqrt{P(1-P)/n}$ with $n = 1000$):

- **Stable configurations are robust.** For $N = 6$ at $\rho = 0.5a$, the instability probability remains below 10% ($\pm 0.9\%$) for $\sigma_\eta \leq 0.2$, covering the full physically relevant range.
- **Curvature-stabilized rings ($N \geq 8$) are similarly robust.** At $N = 8$, $\rho = 2a$ (well inside the stable region on $\mathbf{H}^2$), $P < 0.5$ ($\pm 1.6\%$) for $\sigma_\eta \leq 0.1$. At $\sigma_\eta = 0.2$ (Jupiter’s estimated value), P remains below 15%.

- **The $N = 7$ marginal ring is sensitive.** At $N = 7$, $\rho = 0.05a$ (flat-limit regime), $P(\lambda_{\min} < 0)$ increases monotonically with σ_η : $P \approx 0$ ($\pm 0.3\%$) at $\sigma_\eta = 0.1$ and $P > 0$ for $\sigma_\eta \geq 0.5$. The marginal mode ($\lambda_0 = 0$ at $K = 0$) is the first to be destabilized by disorder, as expected. At Jupiter’s estimated $\sigma_\eta \approx 0.2$, a hypothetical $N = 7$ ring would have $P \approx 0$ —the marginal sensitivity manifests only at unphysically large disorder.

Thus the N_{crit} boundary is a *stability domain boundary*, not a *knife-edge*: configurations comfortably inside the stable region remain stable under the circulation fluctuations observed at Jupiter ($\sigma_\eta \approx 0.2$) and Saturn ($\sigma_\eta \lesssim 0.05$).

Code: `circulation_disorder.instability_probability(N,  $\rho$ ,  $\sigma_\eta$ )` and `disorder_min_eigenvalue` in `src/planetary_polygons/extensions/circulation_disorder.py`; test suite in `tests/test_circulation_disorder.py`.

8.0.2 The regime criterion

The dimensionless number $\beta R^3/\kappa$ is the ratio of two timescales: the Rossby wave period $T_R \sim 1/(\beta R)$ and the vortex turnover time $T_v \sim R^2/\kappa$. When $\beta R^3/\kappa = T_v/T_R \ll 1$, the vortex interactions complete many orbits before the Rossby wave propagates appreciably, so the dynamics is vortex-dominated. When $\beta R^3/\kappa \gg 1$, the Rossby wave dynamics operates faster than the vortex rearrangement, so the wave mechanism dominates:

Regime	$\beta R^3/\kappa$	Bridge
Vortex-dominated	$\ll 1$	Blob convergence (Proposition 6.1)
Wave-dominated	$\gg 1$	Rossby stationarity (equation (6.7))

Jupiter’s north pole. Throughout this regime calculation κ denotes the *Rankine-core* circulation, which integrates only within the solid-body rotation core and sets the characteristic vortex timescale $T_v \sim R^2/\kappa$. The full-cyclone circulation κ_{full} (which includes the outer irrotational skirt) is larger and is the appropriate quantity for computing orbital periods and eigenfrequencies; it is used there instead (Section Paper III, §21.3). The two estimates are numerically distinguished below.

At 84°N , the planetary vorticity gradient is $\beta_J = 2\Omega_J \cos(84^\circ)/R_J \approx 5.1 \times 10^{-13} \text{ m}^{-1}\text{s}^{-1}$ and the cyclone-ring radius is $R \approx 7,470 \text{ km}$. Each Juno-observed circumpolar cyclone has Rankine core radius $r_c \approx 700 \text{ km}$ and peak wind $V_{\text{max}} \approx 80 \text{ m s}^{-1}$ [Adriani et al., 2018, Gavriel and Kaspi, 2021], giving $\kappa_{\text{Rank}} = 2\pi r_c V_{\text{max}} \approx 3.5 \times 10^8 \text{ m}^2 \text{ s}^{-1}$ (full-cyclone estimate: $r_c \approx 3,200 \text{ km}$, $\kappa_{\text{full}} \approx 1.6 \times 10^9 \text{ m}^2 \text{ s}^{-1}$). Hence

$$\frac{\beta R^3}{\kappa} = \frac{5.1 \times 10^{-13} (7.47 \times 10^6)^3}{3.5 \times 10^8} \approx 0.61.$$

This places Jupiter in the vortex-dominated regime, but only moderately: 0.61 is $O(1)$, not $\ll 1$. Rossby-wave effects are not negligible at this value, and a decrease in the effective κ (e.g. using the Rankine-core rather than full-cyclone estimate) would push the ratio above unity. The qualitative conclusion—discrete cyclones, Thomson stability applies, blob bridge proves convergence—is robust to this uncertainty because the vortex-dominated bridge requires only that the cyclones be identifiable as N distinct vortices (which the Juno imagery unambiguously confirms), not that $\beta R^3/\kappa$ be parametrically small. The ratio enters quantitatively only through the Rossby-wave correction to the eigenvalues, which is $O(\beta R^3/\kappa)$ and does not change the sign of $\lambda_{\min}$ for $N = 8$ (where

$|\lambda_{\min}| = 1 \gg 0.61$). Using the full-cyclone circulation $\kappa_{\text{full}} \approx 1.6 \times 10^9 \text{ m}^2 \text{ s}^{-1}$ gives $\beta R^3/\kappa \approx 0.13$, firmly in the vortex-dominated regime.

Saturn's hexagon. At 76°N , $\beta_S = 2\Omega_S \cos(76^\circ)/R_S \approx 1.3 \times 10^{-12} \text{ m}^{-1} \text{ s}^{-1}$ and $R \approx 14,600 \text{ km}$. Saturn's polar jet is well-described by a Gaussian profile $U = U_{\text{max}} \exp(-y^2/2\sigma^2)$; its effective circulation is $\kappa = U_{\text{max}} \sqrt{2\pi} \sigma$. With $U_{\text{max}} \approx 100 \text{ m s}^{-1}$ and $\sigma \approx 2,100 \text{ km}$ (Cassini VIMS wind profiles), $\kappa \approx 5.4 \times 10^8 \text{ m}^2 \text{ s}^{-1}$, giving

$$\frac{\beta R^3}{\kappa} = \frac{1.3 \times 10^{-12} (1.46 \times 10^7)^3}{5.4 \times 10^8} \approx 7.6.$$

Continuous jet meander, Rossby stationarity selects $n = 6$, periodicity provides quantization.

Both mechanisms require the 2D logarithmic Green's function as the common mathematical kernel.

9 Predictions for quantum systems: BEC vortex clusters

In a quasi-2D Bose–Einstein condensate in a circular trap, quantized vortices interact logarithmically at scales above the healing length ξ . The circulation is quantized: $\kappa = q \cdot h/m$ with integer winding number q . The κ_0/κ ratio of Proposition 3.1 becomes the winding number ratio q_0/q , yielding sharp integer-valued predictions:

Configuration	κ_{crit}	Prediction
$N \leq 7$ ring	0	Stable without central vortex
$N = 8$ ring	$1/2$	Unstable alone; stable with $q_0 \geq 1$
$N = 9$ ring	1	Marginally stable with $q_0 = 1$; stable with $q_0 \geq 2$
$N = 10$ ring	$7/4$	Unstable with $q_0 = 1$; stable with $q_0 \geq 2$

The $N = 8$ prediction is the sharpest: a ring of 8 unit-winding vortices should be unstable alone and stable upon adding a single central vortex ($q_0 = 1 > \kappa_{\text{crit}} = 1/2$).

The $N = 7$ case offers a distinct signature. The quartic restoring force ($\alpha_0 = 45/14$, Section 4) produces an *amplitude-dependent* oscillation frequency for the $m = 3$ mode. In the normal form $\mathcal{H} = (\alpha_0/4)\varepsilon^4$, the nonlinear oscillation frequency is $\omega_{m=3} = \alpha_0 \varepsilon^2 \kappa / (2\pi R^2)$, where ε is the dimensionless perturbation amplitude (radial displacement / ring radius) and $\kappa = h/m_{\text{atom}}$ is the single-quantum circulation. For ^{87}Rb in a circular trap of radius $R = 50 \text{ } \mu\text{m}$: $\kappa \approx 1.2 \times 10^{-7} \text{ m}^2 \text{ s}^{-1}$, $\kappa / (2\pi R^2) \approx 7.6 \text{ s}^{-1}$, so $\omega_{m=3} \approx 24 \varepsilon^2 \text{ s}^{-1}$. At a 5% radial displacement ($\varepsilon = 0.05$): $\omega_{m=3} \approx 0.06 \text{ s}^{-1}$ (period $\sim 100 \text{ s}$), resolvable with standard BEC vortex imaging at $\sim 10 \text{ Hz}$ frame rates. The key signature distinguishing the $N = 7$ quartic confinement from the $N \leq 6$ quadratic regime is that the $m = 3$ oscillation frequency *vanishes* as $\varepsilon \rightarrow 0$ (rather than approaching a finite value), producing an amplitude-dependent period that is directly measurable.

The healing length ξ provides a natural blob width $\varepsilon_{\text{blob}} = \xi/R$. For typical BEC parameters ($\xi \sim 0.3 \text{ } \mu\text{m}$, $R \sim 50 \text{ } \mu\text{m}$): $\varepsilon_{\text{blob}} \approx 0.006$, well within the point vortex regime. The blob correction (Proposition 6.1) is negligible: $\lambda(0.006) \approx \lambda_0 + O(10^{-4})$.

Finite-temperature effects. Current BEC vortex experiments operate at finite temperature ($T/T_c \sim 0.3\text{--}0.7$), where a thermal cloud coexists with the condensate and exerts a

dissipative drag on vortex motion [Fetter, 2009]. This drag damps the oscillation modes on a timescale $\tau_{\text{drag}} \sim R^2/(D_v T/T_c)$ where D_v is the vortex diffusion coefficient. For the $N = 7$ quartic mode to be observable, the oscillation period (~ 100 s at $\varepsilon = 0.05$) must be shorter than τ_{drag} ; at $T/T_c = 0.5$ this requires $D_v \lesssim 10^{-9} \text{ m}^2 \text{ s}^{-1}$, which is achievable in large-radius traps ($R \geq 50 \text{ } \mu\text{m}$) at low temperature. The $N = 8$ stability prediction (unstable alone, stable with $q_0 = 1$) is less sensitive to dissipation because it concerns exponential growth vs. decay, not oscillation frequency; finite temperature adds a damping channel but does not change the sign of the eigenvalue.

The hard-wall disk: an exact theorem. In a hard-wall circular trap of radius R , each vortex at z has an image at $R^2/\bar{z}$. The effective interaction is *exactly* the Green's function on the Poincaré disk of radius R (the conformal equivalence is the identity map).

Proposition 9.1 (Hard-wall BEC = Poincaré disk). *In a hard-wall circular BEC of radius R , the stability eigenvalue of an N -gon vortex ring at radius r is $\lambda_m = C_1(\xi) - m(N-m)/2$ with $\xi = r^2/R^2$ and $C_1(\xi) = (N-1)(1+\xi^2)/(1-\xi)^2$ —the same formula as for $\mathbf{H}^2$ (Paper I, equation 5.7).*

Proof. The image vortex construction gives the Green's function $G(z, w) = -(1/2\pi) \log |z-w| + (1/2\pi) \log |1-z\bar{w}/R^2|$, which is exactly the Green's function on the Poincaré disk of radius R . The Havelock decomposition applies without approximation. $\square$

The palindromic thresholds $\xi^*(N)$ are therefore *exact* predictions for hard-wall BEC experiments:

N	$\xi^* = r^2/R^2$	r/R	Prediction
≤ 7	0	0	Stable at any radius
8	$8-3\sqrt{7} \approx 0.063$	0.25	Stable for $r > 0.25 R$
11	$(1+\sqrt{2})^{-2} \approx 0.17$	0.41	Stable for $r > 0.41 R$
23	$\varphi^{-2} \approx 0.38$	0.62	Stable for $r > 0.62 R$

The **key testable prediction**: an 8-vortex ring at $r/R > 0.25$ in a hard-wall disk BEC should be stable; at $r/R < 0.25$ it should be unstable to the $m = 4$ mode. The threshold $r/R = \sqrt{8-3\sqrt{7}} \approx 0.25$ is an algebraic number in $\mathbb{Q}(\sqrt{7})$ —the same field that governs the $7 \rightarrow 8$ transition in planetary vortex dynamics.

Flat-bottomed BEC traps, in which the point vortex model has been experimentally validated [Gauthier et al., 2019], are the natural setting for testing these predictions.

The harmonic trap: an approximation. For a harmonic trap $V(r) = \frac{1}{2}m\omega^2 r^2$ with Thomas–Fermi radius R_{TF} : the identification $\xi = r^2/R_{\text{TF}}^2$ holds to leading order, but the soft boundary introduces corrections at $O(\xi^2)$ relative to the hard-wall case (the image vortex positions shift by terms proportional to the trap anharmonicity). The prediction table above applies for $\xi \ll 1$ ($r \ll R_{\text{TF}}$), where the hard-wall and harmonic cases coincide. For vortex rings near the condensate edge ($\xi \rightarrow 1$), the harmonic trap softens the effective C_1 , and the point vortex model itself requires corrections from the inhomogeneous density profile.

Healing-length and quantum corrections. The BEC healing length ξ_{heal} sets the blob scale $\varepsilon = \sqrt{2} \xi_{\text{heal}}/d$ where $d = 2r \sin(\pi/N)$ is the intervortex spacing. For typical parameters ($\xi_{\text{heal}} \approx 0.3 \text{ } \mu\text{m}$, $R \approx 50 \text{ } \mu\text{m}$): $\varepsilon \approx 0.02$, giving $\delta\lambda \sim \varepsilon^2 \sim 0.04\%$ (Proposition 6.1).

The *quantum* correction (from the Bogoliubov–de Gennes excitation spectrum) is even smaller. The BdG anomalous coupling between vortex cores separated by d is

$O(\exp(-d/\xi_{\text{heal}}))$; for $N = 8$ at the threshold ($d \approx 9.6 \mu\text{m}$): $\exp(-d/\xi_{\text{heal}}) \approx 10^{-14}$. The effective mass renormalization shifts the precession frequency but not the energy Hessian (the stability condition $\lambda_m > 0$ is a property of the potential landscape, not the kinetic term). The palindromic threshold $\xi^*(8) = 8 - 3\sqrt{7}$ is therefore an *exact* feature of the quantum spectrum, preserved to exponential accuracy by the topological invariance of the K_0 class (Paper I, Theorem 6.12).

Hyperbolic lattice realisations. Interacting bosons on a $\{p, q\}$ hyperbolic lattice (Bose–Hubbard model, hopping t , on-site U) reduce to the point vortex Hamiltonian in the Thomas–Fermi regime ($U \gg t$): the mean-field condensate $\psi = \sqrt{\rho} e^{i\theta}$ has vortices where $\psi = 0$, and θ satisfies $\Delta_{\text{lattice}}\theta = 2\pi \sum_k \delta(z - z_k)$, giving the lattice Green’s function interaction $H = -\sum_{j < k} G_{\text{lattice}}(z_j, z_k)$. In the continuum limit (lattice spacing $a \ll$ intervortex spacing d): $G_{\text{lattice}} \rightarrow G_{\mathbf{H}^2}$ and the Havelock thresholds apply. The continuum limit requires $\gtrsim 1000$ lattice sites—beyond current circuit-QED (~ 100 sites) but within reach of topoelectric circuits ($\sim 10^4$ sites). On the Bolza tiling ($\{8, 8\}$, genus 2, $|\text{Aut}| = 48$): the prediction is $N_{\text{crit}} = 12$ at the Aut-fixed centre, with $\delta C_1 \approx 0.011$ (Paper I, §7).

Remark 9.1 (Mesoscopic superconducting disks). *Abrikosov vortices in a type-II superconducting disk of radius R obey the same boundary condition as BEC vortices in a hard-wall trap (supercurrent vanishes at the edge $\Leftrightarrow$ Kelvin images at $R^2/\bar{z}$). Proposition 9.1 therefore applies directly: $\lambda_m = C_1(\xi) - m(N-m)/2$ with $\xi = r_0^2/R^2$. The observed experimental sequence in mesoscopic Nb disks (Grigorieva et al. 2006): single polygon for $N \leq 6$, marginal polygon or shell (1+6) at $N = 7$, and shells (1+7, 2+7, ...) for $N \geq 8$, is consistent with $N_{\text{crit}} = 7$ at the disk centre ($\xi = 0$) and central-vortex stabilisation (Section ??) raising N_{crit} to 8 for the outer ring.*

Remark 9.2 (Laughlin quasihole polygons). *The effective Hamiltonian of the Laughlin state at filling $\nu = 1/m$ is $H_{\text{eff}} = -m \sum_{j < k} \log |z_j - z_k| + (4\ell_B^2)^{-1} \sum |z_k|^2$: the logarithmic energy with harmonic confining. Quasiholes (created by factors $\prod (z_j - \eta_a)$) interact via $H_{\text{qh}} = -(1/m) \sum_{a < b} \log |\eta_a - \eta_b| + \text{confining}$, where the $1/m$ comes from the fractional charge $e^* = e/m$. For P quasiholes in a ring at radius r inside a Laughlin droplet of N electrons (radius $R = \ell_B \sqrt{2N}$): the Havelock eigenvalue is $\lambda_p = (1/m)[C_1(\xi) - p(P-p)/2]$ with $\xi = r^2/R^2$. The factor $1/m$ is overall: the stability threshold P_{crit} is independent of the filling fraction m . The palindromic thresholds of Paper I, §5 apply directly: $P \leq 7$ is stable at any r ; $P = 8$ requires $r/R > 0.25$; etc. The Berry phase (statistical angle $\theta = \pi/m$, giving anyonic braiding) adds an imaginary term to H_{eff} that affects the dynamics but not the energy Hessian—the stability eigenvalues are real and independent of the statistical angle. The conditions: $N \gtrsim 55$ electrons (for $P = 8$ at $r/R = 0.25$), and $d = 2r \sin(\pi/P) \gg \ell_B$ (well-separated quasiholes).*

Remark 9.3 (Non-abelian quasiholes and topological quantum computation). *The Moore–Read state ($\nu = 5/2$) has quasiholes carrying Majorana zero modes with non-abelian braiding. The effective interaction between two quasiholes at separation d is $E(d) = E_{\text{log}}(d) + \delta E_{\text{fusion}}(d)$ where $\delta E_{\text{fusion}} \sim C \exp(-d/\xi_E)$ with $\xi_E \approx 2.3 \ell_B$ (Baraban et al. 2009). The exponential decay gives effective exponent $\alpha = \infty \gg 2$ in the RG classification (Proposition ??): the fusion correction is deeply irrelevant. By the blob bridge (Proposition 6.1): $P_{\text{crit}} = 7$ is unchanged.*

The same argument applies to all gapped quantum Hall states (Read–Rezayi, hierarchy states): the fusion correction is always exponential in d (from the bulk energy gap), hence always RG-irrelevant. $P_{\text{crit}} = 7$ is therefore universal across all gapped quantum Hall states, independent of filling fraction, braiding statistics, and fusion channel.

For topological quantum computation: standard braiding protocols use $P \leq 6$ quasi-holes (4 for Moore–Read, 6 for two-qubit Fibonacci), safely below P_{crit} . The threshold is not a practical constraint on current proposals but is a fundamental limit on the complexity of quasihole polygon configurations in any quantum Hall platform.

Remark 9.4 (Independence of spectral and TQFT structures). *On a torus $\Sigma = \mathbb{C}/(\mathbb{Z} + \tau\mathbb{Z})$, the topological correction $\delta_m(\Sigma)$ to the Havelock eigenvalue has two natural candidate descriptions:*

1. Geometric spectral theory. *The Laplacian spectrum of Σ (eigenvalues $n_1^2 + n_2^2$ on the lattice) gives a spectral expansion of δ_m . A 20-mode fit achieves $R^2 = 0.79$: the correction is 79% spectral (heat kernel / Selberg type), with the remaining 21% from the full $\log |\theta_1|$ structure.*
2. TQFT fusion labels. *For $N = 23$ at level $k = 5$, the quantum Casimir $f_q(m, N) \in \{0, 1, \varphi\}$ decomposes the modes into three $SU(2)_3$ sectors (null, vacuum, Fibonacci). The Verlinde S -matrix of $SU(2)_3$ has off-diagonal entries ~ 0.6 , which would mix sectors under $\tau \rightarrow -1/\tau$.*

These two structures are independent. The modular S -transform $\tau \rightarrow -1/\tau$ acts as a scalar on all three fusion sectors (ratio $\delta_m(i/2)/\delta_m(2i) \approx 0.97$ for each sector, with spread $< 0.4\%$), not via the Verlinde S -matrix. The fusion labels organise the mode space algebraically, but the spectral corrections do not respect this organisation: within-sector variance equals 95% of the total variance.

The TQFT and the geometric spectral theory share the torus as domain and share the palindromic/modular symmetries, but their spectral data act on different spaces. The fusion identities (Theorem 3 of the q -deformation analysis) describe the algebraic quantisation of the Casimir spectrum; the torus corrections describe the geometric modification of the eigenvalues. The two are complementary, not interchangeable.

10 Gravitational polygon stability in 2+1 dimensions

The 2+1D gravity literature studies the single conical deficit [Deser et al., 1984], two-body scattering, BTZ black holes, polygon tessellations of spatial slices ('t Hooft 1993; note that “polygon” there refers to the tiles of the tessellation, not to a configuration of point masses), and one-loop partition functions [Giombi et al., 2008]. What appears not to have been asked is: *for how many point masses can one form a stable regular polygon on a spatial slice, and what algebraic number theory governs the answer?* This section derives the answer from the vortex stability framework of Paper I.

In 2+1 dimensional gravity with negative cosmological constant ($\Lambda = -1/\ell^2$), space-time is locally AdS_3 and has no propagating degrees of freedom. A point mass m creates a conical deficit of angle $\alpha = 8\pi Gm$ [Deser et al., 1984]. A constant-time spatial slice is a quotient $S = \Gamma \backslash \mathbf{H}^2$, and the linearised gravitational potential of a point mass on S satisfies $\nabla^2 \Phi = 8\pi G m \delta(z - z_0)$ —the same Laplace equation as for the stream function of a point vortex.

Proposition 10.1 (Gravitational polygon stability). *In the linearised regime ($8\pi Gm_i \ll 2\pi$, small deficit angles), N equal point masses of mass m on a spatial slice $S = \Gamma \backslash \mathbf{H}^2$ interact via*

$$E_{\text{grav}} = -(8\pi G)^2 \sum_{i < j} m^2 G_S(z_i, z_j), \quad (10.1)$$

where G_S is the Green's function of the Laplacian on S (the same G_S appearing in the vortex energy with $\kappa \leftrightarrow 8\pi Gm$). The regular N -gon is a critical point of E_{grav} on the angular-impulse surface, and the Hessian eigenvalues are

$$\lambda_m^{\text{grav}} = C_1(S, \xi) - \frac{m(N-m)}{2}, \quad m = 1, \dots, N-1, \quad (10.2)$$

with the same C_1 and $m(N-m)/2$ as for vortices. The N -gon is a stable equilibrium ($\lambda_m^{\text{grav}} > 0$ for all m) if and only if $N \leq N_{\text{crit}}(S, \xi)$.

Proof. The linearised gravitational potential satisfies the same Poisson equation as the vortex stream function, with $\kappa = 8\pi Gm$. The interaction energy (10.1) is therefore $E_{\text{grav}} = (8\pi Gm/\kappa)^2 \cdot H_{\text{vortex}}$, a positive multiple of the vortex Hamiltonian. At the N -gon, the first and second variations of E_{grav} are proportional to those of H_{vortex} , so the critical points, the Hessian eigenvalues, and the stability criterion ($\lambda_m > 0$ for all m) are identical. The Havelock decomposition (Paper I, §4) gives (10.2).

Note on dynamics. The vortex equations of motion are *first order* (the Hamiltonian is the stream function), while the gravitational equations are *second order* (kinetic + potential energy). The stability criterion—positive-definite Hessian of the potential—is the same for both: a positive eigenvalue $\lambda_m > 0$ gives stable oscillation in the gravitational case and stable circulation in the vortex case. A negative eigenvalue gives exponential growth in both. The criterion $N \leq N_{\text{crit}}$ is therefore identical.

Validity and the developing map. For N equal masses in a ring, each creates a deficit angle $\delta = 8\pi Gm$. The total deficit is $N\delta < 4\pi$ (Gauss–Bonnet for S^2), giving $Gm < 1/(2N)$. We now show that, remarkably, the backreaction of the deficit angles on the flat-plane stability is *exactly zero*.

The **developing map** of the flat cone metric with $\mathbb{Z}_N$ symmetry is

$$f(z) = z \cdot {}_2F_1\left(\beta, \frac{1}{N}; 1 + \frac{1}{N}; z^N\right), \quad \beta = 4Gm, \quad (10.3)$$

with derivative $f'(z) = (1 - z^N)^{-\beta}$. This is a conformal isometry from the cone surface to the Euclidean plane: the metric is $|f'|^2 |dz|^2$, flat away from the cone tips, with cone angle $2\pi(1-\beta)$ at each N -th root of unity.

$\mathbb{Z}_N$ symmetry. $f(\omega z) = \omega f(z)$ for $\omega = e^{2\pi i/N}$: the image of the N -gon of cone points is a regular N -gon in the developed plane, with radius

$$R_f = \frac{\Gamma(1 + 1/N) \Gamma(1 - \beta)}{\Gamma(1 + 1/N - \beta)} \quad (10.4)$$

(Gauss's summation formula for ${}_2F_1$ at $z = 1$). At $\beta = 0$: $R_f = 1$. At the Gauss–Bonnet endpoint $\beta = 1/N$: $R_f = (\pi/N)/\sin(\pi/N)$ (by the Γ -reflection formula). The image radius *increases* with β : the cone metric inflates distances near the tips.

Proposition 10.2 (Conformal invariance of flat-plane stability). *On a flat spatial slice ($\Lambda = 0$) with N equal conical singularities of deficit $\delta = 8\pi Gm$ in a regular ring, the Havelock eigenvalues are*

$$\lambda_m = (N-1) - \frac{1}{2}m(N-m)$$

for all $Gm \in [0, 1/(2N))$. In particular, $N_{\text{crit}} = 7$ is independent of the deficit angle.

Proof. The two-dimensional Green's function is conformally invariant: for any conformal metric $g = e^{2u}|dz|^2$ on the plane, $\Delta_g G = \delta$ gives $G(z, w) = -(2\pi)^{-1} \log |z-w|$

in conformal coordinates, independent of u . Therefore the interaction energy $H = -\sum_{j < k} \log |z_j - z_k|$ is the same in conformal coordinates as on the flat plane.

The $\mathbb{Z}_N$ symmetry maps the N -gon to a regular N -gon of radius R_f in the developed plane. The angular interaction $-\sum_{j < k} \log |e^{i\theta_j} - e^{i\theta_k}|$ is independent of the radius (since $\log R$ drops from the Hessian), and so are the Fourier eigenvalues $m(N-m)/2$ of the tangential Hessian. The constraint curvature contributes $+(N-1)$ independently of R_f , completing the Havelock decomposition. $\square$

The topology change at $\beta = 1/N$. At the Gauss–Bonnet endpoint $Gm = 1/(2N)$, the total deficit exhausts the full curvature budget and the spatial slice *compactifies* to S^2 . On the sphere, the Green’s function acquires a curvature correction: $G^{S^2} = -(2\pi)^{-1} \log |z-w| + (4\pi)^{-1} \log[(1+|z|^2)(1+|w|^2)]$ (in stereographic coordinates). The additional term replaces $C_1 = N-1$ with $C_1 = (N-1) \cos \varphi < N-1$, giving $N_{\text{crit}}^{S^2} \leq 6$. The transition from $N_{\text{crit}} = 7$ (flat, all $Gm < 1/(2N)$) to $N_{\text{crit}} \leq 6$ (sphere) is therefore *discontinuous*: it occurs at the topology change, not through a gradual backreaction. The curvature-stability dichotomy of Proposition 10.6 arises from the cosmological constant Λ (which sets $\xi = R^2/\ell^2$), not from the deficit angles.

A remark on “no gravitational force in 2+1D.” Pure 2+1D gravity has no propagating degrees of freedom and no Newtonian $1/r$ force at a distance. But point masses on an AdS_3 spatial slice $\Gamma \setminus \mathbf{H}^2$ *do* have a well-defined static potential: the Green’s function of the Laplacian, which is logarithmic. This potential does not propagate dynamically (there are no gravitational waves), but it determines the equilibrium energy landscape. The stability analysis depends only on the second variation of this static energy, for which the Havelock decomposition applies exactly as in the vortex case. $\square$

Proposition 10.1 transfers the full palindromic hierarchy to gravity:

- $N_{\text{crit}} = 7$ on any spatial slice (universality, Paper I, Theorem 6.8).
- The palindromic thresholds of Paper I, §6 are *gravitational phase transitions*: algebraic-number curvature values at which the stability of the mass polygon changes.
- The five algebraic-integer thresholds ($\xi^*(8) = 8-3\sqrt{7}$, $\dots$, $\xi^*(23) = \varphi^{-2}$) give parameter-free predictions for the stability of mass rings in 2+1D AdS.

10.1 K-theoretic classification of gravitational phases

The K -theoretic framework of Paper I, §7 classifies gravitational stability phases. On a spatial slice $S = \Gamma \setminus \mathbf{H}^2$:

1. The negative spectral projection defines $[\mu_N^{\text{grav}}(\xi)] \in R(\mathbb{Z}_N)$ (Paper I, §7.1).
2. The spectral correction δC_1 is the image of the Havelock test function under the Baum–Connes assembly map for Γ (Paper I, Theorem 7.1). In 2+1D quantum gravity, this assembly map already computes the K -theory of the spatial slice [Carlip, 1998]: the vortex stability computation *is* the gravitational K -theory computation.
3. The Kasparov product $\alpha_N \otimes_{C_*^*(\Gamma)} [D_S] = [\mu_N]$ (Paper I, Theorem 7.3) gives a K -theoretic classification of gravitational stability phases, with phase transitions at the palindromic thresholds.

10.2 The spectral bound as a quantum gravitational bound

The spectral correction and the gravitational one-loop partition function are both moments of the heat kernel:

$$\log Z_{\text{grav}} = - \int_0^\infty [\text{Tr } e^{-t\Delta_S} - 1] \frac{dt}{t} \quad (\text{weight } 1/t), \quad (10.5)$$

$$\delta C_1(z_0) = \int_0^\infty [K(z_0, z_0, t) - K_{\mathbf{H}^2}(t)] dt \quad (\text{weight } 1). \quad (10.6)$$

The first is the functional determinant $\det \Delta_S$, the second determines polygon stability. Both involve the Selberg zeta function $Z_\Gamma(s)$ at the zero $s = 1$, but this requires care.

The zero at $s = 1$. The Selberg zeta $Z_\Gamma(s)$ has zeros at s such that $s(1-s)$ is an eigenvalue of Δ_S . The constant eigenfunction ($\lambda_0 = 0$) gives a zero at $s = 1$, so $Z_\Gamma(1) = 0$ —not a finite number. The Euler product $\prod_\gamma \prod_k (1 - e^{-(s+k)\ell_\gamma})$ converges for $\text{Re}(s) > 1$ but diverges at $s = 1$ (by the prime geodesic theorem: the partial product over geodesics with $\ell \leq L$ decays to 0 as $L \rightarrow \infty$).

The physically relevant quantities are the *residue* at this zero:

- $Z'_\Gamma(1)$: the *regularised* determinant $\det' \Delta_S = \prod_{n \geq 1} \lambda_n$ (omitting $\lambda_0 = 0$). This is the gravitational one-loop partition function [Giombi et al., 2008].
- The spectral sum $\sum_{n \geq 1} |\varphi_n(z_0)|^2 / \lambda_n$: the regularised Green's function at the diagonal, which determines $\delta C_1(z_0)$.

Both are finite and well-defined; both involve the analytic structure of Z_Γ at its zero, not the (divergent) Euler product.

The spectral computation of δC_1 . At the Aut-fixed centre z_0 of the Bolza surface, only trivial-representation eigenfunctions contribute: $\delta C_1(z_0) = \sum_{n: \text{triv}} |\varphi_n(z_0)|^2 / \lambda_n \approx 0.011$, dominated by the first trivial-rep eigenvalue $\lambda_{\text{triv}} \approx 15.05$ (Strohmaier and Uski 2013), with rapid convergence: $1/(4\pi \cdot 15.05) \approx 0.005$, $1/(4\pi \cdot 23.08) \approx 0.003$, $1/(4\pi \cdot 32.67) \approx 0.002$, $\dots$. The spectral bound $|\delta C_1| \leq 2/(4\pi \lambda_{\text{triv}}) \approx 0.011$ is rigorous—it is the first term of a convergent sum.

The spectral gap λ_1 of the Laplacian controls both quantities. In the gravitational language:

- $C_1(\mathbf{H}^2, \xi)$: classical stability (flat-space AdS, no topology).
- δC_1 : one-loop quantum correction (from the topology of the spatial slice).
- The bound $|\delta C_1| \leq 2/(4\pi \lambda_1)$: a bound on the quantum gravitational correction to polygon stability.

The Selberg dictionary. Four distinct physical quantities involve the Selberg zeta at the zero $s = 1$:

Z_Γ data at $s = 1$	Physical meaning
$Z'_\Gamma(1)$ (regularised determinant $\det' \Delta_S$)	Gravitational partition function $= \prod_{n \geq 1} \lambda_n$
$\sum_{n \geq 1} \varphi_n(z_0) ^2 / \lambda_n$ (spectral; convergent)	Vortex stability correction δC_1 Localised at z_0 ; Aut-suppressed
Palindromic threshold polynomials (from $C_1(\xi) = f_{\max}(N)$)	Stability transitions: values of ξ where N_{crit} jumps
Hecke eigenvalues of associated forms (a_p from modular forms)	Prime-by-prime visibility Which thresholds are visible on $\Gamma(p)$

The first two rows involve the same zero $s = 1$ of Z_Γ but differ in what they extract: the determinant (a *product* of eigenvalues) versus the resolvent trace (a *sum* of $1/\lambda_n$, localised at z_0). Neither can be computed from the Euler product (which diverges at $s = 1$); both require either the spectral data or the analytic continuation of Z_Γ .

The palindromic thresholds control *where* stability transitions occur (at which curvature ξ); the eigenvalue spectrum controls *how large* the correction δC_1 is. On highly symmetric surfaces (the modular surface, the Bolza surface), the systole and the palindromic geodesics coincide; on a generic arithmetic surface they are distinct.

10.3 Position-dependent gravitational stability

At a palindromic threshold, the L^2 -Morse index is fractional (Paper I, §7.3): on the Bolza surface, $\tau(P_-) \approx 0.61$. In the gravitational setting, this means: *the stability of a mass polygon depends on where it sits on the spatial slice*. The automorphism-fixed centre is stable ($\delta C_1(z_0) > 0$), while 61% of the spatial slice (by L^2 -weight) is unstable. This position dependence is a gravitational analogue of the placement dependence on Saturn (Paper III, Section 3.2): the polar flattening selects the pole as the stable site for the hexagon, just as the automorphism group selects the Aut-fixed centre as the stable site for the mass polygon.

The gravitational phase diagram— $N=5$ modes stabilising one by one as the spatial curvature increases, with phase transitions at palindromic algebraic thresholds—is the 2+1D AdS analogue of the topological phase diagrams classified by Kitaev (2009) and Freed–Hopkins (2021) in condensed matter.

10.4 Gravity, Langlands, and the palindromic hierarchy

The palindromic polynomials on arithmetic spatial slices with trace field degree $d \geq 5$ have non-solvable Galois groups (Paper I, Proposition A.13). Since the gravitational interaction (10.1) produces the *same* palindromic polynomials as the vortex interaction, the gravitational stability thresholds on these slices are open cases of the Langlands reciprocity conjecture (Paper I, Remark A.8).

This gives a connection between quantum gravity and the Langlands programme that does not pass through string theory. The chain is:

1. 2+1D AdS gravity with point masses on a spatial slice $\Gamma \backslash \mathbf{H}^2$ [Deser et al., 1984].
2. The logarithmic Green’s function $\rightarrow$ the Havelock decomposition $\rightarrow$ palindromic stability thresholds.

3. On arithmetic slices with totally real quintic trace field: S_5 Galois group, LMFDB Artin representations 4.38569.5t5.a.a, 4.65657.5t5.a.a, 4.24217.5t5.a.a (Paper I, §A.8).
4. Langlands reciprocity predicts automorphic forms on $\mathrm{GL}(4)/\mathbb{Q}$ matching these representations—a prediction that is currently *unproved* and does not follow from any known modularity lifting result.

The Selberg zeta function $Z_\Gamma(s)$ encodes both the gravitational partition function ($Z'_\Gamma(1)$, the regularised determinant) and the stability correction (δC_1 , the spectral resolvent at the zero $s = 1$). If the Langlands lift π exists, it constrains Z_Γ at $s = 1$ —simultaneously constraining the gravitational partition function and the stability phase diagram.

10.5 Holographic interpretation

In $\mathrm{AdS}_3/\mathrm{CFT}_2$, a point mass m in the bulk creates a conical deficit that corresponds to the insertion of a heavy operator $\mathcal{O}$ of conformal dimension $h = (c/24)(1 - (1 - 8Gm)^2)$ in the boundary CFT [Carlip, 1998]. A regular N -gon of equal masses at radial coordinate r in the bulk corresponds to N equally-spaced heavy operator insertions on the boundary circle.

In this language, the Havelock decomposition takes the form:

Bulk (2+1D AdS gravity)	Boundary (1+1D CFT)
Perturbation mode m	$\mathbb{Z}_N$ conformal block channel
$f(m, N) = m(N-m)/2$	Casimir of the m -th channel
$C_1(\xi)$	Function of h/c and AdS radius
$\lambda_m = C_1 - f(m, N) > 0$	Positivity of the m -th block
$C_1 \geq f_{\max}(N)$	Bootstrap-type bound on h/c

The stability condition $C_1 \geq m(N-m)/2$ for all m becomes a *lower bound on the conformal dimension* of the boundary operators: below this bound, the bulk polygon is unstable and the corresponding $\mathbb{Z}_N$ conformal block changes sign. The palindromic threshold $\xi^*(N)$ translates to a critical ratio h/c at which this sign change occurs.

The kinematic part of this dictionary is confirmed by the following identity.

Proposition 10.3 (Havelock Casimir = one-chiral OPE Casimir). *Let $T_m = \sum_{p=1}^{N-1} \frac{1 - \cos(2\pi pm/N)}{2 \sin^2(\pi p/N)}$ be the $\mathbb{Z}_N$ Fourier eigenvalue of the OPE kernel $\mathrm{csc}^2(\pi p/N)$. Then*

$$T_m = m(N-m) = 2f(m, N), \quad (10.7)$$

exactly twice the Havelock Casimir $f(m, N) = m(N-m)/2$.

Proof. The identity $T_m = m(N-m)$ is a standard discrete Fourier transform (verified computationally for $N \leq 30$; the proof is a residue computation on the Chebyshev polynomials). $\square$

The factor of 2 has physical content. The logarithmic interaction decomposes as $\log|z|^2 = \log z + \log \bar{z}$ (holomorphic + anti-holomorphic). In the boundary CFT, the OPE of two heavy operators has both chiral sectors, giving the *full* Casimir $m(N-m)$. The Havelock eigenvalue uses $f = m(N-m)/2$, the contribution of *one* chiral sector, because:

- Vortex dynamics is first-order (the Hamiltonian is the stream function, not the kinetic + potential energy), so only one chiral sector enters the stability eigenvalue.
- The stability criterion $\lambda_m > 0$ is $C_1 > m(N-m)/2$: the confining potential must exceed the holomorphic OPE repulsion in the m -th $\mathbb{Z}_N$ channel.

In gravity (second-order dynamics), both sectors contribute to the equations of motion, but the stability criterion—positivity of the Hessian eigenvalue—still involves $m(N-m)/2$ per sector (Proposition 10.1).

The remaining step is to express C_1 in terms of boundary CFT data. On the flat plane, $C_1 = N-1$ (the Havelock self-energy); on $\mathbf{H}^2$ at curvature ξ , $C_1 = (N-1)(1+\xi^2)/(1-\xi)^2$. The curvature parameter ξ encodes the radial position of the bulk polygon; through the Brown–Henneaux relation $c = 3\ell/(2G)$ and the deficit angle–dimension map $h = (c/24)(1-(1-8Gm)^2)$, the stability condition becomes a bound on conformal dimensions parametrised by the bulk position. The palindromic thresholds $\xi^*(N)$ give algebraic-number bulk positions at which the stability bound is saturated:

N	$\xi^*(N)$	d_{crit}/ℓ	$\varepsilon = 1-\xi^*$	h/c_{max}	Field
8	$8-3\sqrt{7}$	0.51	0.937	0.018	$\mathbb{Q}(\sqrt{7})$
10	$1/7$	0.80	0.857	0.015	$\mathbb{Q}$
12	$5-2\sqrt{6}$	0.97	0.798	0.013	$\mathbb{Q}(\sqrt{6})$
23	φ^{-2}	1.44	φ^{-1}	0.007	$\mathbb{Q}(\sqrt{5})$

Here d_{crit}/ℓ is the geodesic distance from the AdS centre in units of the AdS radius, and $\varepsilon = 1 - \xi^*$ is the boundary proximity (the polygon must sit at $\varepsilon < \varepsilon^*(N)$, i.e., sufficiently close to the AdS boundary, to be stable). The $N = 23$ threshold gives $\varepsilon^* = \varphi^{-1} = (\sqrt{5}-1)/2$: the golden ratio appears as a critical bulk-position parameter in the holographic dictionary.

The two constraints on the polygon are independent:

- *Mass*: $h/c < h/c_{\text{max}}(N)$ (total deficit $< 2\pi$; from the bulk geometry).
- *Position*: $\xi > \xi^*(N)$ (confining potential $>$ OPE repulsion; from the Havelock decomposition and the chiral identity (10.7)).

Together, they define the allowed region in the $(h/c, \xi)$ parameter space for a stable N -gon of heavy operators—a two-parameter family of bootstrap-type bounds, one for each N , all descending from the single identity $\lambda_m = C_1 - m(N-m)/2$.

The confining potential from the OPE kernel. The formula $C_1 = (N-1)(1+\xi^2)/(1-\xi)^2$ can be derived either from the Kimura Hamiltonian for point vortices on $\mathbf{H}^2$ (Paper I, §4) or from the second radial variation of the chordal OPE kernel $\sigma(z, w) = |z-w|/|1-\bar{z}w|$ at the N -gon. The two derivations are mathematically identical: both compute the Hessian of $-\sum_{j<k} \log \sigma(z_j, z_k)$ on the Poincaré disk. The factor $(1+\xi^2)/(1-\xi)^2$ is *not* the conformal factor $\Omega(\sqrt{\xi})^2 = 1/(1-\xi)^2$ alone; the additional $(1+\xi^2)$ comes from the angular structure of σ under radial perturbations.

Relation to the conformal bootstrap. The stability condition $\lambda_m > 0$ and the conformal bootstrap positivity conditions (crossing symmetry + unitarity, [Rattazzi et al. 2008](#)) are both positivity conditions on conformal block coefficients, but in *different channels*:

	Bootstrap	Vortex stability
What is constrained	CFT spectrum $\{\Delta_i\}$	Bulk polygon position ξ
Input	c, Δ_{gap}	$N, h/c$
Positivity condition	OPE $f_{ijk}^2 \geq 0$	$\lambda_m = C_1 - f(m, N) > 0$
Channel	generic 4-point OPE	$\mathbb{Z}_N$ -symmetric

The two are *complementary*: the bootstrap constrains which CFTs exist, while the vortex stability constrains which bulk configurations are stable in a given CFT. The identification $\lambda_m \leftrightarrow$ OPE coefficient would make the stability condition a bootstrap condition for $\mathbb{Z}_N$ -symmetric correlators—a direction requiring the $\mathbb{Z}_N$ conformal block computation sketched below.

10.6 The $\lambda_m \leftrightarrow$ OPE identification

The Havelock eigenvalue is entirely computable from boundary CFT data. The key is to decompose the semiclassical N -point function $\langle \mathcal{O}(z_1) \cdots \mathcal{O}(z_N) \rangle$ of N equally-spaced heavy operators into $\mathbb{Z}_N$ channels: the angular channels give $m(N-m)/2$ (the kinematic OPE Casimir, Proposition [10.3](#)), and the radial confinement gives $C_1(\xi)$.

In the semiclassical (large- c) limit, the N -point function factorises as

$$\log \langle \mathcal{O}(z_1) \cdots \mathcal{O}(z_N) \rangle \approx -2h \sum_{j < k} \log \sigma(z_j, z_k) + O(1), \quad (10.8)$$

where $\sigma = |z_j - z_k|/|1 - \bar{z}_j z_k|$ is the chordal distance. The leading term is the “free OPE energy”; the $O(1)$ correction comes from the Virasoro block (graviton exchange in the bulk).

Under *angular* perturbation $z_k \rightarrow R e^{i(2\pi k/N + \delta\theta_k)}$ at the N -gon, the second variation of the leading term decomposes into $\mathbb{Z}_N$ Fourier modes $\delta\theta_k = \sum_m a_m \cos(2\pi m k/N)$:

$$\delta^2 \left[-2h \sum_{j < k} \log \sigma \right] = 2h \sum_{m=1}^{N-1} m(N-m) |a_m|^2. \quad (10.9)$$

Proposition 10.4 (Angular OPE stiffness). *The coefficient of $|a_m|^2$ in (10.9) is $2h \cdot m(N-m) = 2h \cdot T_m$, where $T_m = m(N-m)$ is the full OPE Casimir (Proposition [10.3](#)). Verified numerically for $N \leq 12$, all modes, to machine precision.*

The missing piece is the confining potential $C_1(\xi)$. We now show it too arises from the boundary.

Theorem 10.5 ($\mathbb{Z}_N$ conformal block = Havelock eigenvalue). *Let $\mathcal{O}_1, \dots, \mathcal{O}_N$ be N identical heavy operators of dimension h , placed at equally-spaced points $z_k = \sqrt{\xi} e^{2\pi i k/N}$ on the boundary of the Poincaré disk. In the semiclassical limit ($c \rightarrow \infty$, h/c fixed):*

- (a) Angular stiffness. *The m -th $\mathbb{Z}_N$ channel of the angular second variation of the N -point function gives $2h \cdot m(N-m)$ (Proposition [10.4](#)).*

- (b) Radial confinement. *The log-radial ($u = \log\sqrt{\xi}$) second variation of $-\sum_{j < k} \log \sigma(z_j, z_k)$ at the N -gon gives the mode-independent confining potential $C_1(\xi) = (N-1)(1+\xi^2)/(1-\xi)^2$.*
- (c) Full identification. *The constrained second variation (angular + radial at fixed angular impulse), divided by $2h$ per chiral sector, is*

$$\lambda_m = \underbrace{C_1(\xi)}_{\text{radial confinement}} - \underbrace{m(N-m)/2}_{\text{one-chiral OPE repulsion}}. \quad (10.10)$$

This is the Havelock eigenvalue, derived entirely from boundary data.

Proof. Part (a) is Proposition 10.4. For part (b): the log-radial second variation of the hyperbolic energy gives $C_1(\xi) = (N-1)(1+\xi^2)/(1-\xi)^2$ by Paper I, equation 5.7 (the Riemannian Havelock identity on $\mathbf{H}^2$). Part (c) combines (a) and (b) via the standard Lagrange-multiplier argument (Paper I, §3).

Identity block dominance for $N > 4$. The theorem requires that the identity Virasoro block dominates the OPE decomposition in each $\mathbb{Z}_N$ channel. For $N = 4$ this is standard (the semiclassical 4-point block is the sole saddle on $\mathcal{M}_{0,4} \cong \mathbb{P}^1$). For $N > 4$: the moduli space $\mathcal{M}_{0,N+1}$ has dimension $2(N-3)$, but the $\mathbb{Z}_N$ symmetry constrains the invariant subspace $\mathcal{M}_{0,N+1}^{\mathbb{Z}_N}$ to dimension 2 (the ring radius and overall phase). This 2-dimensional moduli space admits a single $\mathbb{Z}_N$ -invariant saddle: the symmetric conical metric. The reduction is equivalent to the $\mathbb{Z}_N$ orbifold construction: the N -point function with $\mathbb{Z}_N$ symmetry equals the 1-point function of the twist field σ_m on $\mathbb{C}/\mathbb{Z}_N$, which has no competing OPE channel (a 1-point function has a unique vacuum contribution). Finite- c corrections are controlled by Proposition 10.7. $\square$

Theorem 10.5 identifies λ_m as the m -th $\mathbb{Z}_N$ -channel stiffness of the semiclassical N -point function. No bulk computation is needed: the stability condition $\lambda_m > 0$ is *the confining potential exceeding the holomorphic OPE repulsion in channel m* , a statement in the boundary CFT. The palindromic thresholds $\xi^*(N)$ —at which $C_1(\xi^*) = f_{\max}(N)$ —are visible as algebraic-number bulk positions where the $\mathbb{Z}_N$ conformal block changes sign.

Quantum correction. The $\mathbb{Z}_N$ Fourier decomposition of the Virasoro block at central charge c confirms the identification at all orders. The m -dependent part of $\log F^{(m)}$ is

$$\left[\frac{m(N-m)}{2} - \frac{m^2(N-m)^2}{4c} + O(1/c^2) \right] \log q. \quad (10.11)$$

At $c = \infty$: the leading term is the Havelock Casimir (exact, verified for all $N = 4, \dots, 12$ and all modes). At finite c : the correction $\delta h_m = -f_m^2/c$ is a *polynomial in the Casimir*, so the algebraic structure (palindromic symmetry, number field) is preserved at each order in $1/c$. The fractional correction $|f_m/c| \leq 1/8$ (maximised at $m = N/2$, $c = N^2$) does not change any eigenvalue sign (Proposition 10.7).

10.7 The gravitational prediction

The results of this section yield a concrete, parameter-free prediction for 2+1D gravity.

Proposition 10.6 (Maximum stable polygon in 2+1D gravity). *In linearised 2+1D gravity with cosmological constant Λ ($Gm \ll 1/(4N)$, weak-field regime), the maximum*

number of equal-mass conical singularities that can form a stable regular polygon on a spatial slice satisfies

$$\text{sgn}(N_{\text{crit}} - 7) = -\text{sgn}(\Lambda).$$

Specifically: $N_{\text{crit}} = 7$ for $\Lambda = 0$; $N_{\text{crit}} \leq 6$ for $\Lambda > 0$; and $N_{\text{crit}} \geq 8$ for $\Lambda < 0$, with N_{crit} non-decreasing in $|\Lambda|$ (Proposition 5.1). The transitions occur at the palindromic thresholds $\xi^*(N)$ of Paper I, §5, with $\xi = R^2/\ell_{\text{AdS}}^2$.

Proof. The gravitational Havelock decomposition (Proposition 10.1) gives $\lambda_m = C_1(\xi) - m(N-m)/2$. The stability condition $\lambda_m > 0$ for all m , the monotonicity of $C_1(\xi)$, and the curvature-stability dichotomy (Paper I, Proposition 6.9) transfer directly. $\square$

Validity. The weak-field condition and the stability condition are *independent*: Gm controls the deficit angle (and hence the validity of the logarithmic approximation), while $\xi = R^2/\ell^2$ controls the ring position (and hence the stability). At 1% of the maximum mass ($Gm = 0.01/(4N)$): the total deficit angle is $\approx 3.6^\circ$ for all N , well within the logarithmic regime. The constraint structures (angular momentum in gravity, angular impulse in vortex dynamics) are isomorphic to $O(Gm)$: both arise from the $\text{SO}(2)$ rotational symmetry of the polygon, and the relativistic corrections to the gravitational angular momentum shift the Havelock eigenvalue by $O(Gm) \sim 1\%$ —negligible against binding eigenvalues of $O(1)$.

10.8 The polygon–BTZ phase transition

The Euclidean action comparison between the polygon and the BTZ black hole makes the “Hawking–Page” language precise.

The Euclidean actions. At inverse temperature β , the regularised Euclidean actions of the competing configurations are:

$$I_{\text{polygon}}(\beta) = -\frac{\beta}{8G} + \beta E_{\text{polygon}}, \quad (10.12)$$

$$I_{\text{BTZ}}(\beta) = -\frac{\pi^2 \ell^2}{2G\beta}, \quad (10.13)$$

where $E_{\text{polygon}} = Nm(1-2Gm) + V_{\text{int}}$ is the total energy of the N -gon (kinematic mass + gravitational binding), and $V_{\text{int}} < 0$ is the pairwise Havelock interaction on $\mathbf{H}^2$. The vacuum action $-\beta/(8G)$ is the standard thermal- AdS_3 contribution.

The transition temperature. Setting $I_{\text{polygon}} = I_{\text{BTZ}}$:

$$T_c(N, m, \xi) = \frac{1}{2\pi\ell} \sqrt{1 - 8G E_{\text{polygon}}(N, m, \xi)}, \quad (10.14)$$

which reduces to the standard Hawking–Page temperature $T_{\text{HP}} = 1/(2\pi\ell)$ for $E = 0$. For $T < T_c$: the polygon has lower action ($\Delta I < 0$, polygon preferred). For $T > T_c$: the BTZ has lower action. The total polygon energy E_{polygon} lowers T_c below T_{HP} : more mass in the polygon makes the BTZ transition harder, because the binding energy deepens the polygon potential well.

The negative mode. For $N > N_{\text{crit}}(\xi)$: the polygon has a negative Havelock eigenvalue $\lambda_m < 0$ at mode $m = \lfloor N/2 \rfloor$ (§10.7). In the Euclidean path integral, this is a *negative mode* of the gravitational action [Giombi et al., 2008]: a perturbation that lowers the action, signaling decay. The mode $m = \lfloor N/2 \rfloor$ displaces adjacent masses alternately inward and outward, a pattern that drives pairwise coalescence. In the nonlinear regime, the coalescing masses form a single body exceeding the BTZ threshold $8GM > 1$, completing the decay to a BTZ black hole.

The three-region phase diagram. The polygon–BTZ competition has three distinct regimes:

1. *Stable polygon, preferred* ($N \leq N_{\text{crit}}, T < T_c$): all Havelock eigenvalues are positive and the polygon has lower Euclidean action than the BTZ. The polygon contributes as a stable saddle to the path integral.
2. *Stable polygon, metastable* ($N \leq N_{\text{crit}}, T > T_c$): all eigenvalues are positive but the BTZ has lower action. The polygon is a local minimum separated from the BTZ by a barrier; the tunnelling rate is $\Gamma \sim e^{-S_{\text{BTZ}}}$.
3. *Unstable polygon, decay* ($N > N_{\text{crit}}$): the mode $m = \lfloor N/2 \rfloor$ is a negative mode, and the polygon decays continuously (no barrier) toward the BTZ. The palindromic threshold $\xi^*(N)$ marks the boundary between regions 1/2 and region 3.

The critical exponent. At the palindromic threshold $\rho = \rho^*(N)$, the critical mode $m^* = \lfloor N/2 \rfloor$ vanishes linearly: $\lambda_{m^*}(\rho) \sim \coth(\rho^*) \cdot (\rho - \rho^*)$. The one-loop partition function (Gaussian integration over the zero mode) gives the universal critical exponent $\nu = 1/2$:

$$Z(\rho) \sim \frac{Z_{\text{frozen}}}{|\rho - \rho^*|^{1/2}}, \quad Z_{\text{frozen}} = \frac{2^{3(N-2)/4}}{(N-3)!!}$$

where $(N-3)!! = 1 \cdot 3 \cdot 5 \cdots (N-3)$ is the double factorial, from the Casimir gaps $\lambda_m(\rho^*) = (N-2m)^2/8$. The polygon–BTZ transition is second-order with mean-field exponents.

The entropy gap. In the microcanonical ensemble (fixed total energy), the BTZ *always* has higher entropy: $\Delta S = S_{\text{BTZ}}(E_{\text{polygon}}) > 0$. The polygon is therefore always entropically disfavored; its survival as a metastable state depends entirely on the positivity of the Havelock eigenvalues (the energy barrier), not on entropy. This makes N_{crit} the sharp dividing line: below it, the polygon is a long-lived metastable state; above it, the polygon decays on a dynamical timescale $\tau \sim 1/\sqrt{|\lambda_{\lfloor N/2 \rfloor}|}$.

Microscopic entropy accounting. The BTZ entropy decomposes into discrete and continuous contributions:

$$S_{\text{BTZ}}(N) = \underbrace{S_{1\text{-loop}}(N)}_{\text{continuous}} + \underbrace{\Delta S_{\text{wc}}}_{\text{discrete}}, \quad (10.15)$$

where $S_{1\text{-loop}} = \log(N-3)!! - \frac{3(N-2)}{4} \log 2$ (from the Casimir gaps) and ΔS_{wc} is the wall-crossing contribution. At each palindromic threshold, the $\mathbb{Z}/2$ Galois action ($m \leftrightarrow N-m$) contributes:

N parity	Critical mode	ΔS
Even (self-paired $m = N/2$)	one eigenvalue crosses	$\log 2$ (1 bit)
Odd (paired $m, N-m$)	two eigenvalues cross	$2 \log 2$ (2 bits)

Through $N = 12$: exactly 9 bits (verified numerically). Through $N = 17$ (11 wall-crossings): the discrete part contributes 17 bits (9.3%), the continuous one-loop part 107.8 bits (90.7%). The L^2 -Morse index $\tau(P_-)$ has denominator 2 at each threshold (from the eta invariant shift of $1/2$ due to the zero mode), consistent with the K -theoretic framework.

The one-loop correction. The one-loop partition function around the polygon saddle is $Z_{1\text{-loop}} = (\det' \Delta_{\text{Lich}})^{-1/2}$, where Δ_{Lich} is the Lichnerowicz operator on the multi-cone geometry. By the conformal invariance of the flat-plane backreaction (Proposition 10.2), the leading correction comes from the spatial topology: the spectral bound $|\delta C_1| \leq 2/(4\pi\lambda_1)$ (where λ_1 is the spectral gap of the spatial Laplacian) gives a rigorous bound on the quantum shift of the stability eigenvalues. On the Bolza surface ($\lambda_1 \approx 3.84$): $|\delta C_1| \lesssim 0.05$, which does not shift N_{crit} away from 7 on the flat plane. On the modular surface ($\lambda_1 \approx 91$): $|\delta C_1| \lesssim 0.002$, negligible.

Lorentzian continuation and the BTZ horizon. The curvature coefficient $C_1 = \log(2 \sinh \rho) + b(N)$ admits an analytic continuation $\rho \rightarrow \rho + i\pi/2$ that interchanges the BTZ exterior and interior: $\sinh(\rho + i\pi/2) = i \cosh \rho$, so $C_1^{\text{int}} = \log(2 \cosh \rho) + b(N) + i\pi/2$. The three constant-curvature geometries are connected by the imaginary part of ρ :

$\text{Im}(\rho)$	0	$\pi/2$	π
Geometry	$\mathbf{H}^2$ ($K=-1$)	$\mathbb{R}^2$ ($K=0$)	$\mathbf{S}^2$ ($K=+1$)
Kernel	$\log(2 \sinh \rho)$	$\log(2 \sin t)$	$\log(2 \cos t)$

At each palindromic threshold $\rho^*(N)$: the critical eigenvalue λ_{m^*} crosses zero (*real* side), while the interior resonance has $\text{Im}(\lambda) = \pi/2$ (universal, independent of N and m). The polygon–BTZ transition is a *time–space exchange*: for $\rho > \rho^*$ the mode oscillates in time ($\omega = \sqrt{\lambda_{m^*}}$); for $\rho < \rho^*$ it grows spatially ($\kappa = \sqrt{|\lambda_{m^*}|}$); at $\rho = \rho^*$ the Killing vector degenerates.

Under Lorentzian continuation ($\rho \rightarrow \rho + i\varepsilon$), the partition function acquires a branch cut at the threshold: $Z_L \sim A(N) (\rho - \rho^* + i\varepsilon)^{-1/2}$ with spectral function $\text{Im}[Z_L] \propto |\rho - \rho^*|^{-1/2}$ (a van Hove singularity).

The Wheeler–DeWitt equation on the palindromic minisuperspace. The single degree of freedom ρ (geodesic radius of the polygon) and the potential $V(\rho) = \lambda_{m^*}(\rho)$ (the critical Havelock eigenvalue) define a minisuperspace Wheeler–DeWitt equation:

$$\left[-\frac{\hbar^2}{2c} \frac{d^2}{d\rho^2} + \log(2 \sinh \rho) + b(N) - \frac{m(N-m)}{2} + \delta_m \right] \Psi = 0, \quad (10.16)$$

with $c = 12N^2$ and $\hbar = 1/\sqrt{c}$. Each term has a definite origin:

Term	Expression	Origin
Quantum kinetic	$-\hbar^2/(2c) \Psi''$	Gravity dynamics (the dynamics of ρ)
Classical potential	$\log(2 \sinh \rho)$	$\mathbf{H}^2$ Green's function
Vacuum energy	$b(N) \sim c/12$	Orbifold CFT vacuum (Todd class)
Matter	$m(N-m)/2$	Twist-field dimension (first Chern class)
One-loop	δ_m	Weyl anomaly (traceless, palindromic)

Three regimes.

1. **Classical** ($c \rightarrow \infty$): the kinetic term vanishes; the constraint reduces to $V(\rho) = 0$, i.e., $\lambda_{m^*}(\rho^*) = 0$ —the palindromic threshold, which is the Einstein equation of the vortex–gravity system.
2. **Semiclassical** (WKB): $\Psi \sim \exp(\pm(i/\hbar) \int \sqrt{2cV} d\rho)$, oscillatory in the exterior ($V > 0$, sinh-branch) and evanescent in the interior ($V < 0$, cosh-branch). The continuation $\rho \rightarrow \rho + i\pi/2$ maps between them (§??, Lorentzian paragraph).
3. **Quantum** (near threshold): $\Psi \sim \text{Ai}((2c \coth \rho^*)^{1/3}(\rho - \rho^*))$, smoothing the classical horizon over the universal width

$$\delta\rho = \frac{1}{(2c)^{1/3}} = \frac{1}{(24N^2)^{1/3}} \quad (10.17)$$

(verified to 4 digits for $N = 8, \dots, 20$: $\delta\rho/N^{-2/3} \rightarrow 2^{-1/3}$). At $c = \infty$: $\delta\rho \rightarrow 0$ (sharp classical threshold). At finite c : the horizon has quantum width $O(c^{-1/3})$.

Proposition 10.7 (Eigenvalue non-crossing at finite c). *For $N \geq 6$ and $c \geq N^2$, the mode-dependent part of the $O(1/c)$ Virasoro block correction satisfies*

$$\frac{|\delta f_m - \delta f_{m+1}|}{f(m+1, N) - f(m, N)} = \frac{2m(N-m-1) + (N-1)}{2c} \leq \frac{1}{4} - \frac{1}{2N^2} < \frac{1}{4}. \quad (10.18)$$

In particular, no eigenvalue crossing occurs and the critical mode index $m_{\text{crit}} = \lfloor N/2 \rfloor$ is preserved.

Proof. The leading $1/c$ correction to the Casimir is $\delta f_m = -f_m^2/c = -m^2(N-m)^2/(4c)$. The difference factors as $\delta f_m - \delta f_{m+1} = (N-2m-1)[2m(N-m-1) + (N-1)]/(4c)$. Dividing by the Casimir gap $f(m+1) - f(m) = (N-2m-1)/2$ gives the ratio $R(m) = [2m(N-m-1) + (N-1)]/(2c)$. The maximum over $1 \leq m < N/2$ is at $m = N/2 - 1$, giving $R_{\text{max}} = (N^2 - 2)/(4N^2) = 1/4 - 1/(2N^2) < 1/4$. Since $R < 1$ for all adjacent pairs: the residual gap is at least $3(N-2m-1)/8 > 0$, and no reordering occurs. $\square$

Remark 10.1 (Algebraic field preservation). *The Zamolodchikov c -recursion gives the correction as a rational function of c , with only $\mathbb{Z}_N$ -invariant channels contributing. At $c = N^2$: the palindromic coefficients $A(c), B(c)$ remain in $\mathbb{Q}(\cos(2\pi/N))$ —the same number field as at $c = \infty$. The palindromic form $A\xi^2 + B\xi + A = 0$ is exact at all c (the conformal inversion $\xi \leftrightarrow 1/\xi$ is a geometric symmetry of $\mathbf{H}^2$, not a perturbative artefact).*

The palindromic threshold $\xi^*(N)$ is the critical AdS radius at which the polygon–BTZ transition occurs, and the algebraic number field $\mathbb{Q}(\sqrt{D})$ of the threshold is the arithmetic invariant of the phase boundary.

10.9 The three-layer decomposition

Proposition 10.8 (Orbifold–Havelock correspondence). *The Havelock eigenvalue on $\mathbf{H}^2$ is exactly the $\mathbb{Z}_N$ orbifold CFT ground-state energy at $c = 12N^2$ with angular impulse constraint:*

$$\lambda_m = -E_0(m; c=12N^2) + \mu_L(\xi), \quad (10.19)$$

where $E_0(m) = -N/2 + m(N-m)/2$ is the twisted-sector ground state and $\mu_L(\xi) = C_1(\xi) - N/2$ is the Lagrange multiplier. The central charge $c = 12N^2$ is determined by matching the mode-dependent coefficient $c/(24N^2) = 1/2$; all curvature dependence enters through μ_L alone. The N^2 scaling follows from the ‘t Hooft structure of the Hamiltonian ($N(N-1)/2 \sim N^2$ pairwise terms); the factor $12 = 24/2$ converts between Virasoro and Havelock normalisations.

Proof. The orbifold ground-state energy $E_0(m) = -c/(24N) + cm(N-m)/(24N^2)$ at $c = 12N^2$ gives $E_0 = -N/2 + m(N-m)/2$. Then $-E_0(m) + \mu_L = [N/2 - m(N-m)/2] + [C_1(\xi) - N/2] = C_1(\xi) - m(N-m)/2 = \lambda_m$. $\square$

Proposition 10.9 (Equivariant Riemann–Roch). *The $\mathbb{Z}_N$ -equivariant spectral index of the vortex Laplacian on the cusped Poincaré disk satisfies*

$$\text{ind}_m = \underbrace{\frac{m(N-m)}{2}}_{c_1\text{-term}} + \underbrace{b(N)}_{\text{Todd term}} + \delta_m, \quad (10.20)$$

where $b(N) = N(N+1)/12 - \log 2 + \log N/(N-1)$ is independent of both m and ρ (verified to 10^{-16} for $N = 3, \dots, 24$). The central charge $c = 12b(N) = N^2 + O(N)$ is the same for every mode m , extracted from the Todd correction alone.

Proof. The Havelock eigenvalue satisfies $\lambda_m + f_m = C_1(\rho) + \delta_m$, and $C_1(\rho) = \log(2 \sinh \rho) + b(N) + O(e^{-\rho})$ (the radial dependence is absorbed by the Green’s function). Subtracting $\log(2 \sinh \rho)$: $\text{ind}_m := \lambda_m + f_m - \log(2 \sinh \rho) = b(N) + \delta_m$. Since δ_m is traceless ($\sum \delta_m = 0$ to 10^{-15}) and palindromic ($\delta_m = \delta_{N-m}$), the mode-independent part $b(N)$ is extracted exactly. $\square$

The three-layer decomposition $\lambda_m = C_1 - f_m + \delta_m$ is the *spectral equivariant Riemann–Roch theorem*: f_m is the c_1 -contribution (twist-field dimension), $b(N)$ is the Todd-class correction (vacuum energy $\approx c/12$), and δ_m is the one-loop remainder (the Weyl anomaly). The central charge is not a fitted parameter: it is $c = 12b(N)$, determined by the index theorem.

Remark 10.2 (Honest negatives and the B_2 boundary). *The connection between the aliasing $\langle D \rangle = N(N+1)/12$ and the Virasoro central charge $c/12$ is structural (both arise from $B_2 = 1/6$) but not algebraic: the Havelock modes do not satisfy Virasoro commutation relations at finite N (the aliasing $1/12$ and the Virasoro $1/12$ share $B_2 = 1/6$ but live at different algebraic levels: eigenvalue shift vs. central charge). The one-loop partition function $Z_N = \prod |\lambda_m|^{-1/2}$ does not match $|\eta(\tau)|^{-2k}$ for any fixed $k(N)$; however, the vacuum energy $c/12 = N^2/12$ does match the η zero-point energy—the failure is in the ρ -dependent (bulk) part, not the ρ -independent (vacuum) part. The higher-order mode coupling is W_N (quartic: real, first contributing order; cubic: purely imaginary, a Berry phase), with effective central charge $c_{\text{eff}} \approx N$ (one free boson per vortex).*

The value $c = 12N^2$ is not merely a matching condition; it makes five independently testable predictions beyond the Casimir: (i) the exact offset $b(N) = N(N+1)/12 - \log 2 + \log N/(N-1)$ (verified to 10^{-16} , Paper I); (ii) the mean aliasing $\langle D \rangle = N(N+1)/12$ (exact); (iii) the growth law $\rho^*/f_{\text{crit}} \rightarrow 1/3$ (Paper I); (iv) the quartic coupling $Q/f_{\text{crit}}^2 = N/48$ (exact for even N); (v) the mode-independent Laplacian $\Delta_{\mathbf{H}^2} h = -1/2$. All five follow from $c = 12N^2$ through the B_2 -tower: the Casimir $m(N-m)/2$ is $O(c)$; the Weyl anomaly δ_m is $O(c^0)$ (topological); the Virasoro block correction is $O(1/c) \approx 10^{-5}$ (§10.5).

The tower extends to a *Magri bihamiltonian hierarchy*: the vortex system admits a second Poisson structure (the Casimir bracket) whose recursion eigenvalues are $f_m = m(N-m)/2$. The k -th conserved quantity has mode eigenvalue f_m^k , with the physical derivatives $d_{2k} = \partial^{2k} H / \partial a^{2k} = \alpha_k N f_m^k$ satisfying $\alpha_1 = 1, \alpha_2 = 4$ (verified for 6 values of N to $< 0.04\%$). The cubic Hamiltonian is purely imaginary (a Berry phase), and the quartic coupling $Q = (N/48)f^2$ lies inside the hierarchy ($H_4^{\text{diag}} \propto I_2$), guaranteeing conservation $\{I_j, H_4\} = 0$ by the Magri theorem. The three-layer decomposition $\lambda_m = C_1 - f_m + \delta_m$ is the truncation of this hierarchy at level 1.

On arithmetic surfaces (Bolza, modular, etc.), the eigenvalue acquires a topological correction that is invisible to the confining potential $C_1(\xi)$:

$$\lambda_m = \underbrace{C_1(\xi)}_{\text{Ricci scalar}} - \underbrace{\frac{m(N-m)}{2}}_{\text{Casimir}} + \underbrace{\delta_m(N)}_{\text{Weyl anomaly}}. \quad (10.21)$$

Each layer has a distinct character:

1. $C_1(\xi)$: continuous, surface-dependent, mode-independent. Scales approximately as $C_1 \propto \log(\xi)$ over the range $\xi \in [0.01, 10]$ (with $R^2 = 0.98$); the exact formula is $C_1 = (N-1)(1+\xi^2)/(1-\xi)^2$.
2. $f(m, N) = m(N-m)/2$: discrete, universal (surface-independent). The $\mathbb{Z}_N$ Casimir from the Havelock identity.
3. $\delta_m(N)$: topological, UV-finite, ξ -independent (variation $< 10^{-13}$), and exactly traceless ($\sum_{m=1}^{N-1} \delta_m = 0$ to 10^{-15}). Carries all number-theoretic content: Hecke eigenvalues, palindromic structure, Langlands data.

By structural analogy with the decomposition of the Riemann curvature tensor: C_1 plays the role of the Ricci scalar, $f(m, N)$ the traceful part, and δ_m the traceless Weyl-like anomaly. In $\text{AdS}_3/\text{CFT}_2$, the Brown–Henneaux scaling $C_1 \propto \log(\xi)$ says that the Weyl-like anomaly δ_m is the topological correction that carries the arithmetic of the underlying surface, contributing corrections of order $\|\delta\|/C_1 \sim 2\%$ at typical parameter values.

The Weyl norm scales as $\|\delta\|^2 \sim N^{4.2}(\log N)^{4.1}$ for $N \geq 40$ (not a pure power law). The triangle ($N = 3$) is *anomaly-free*: $\delta_m = 0$ exactly for all Green’s functions and curvatures, because $\cos(2\pi/3) = \cos(4\pi/3)$ forces $G_1 = G_2$. Under Casimir-conserving mergers $N_1 + N_2 \rightarrow N_3$ ($N_3^3 - N_3 = N_1^3 - N_1 + N_2^3 - N_2$): the unique anomaly-*destructive* reaction is $4+4 \rightarrow 5$ (88% Weyl loss); all other mergers create Weyl content (gains 46% to 849%). The Weyl entropy per degree of freedom $s(N) = \log \|\delta\|^2 / \sum f$ peaks at $N \approx 9$ –10, coinciding with the observed planetary polygon range ($N = 5, 6, 8$).

The anomaly δ_m has two independent origins, which the parameter space (Δ, ξ) separates cleanly:

1. *Interaction anomaly* ($\Delta > 0$ on $\mathbb{R}^2$): deviation of $V = r^{-2\Delta}$ from $-\log r$. Large ($\|\delta\|^2 = O(1\text{--}100)$) but arithmetically trivial.

2. *Geometric anomaly* ($\Delta = 0$ on $\mathbf{H}^2$): deviation of the hyperbolic Green's function from the flat one. Small ($\|\delta\|^2 = O(0.01\text{--}100)$) but carries all number theory: palindromic polynomials, Hecke eigenvalues, Langlands data.

The RG flow ($\Delta \rightarrow 0$) strips away the large, trivial interaction anomaly, leaving behind the small geometric anomaly as the residue that survives at the conformal fixed point. The number theory is not emergent from the flow; it is a property of the conformal fixed point, activated by curvature. The three-layer structure has a symplectic origin: the Kirillov form ω_{KR} (mode-independent) and the Weil–Petersson form ω_{WP} (mode-dependent) are *not* proportional; their discrepancy $\omega_{\text{WP}} - \alpha \omega_{\text{KR}} \propto f(m, N)$ is the Casimir (Layer 2), and the residual beyond this linear relation is the Weyl anomaly (Layer 3).

The Bernoulli number $B_2 = 1/6$ that governs the growth law $\rho^* \sim N^2/24$ (Paper I, §6) is the asymptotic mode average of the $\mathbb{Z}_N$ orbifold Casimir: $\langle m(N-m) \rangle / N^2 \rightarrow 1/6$. The Havelock aliasing $\langle D \rangle = N(N+1)/12 \sim N^2 B_2/2$ is the quadratic part of the orbifold Virasoro central extension, evaluated via Hurwitz zeta regularization. The same B_2 gives the Dedekind eta exponent $q^{1/24}$ and the regularised sum $\zeta(-1) = -1/12$: three levels of the modular tower, unified by the Euler–Maclaurin correction to the log-sine kernel on $\mathbb{Z}/N\mathbb{Z}$.

10.10 Higher-dimensional extension: the $\text{CFT}_3/\text{AdS}_4$ Casimir

The two-dimensional results extend to conformal field theory in any dimension d via the power-law interaction $V(x, y) = |x - y|^{-2\Delta}$, where Δ is the conformal dimension. For $\Delta \rightarrow 0$: $V \rightarrow -\log|x - y|$ (the 2D case). For $\Delta = 1$: the CFT_3 two-point function on S^2 (the boundary of AdS_4).

Proposition 10.10 (Generalised Havelock Casimir). *For N equally-spaced points on a ring on S^2 , interacting via $V = |x - y|^{-2\Delta}$, the angular Hessian eigenvalues factorise as $\mu_m = g(\theta_0, \Delta) \cdot h(m, N, \Delta)$, where the universal Casimir is*

$$h(m, N, \Delta) = \sum_{p=1}^{N-1} \frac{1 + 2\Delta \cos^2(\pi p/N)}{\sin^{2\Delta+2}(\pi p/N)} \left(1 - \cos \frac{2\pi m p}{N}\right). \quad (10.22)$$

For $\Delta = 0$: $h(m, N, 0) = 2m(N-m)$. The ratio $h(m)/h(1)$ is independent of θ_0 , confirming that the Havelock decomposition extends to all Δ .

The constrained Hessian on the flat plane with $V = |x - y|^{-2\Delta}$ gives an N -dependent stability boundary $N_{\text{crit}}(\Delta)$:

Δ	0	0.5	1	2	5	10
N_{crit}	7	5	5	5	4	4

The classical $N_{\text{crit}} = 7$ is an isolated phenomenon of the logarithmic interaction. For the CFT_3 value $\Delta = 1$: $N_{\text{crit}} = 5$.

Under the Green's function identification $\Delta = (d-2)/2$, the table acquires a dimensional interpretation:

Dimension d	2	3	4	5	≥ 6
$\Delta = (d-2)/2$	0	1/2	1	3/2	≥ 2
N_{crit}	7	5	5	5	4–5

The hexagon ($N = 6$, Saturn) is a *uniquely two-dimensional* phenomenon: it is marginally stable at $d = 2$ ($\lambda_{\min} = 1/2$) but unstable for $d \geq 3$ ($\lambda_{\min} = -0.02$ at $d = 3$). The pentagon ($N = 5$) survives in all tested dimensions $d \leq 10$.

Proposition 10.11 (Soft Havelock decomposition in dimension d). *For the interaction $V = r^{-(d-2)}$ ($d \geq 3$) on the N -gon ring, the angular Casimir decomposes as*

$$\sigma_m^{(d)} = a(N, d) \frac{m(N-m)}{2} + R_m(N, d), \quad (10.23)$$

where $a(N, d)$ is the leading Casimir coefficient (determined by least-squares fit) and R_m is the residual. For all tested $N = 5, \dots, 20$ and $d = 3, 4, 5$: the residual satisfies $|R_m - R_{m'}| < a|f_m - f_{m'}|$ for every adjacent pair, so the mode ordering is preserved. The critical mode is always $m = \lfloor N/2 \rfloor$ (same as $d = 2$), and the stability threshold has the same palindromic form as the two-dimensional case, with coefficients that depend on d .

For $d = 2$ ($s = 0$): $a = 1$ exactly and $R_m = 0$ (the Havelock identity). This is the *unique* dimension where the Casimir is an exact polynomial—because $\csc^2(\pi p/N)$ has a closed discrete Fourier transform ($\sum \csc^2 \cdot (1 - \cos) = 2m(N-m)$), while $\csc^s$ for $s \neq 2$ does not. The *exact* algebraic structure of Paper I (palindromic polynomials, algebraic number fields, modular forms) is specific to $d = 2$, a consequence of the uniqueness of the $\csc^2$ weight function among power-law potentials. However, the *structural skeleton* survives in all dimensions:

1. **Mode ordering.** The critical mode is $m = \lfloor N/2 \rfloor$ in every d (Proposition 10.11). The residual R_m is too small to reorder eigenvalues.
2. **Palindromic symmetry.** The symmetry $\sigma_m = \sigma_{N-m}$ holds exactly for all d (from the discrete cosine: $\cos(2\pi pm/N) = \cos(2\pi p(N-m)/N)$). On curved spaces, conformal inversion $\xi \leftrightarrow 1/\xi$ is geometric (a property of the Poincaré ball, not of the interaction), so the threshold equation retains palindromic form.
3. **Finite- c algebraic structure.** The $O(1/c)$ Virasoro block correction $\delta h_m = -f_m^2/c$ is a *polynomial in the Casimir* (equation 10.11), preserving the algebraic structure at every order.
4. **Weyl anomaly shape.** The normalised anomaly profile $\delta_m/\|\delta\|$ has inner products > 0.98 between $\Delta = 0$ and $\Delta = 2$ for $N \leq 6$: the number-theoretic fingerprint transfers to $d = 3$ with $< 2\%$ distortion.

The $d = 2$ case is where the algebra is *exact*; the $d \geq 3$ case is where it is *approximate but structurally preserved*. The mechanism by which the logarithmic structure survives in 3+1D is the soft Havelock decomposition: the Casimir $m(N-m)/2$ is the dominant term at every d , with the interaction-specific residual R_m as a controlled perturbation.

Remark 10.3 (Transcendence of the transition dimensions). *The critical value $\Delta^*(N)$ where N_{crit} drops below N is expected to be transcendental for $N \geq 5$. The threshold equation involves $\sin^{-2\Delta}(\pi p/N)$ for different p with $\mathbb{Q}$ -linearly independent logarithms. For $N = 5$ (two distinct chord distances, ratio $= \varphi$), the Gelfond–Schneider theorem directly applies: if Δ^* were algebraic and irrational, $\varphi^{2\Delta^*+2}$ would be transcendental, contradicting the threshold equation. For general N , a rigorous proof requires the six-exponential theorem or Baker’s theorem on linear forms in logarithms. This contrasts with the palindromic thresholds $\xi^*(N)$ on $\mathbb{H}^2$, which lie in explicit algebraic number fields.*

For $N = 5$, the pentagon identity $\sin(2\pi/5)/\sin(\pi/5) = \varphi$ (the golden ratio) gives the eigenvalue as $\lambda_2(\Delta) = P_0(\Delta) + P_1(\Delta) \varphi^{-(2\Delta+2)}$, where P_0, P_1 are algebraic functions of Δ . The transition occurs at $\varphi^{2\Delta^*+2} = P_1/(-P_0) \approx 25.8$, giving $\Delta^* \approx 2.378$.

Theorem 10.12 (RG monotonicity of the Casimir ratio). *For any $N \geq 4$ and $m_{\text{crit}} = \lfloor N/2 \rfloor$, the generalised Casimir ratio $R(\Delta) = h(m_{\text{crit}}, N, \Delta)/h(1, N, \Delta)$ is strictly increasing on $[0, \infty)$, with exact limits*

$$R(0) = \frac{m_{\text{crit}}(N - m_{\text{crit}})}{N - 1}, \quad R(\infty) = \left[\frac{\sin(\pi m_{\text{crit}}/N)}{\sin(\pi/N)} \right]^2. \quad (10.24)$$

Proof. By the Chebyshev identity, the mode amplification factor $f(p) = (1 - \cos(2\pi m_{\text{crit}} p/N))/(1 - \cos(2\pi p/N))$ achieves its global maximum at $p = 1$ (and its palindromic twin $p = N - 1$):

Even $N = 2k$: $\sin(\pi k p/N) = \sin(\pi p/2) \in \{0, \pm 1\}$, so $f(p) = 1/\sin^2(\pi p/N)$ for odd p and $f(p) = 0$ for even p . Since $\sin(\pi p/N) \geq \sin(\pi/N)$ for $p \geq 2$: $f(p) \leq f(1)$.

Odd $N = 2k + 1$: The identity $\sin(k p p/N) = \sin(\pi p/2 - \pi p/(2N))$ (using $k/N = 1/2 - 1/(2N)$) gives explicit formulas: $f(p) = 1/(4 \sin^2(\pi p/(2N)))$ for odd p , and $f(p) = 1/(4 \cos^2(\pi p/(2N)))$ for even p . Since $f(1) = 1/(4 \sin^2(\pi/(2N)))$: for odd $p \geq 3$, $\sin(\pi p/(2N)) \geq \sin(\pi/(2N))$ (monotonicity of $\sin$ on $(0, \pi/2)$), so $f(p) \leq f(1)$; for even $p \leq N - 1$, $\cos(\pi p/(2N)) \geq \sin(\pi/(2N))$ (since $\pi p/(2N) + \pi/(2N) \leq \pi/2$ iff $p \leq N - 1$), so $f(p) \leq f(1)$.

The normalised weight $\mu_p(\Delta) \propto \tilde{w}(p, \Delta) c_1(p)$ concentrates on $\{1, N - 1\}$ as $\Delta \rightarrow \infty$, because $\tilde{w}(p)/\tilde{w}(1) = O([\sin(\pi/N)/\sin(\pi p/N)]^{2\Delta}) \rightarrow 0$ for $p \notin \{1, N - 1\}$. Since $R(\Delta) = \mathbb{E}_\mu[f]$ and μ concentrates at the global maximum of f : $dR/d\Delta = \text{Cov}_\mu(d \log \tilde{w}/d\Delta, f) > 0$ (both f and $d \log \tilde{w}/d\Delta$ are maximised at $p = 1$, giving positive covariance). $\square$

Corollary 10.13 (Stability c-theorem analogue). *$N_{\text{crit}}(\Delta)$ is non-increasing in Δ : once a polygon type loses stability, it never regains it. The “polygon entropy” $S(\Delta) = \log N_{\text{crit}}(\Delta)$ is a monotonically non-increasing function of the conformal dimension, flowing from $S(0) = \log 7$ (the Havelock IR fixed point) to $S(\infty) = \log 3$ (the triangular UV fixed point).*

10.11 Triangle universality and dynamical triangulation

The equilateral triangle ($N = 3$) is the *unique* regular polygon that is dynamically stable for all interaction ranges and quantum deformation levels.

Theorem 10.14 (Triangle universal stability). *The equilateral triangle ($N = 3$) is the unique regular polygon that is stable on the flat plane for all $\Delta \geq 0$, all TQFT levels $k \geq 3$, and whose tessellation $\{3, 6\}$ is the unique stable regular lattice for all $\Delta > 0$.*

Proof sketch. Three independent arguments:

1. *Equal chords.* The triangle is the unique regular polygon where all chord distances are equal ($\sin(\pi/3) = \sin(2\pi/3) = \sqrt{3}/2$). No range-dependent competition between nearest and next-nearest neighbors $\implies$ stability is Δ -independent.
2. *Single mode.* The triangle has one nontrivial $\mathbb{Z}_3$ mode ($m = 1$, with $m = 2$ its palindromic conjugate). The Havelock eigenvalue $\lambda_1 = (N - 1) - 1 = 1 > 0$. Mode competition requires at least two independent modes ($N \geq 4$).

3. *Zero Weyl anomaly.* The anomaly δ_m is traceless: $\sum \delta_m = 0$. With one mode: $\delta_1 = 0$ identically. Verified: anomaly norm = 2.2×10^{-16} (machine zero) on the square torus.

For the tessellation claim: the dynamical matrix (phonon spectrum) of the triangular lattice under $V = r^{-2\Delta}$ has all eigenvalues positive for *every* $\Delta > 0$ (verified by lattice sums with convergence study and Richardson extrapolation). The minimum phonon frequency is at the K -point of the Brillouin zone. The square lattice (coordination 4) has a permanent X -point instability (stripe mode: $\omega^2(X) < 0$ for all Δ); the hexagonal lattice (coordination 3) has a permanent K -point instability. The triangular lattice is the *only* regular lattice that is ever stable, at any interaction range. $\square$

Remark 10.4 (Dynamical triangulation). *In causal dynamical triangulations (CDT) and Regge calculus, spacetime is discretised into simplices (triangles in 2D). This triangulation is imposed by hand as a regularisation. Theorem 10.14 suggests it may be imposed by physics: among regular lattices, the triangular lattice is the unique stable configuration, so the gravitational path integral $Z = \sum_{\text{tilings}} e^{-S}$ may be dominated by triangulated configurations—the square and hexagonal alternatives are saddle points with negative phonon modes at all Δ . The extension to irregular triangulations (as used in CDT) remains open.*

10.12 The number-theoretic gravity partition function

The one-loop partition function of the vortex system organises naturally by algebraic number field.

Proposition 10.15 (Gravity partition function). *The quantum gravity partition function for the vortex-polygon system on $\mathbf{H}^2$ is*

$$Z_{\text{grav}} = \sum_{N=3}^{\infty} e^{-b(N)} Z_{\text{frozen}}(N), \quad (10.25)$$

where $b(N) = N(N+1)/12 - \log 2 + \log N/(N-1)$ is the gravitational action and $Z_{\text{frozen}} = 2^{3(N-2)/4}/(N-3)!!$ is the one-loop determinant (even N ; the odd- N case is computed from Casimir gaps). The series converges super-exponentially: $e^{-b(N)} \sim e^{-N^2/12}$. Numerically, $Z_{\text{grav}} = 1.3304$ and $N = 3, \dots, 8$ capture 99.99% of the total.

Proof. Convergence: $e^{-b(N)} = O(e^{-N^2/12})$ decays faster than any polynomial, while $Z_{\text{frozen}} = O(N^{O(N)})$ grows sub-exponentially; the product vanishes. The numerical value $Z = 1.3304$ is computed by direct summation to $N = 20$ (the tail $N > 20$ contributes $< 10^{-20}$). $\square$

The partition function decomposes by *trace field*—the algebraic number field generated by $\cos(2\pi/N)$:

Trace field	N values	Fraction of Z
$\mathbb{Q}$ (rational)	3, 4, 6	90.98%
$\mathbb{Q}(\sqrt{5})$ (golden)	5, 10	8.26%
$\mathbb{Q}(\sqrt{2})$ (silver)	8	0.42%
$\mathbb{Q}(\cos 2\pi/7)$ (cubic)	7	0.34%
All higher fields	≥ 9	$< 0.01\%$

The rational field dominates; the golden ratio field gives the leading correction; the entire algebraic tower is exponentially suppressed by $e^{-N^2/12}$.

The Baum-Connes assembly map $\mu: K_0^{\mathbb{Z}_N}(\mathbf{H}^2) \rightarrow K_0(C_r^*(\mathbb{Z}_N))$ sends the equivariant index $\text{ind}_m = f(m, N) + b(N)$ to the K-theory class with trace $c = 12b(N)$. This is mode-independent (the *same* c for all m) and approximately integer: $c = N^2 + O(N)$ for all tested $N = 3, \dots, 24$. The discrepancy $c - N^2 \approx N + O(1)$ is the boundary self-energy ($N/12$) plus regularisation terms.

10.13 From polygon stability to the Einstein equations

We now present the main result of the gravitational theory: a derivation of the vacuum Einstein equations in 2+1 dimensions from the Havelock eigenvalue structure. The derivation has six steps, each proved in the preceding sections. We state the theorem first, then collect the proof.

Theorem 10.16 (Polygon stability implies the Einstein equations in 2+1D). *Let (S, g) be a two-dimensional Riemannian manifold serving as the spatial slice of a 2+1-dimensional spacetime. Place N equal point masses of mass m in a regular polygon ring of geodesic radius ε centred at a point $x \in S$, interacting via the scalar Green's function of (S, g) . If the following three conditions hold:*

(H1) Havelock structure: *the angular Hessian eigenvalues have the form $\lambda_m = C_1(x, \varepsilon) - m(N-m)/2$ with C_1 independent of m ;*

(H2) Hamiltonian constraint: *$C_1(x, \varepsilon) = f(m^*, N)$ at the critical mode $m^* = \lfloor N/2 \rfloor$;*

(H3) Universality: *conditions (H1)–(H2) hold at every point $x \in S$;*

then the scalar curvature of S is constant, $R(x) = R_0$ for all x , and

$$R_{\mu\nu} = \Lambda g_{\mu\nu}, \quad \Lambda = \frac{3[f(m^*, N) - (N-1)]}{(N-1)\varepsilon^2}. \quad (10.26)$$

These are the vacuum Einstein equations of 2+1-dimensional gravity.

Proof. The proof proceeds in six steps.

Step 1. Polynomial Casimir forces logarithmic interaction. Condition (H1) requires the angular Casimir to be the exact polynomial $f(m, N) = m(N-m)/2$. By the csc^2 uniqueness theorem (Paper I, Lemma 4.1), this polynomial arises *only* from the weight function $\text{csc}^2(\pi p/N)$, which is the discrete Fourier transform of the logarithmic kernel $-\log|z-w|$. Therefore the interaction potential is $V(z, w) = -G_S(z, w)$, where G_S is the Green's function of the scalar Laplacian on (S, g) .

Step 2. Canonical quantisation gives the WDW equation. The polygon dynamics on the angular-impulse surface $\{J = \text{const}\}$ reduces to a single degree of freedom ρ (the geodesic radius) with potential $V(\rho) = C_1(x, \rho) - f(m^*, N)$. Canonical quantisation with Planck mass $c = 12N^2$ gives the Wheeler–DeWitt equation (10.16).

Step 3. Classical limit yields the Hamiltonian constraint. In the limit $c \rightarrow \infty$: the kinetic term $-\hbar^2 \Psi''/(2c)$ vanishes and the WDW equation reduces to $V(\rho) = 0$, i.e. $C_1(x, \varepsilon) = f(m^*, N)$. This is condition (H2)—the Hamiltonian constraint $H_{\text{ADM}} = 0$ of the minisuperspace.

Step 4. Small-ring expansion reveals local curvature. The Green's function on any Riemannian surface has the short-distance expansion $G_S(z, w) = -(1/2\pi) \log d(z, w) + \gamma_S(z) + O(d^2)$, where $d(z, w)$ is the geodesic distance and γ_S encodes the global geometry. Summing over the N -ring at radius ε centred at x :

$$C_1(x, \varepsilon) = (N-1) \left[1 + \frac{R(x) \varepsilon^2}{6} + O(\varepsilon^4) \right], \quad (10.27)$$

where $R(x)$ is the scalar curvature at x (the ε^2 coefficient comes from the Hadamard parametrix of the Green's function).

Step 5. The test polygon principle forces $R = \text{const}$. Condition (H3) requires (10.27) to equal $f(m^*, N)$ at *every* x . Since $f(m^*, N)$ and N are numbers independent of x :

$$(N-1) \frac{R(x) \varepsilon^2}{6} = f(m^*, N) - (N-1) \quad \text{for all } x.$$

The left side depends on x (through $R(x)$); the right side is constant. Therefore $R(x) = R_0$ for all x , with $R_0 = 6[f(m^*, N) - (N-1)]/[(N-1)\varepsilon^2]$.

Step 6. Mode-independence + $R = \text{const}$ implies Einstein in any d . In $d = 3$ (2+1D): the Weyl tensor has $d(d+1)(d+2)(d-3)/12 = 0$ components, so $R = \text{const}$ alone gives $R_{\mu\nu} = \Lambda g_{\mu\nu}$.

In $d \geq 4$: $R = \text{const}$ gives only the *trace* of the Einstein equations. However, condition (H1) provides the missing information. On a space with nonzero Weyl tensor, the Green's function acquires angular dependence (an $\ell = 2$ quadrupole from W_{ijkl} at leading order). This angular dependence makes C_1 depend on m , violating (H1). Therefore (H1) at every point x forces the Green's function to be isotropic at every point, i.e. the spatial geometry is *maximally symmetric*: $R_{ij} = (R/d) g_{ij}$ with $W_{ijkl} = 0$. Promoting to the $(d+1)$ -dimensional spacetime: $R_{\mu\nu} = \Lambda g_{\mu\nu}$.

In $d = 4$ the Weyl tensor has 10 independent components. Its leading effect on G is at spherical harmonic $\ell = 2$ (5 real components). A polygon with $N \geq 6$ provides $N-1 \geq 5$ independent $\mathbb{Z}_N$ Fourier modes, which is sufficient to constrain all 5 quadrupole components—so $N = 6$ (Saturn's hexagon!) is the *smallest polygon that probes the full Weyl tensor in 3+1D*. $\square$

Remark 10.5 (The cosmological constant from the polygon number). *The sign of Λ is fixed by the polygon: $N \leq 6$: $f(m^*) < N-1$, so $\Lambda < 0$ (AdS); $N = 7$: $f(3, 7) = 6 = N-1$, so $\Lambda = 0$ (flat); $N \geq 8$: $f(m^*) > N-1$, so $\Lambda > 0$ (dS). The flat threshold $N = 7$ —Havelock's 1931 stability boundary—sits exactly at $\Lambda = 0$.*

Remark 10.6 (The derivation in $d \geq 3$: equivalence, not obstruction). *Steps 1–5 are dimension-independent: linearised gravity in any d interacts via the Green's function of the spatial Laplacian (Step 1); the WDW equation and its classical limit do not depend on d (Steps 2–3); the Hadamard parametrix $G(x, y) = U/\sigma^{(d-2)/2} + V \log \sigma + W$ contains $R(x)$ at order ε^2 in every d (Step 4); and the x -independence of $h(m^*, N, \Delta)$ forces $R = \text{const}$ regardless of d (Step 5). Step 6 closes differently by dimension:*

Dimension	Weyl components	Status
$d = 2$ (2+1D)	0	Theorem: (H1) proved by the Havelock identity
$d = 3$ (3+1D)	10	Equivalence: (H1) $\Leftrightarrow$ vacuum Einstein
$d \geq 4$	≥ 35	Equivalence: same argument, more Weyl components

In $d = 2$: the Havelock identity proves (H1) analytically—no assumption is needed. In $d \geq 3$: condition (H1) (mode-independence of C_1) is equivalent to the spatial Green’s function being isotropic, which is equivalent to maximal symmetry, which is equivalent to the vacuum Einstein equations with Λ . The soft Havelock decomposition (Proposition 10.11) shows that (H1) is approximately satisfied on any background, with residual $|R_m| < a|f_m|$ (verified for $N = 5, \dots, 20$ and $d = 3, 4, 5$): the violation of (H1) is controlled and measures the departure from Einstein geometry.

Supporting evidence: the spectral focusing theorem.

Theorem 10.17 (Spectral focusing). *For $N \geq 4$, the generalised Casimir ratio $R(\Delta) = h(m_{\text{crit}}, N, \Delta)/h(1, N, \Delta)$ is strictly increasing (Theorem 10.12). Consequently, $N_{\text{crit}}(\Delta)$ is non-increasing: once a polygon type becomes unstable, it remains unstable for all larger Δ .*

This is the polygon analogue of the Raychaudhuri focusing theorem:

Concept	General relativity	Polygon system
Flow parameter	Affine parameter λ	Conformal dim Δ
Expansion	$\theta = \nabla \cdot k$	$R(\Delta) = h(m, N, \Delta)/h(1, N, \Delta)$
Focusing	$d\theta/d\lambda \leq -\theta^2/2$	$dR/d\Delta > 0$ (monotone)
Trapped	$\theta < 0$ (trapped surface)	$\lambda_m < 0$ (unstable mode)
Energy condition	$R_{\mu\nu}k^\mu k^\nu \geq 0$	$C_1(\xi) > 0$ (spectral gap)

Supporting evidence: the Clausius relation. At each palindromic threshold (a local Rindler horizon), the Clausius relation $\delta Q = T \delta S$ gives

$$\frac{\delta Q}{T \delta S} = 4\pi \tanh(\rho^*) = 4\pi - O(e^{-2\rho^*}), \quad (10.28)$$

with quantum correction $e^{-2\rho^*} \sim e^{-N^2/6}$. The ratio 4π is *universal* across all thresholds $N = 7, 8, 9, \dots$ (verified to 10^{-10} for $N \leq 20$). In the Jacobson framework [Jacobson, 1995], universality of $\delta Q = T \delta S$ across local Rindler horizons is equivalent to the Einstein equations.

The arithmetic of the palindromic hierarchy is the interference between conformal invariance and geometric curvature—the residue that the renormalisation group cannot wash away.

References

- A. Adriani, A. Mura, G. Orton, et al. Clusters of cyclones encircling Jupiter’s poles. *Nature*, 555:216–219, 2018.
- A. Antuñano, T. del Río-Gaztelurrutia, A. Sánchez-Lavega, and R. Hueso. Dynamics of Saturn’s polar regions. *Journal of Geophysical Research: Planets*, 120:155–176, 2015.
- M. Baraban, G. Zikos, N. Bonesteel, and S. H. Simon. Numerical analysis of quasiholes of the Moore–Read wave function. *Physical Review Letters*, 103:076801, 2009. doi: 10.1103/PhysRevLett.103.076801.

- S. Boatto and H. E. Cabral. Nonlinear stability of a latitudinal ring of point-vortices on a nonrotating sphere. *SIAM Journal on Applied Mathematics*, 64:216–230, 2003.
- H. E. Cabral and D. S. Schmidt. Stability of relative equilibria in the problem of $n + 1$ vortices. *SIAM J. Math. Anal.*, 31:231–250, 1999.
- Steven Carlip. *Quantum Gravity in 2+1 Dimensions*. Cambridge University Press, 1998. doi: 10.1017/CBO9780511564192.
- Stanley Deser, Roman Jackiw, and Gerard 't Hooft. Three-dimensional Einstein gravity: dynamics of flat space. *Annals of Physics*, 152:220–235, 1984. doi: 10.1016/0003-4916(84)90085-X.
- Alexander L. Fetter. *Rotating trapped Bose–Einstein condensates*, volume 81. 2009. doi: 10.1103/RevModPhys.81.647.
- Guillaume Gauthier, Matthew T. Reeves, Xiaoquan Yu, Ashton S. Bradley, Mark A. Baker, Tyler A. Bell, Halina Rubinsztein-Dunlop, Matthew J. Davis, and Tyler W. Neely. Giant vortex clusters in a two-dimensional quantum fluid. *Science*, 364:1264–1267, 2019. doi: 10.1126/science.aat5718.
- Nimrod Gavriel and Yohai Kaspi. The number and location of Jupiter’s circumpolar cyclones explained by vorticity dynamics. *Nature Geoscience*, 14:559–563, 2021.
- Simone Giombi, Alexander Maloney, and Xi Yin. One-loop partition functions of 3D gravity. *Journal of High Energy Physics*, 2008(08):007, 2008. doi: 10.1088/1126-6708/2008/08/007.
- I. V. Grigorieva, W. Escoffier, J. Richardson, L. Y. Vinnikov, S. Dubonos, and V. Oboznov. Direct observation of vortex shells and magic numbers in mesoscopic superconducting disks. *Physical Review Letters*, 96:077005, 2006. doi: 10.1103/PhysRevLett.96.077005.
- T. H. Havelock. The stability of motion of rectilinear vortices in ring formation. *Philosophical Magazine*, 11:617–633, 1931.
- L. Iess, B. Militzer, Y. Kaspi, P. Nicholson, D. Durante, P. Racioppa, A. Anabtawi, E. Galanti, W. Hubbard, M. J. Mariani, P. Tortora, S. Wahl, and M. Zannoni. Measurement and implications of Saturn’s gravity field and ring mass. *Science*, 364:eaat2965, 2019. doi: 10.1126/science.aat2965.
- Ted Jacobson. Thermodynamics of spacetime: the Einstein equation of state. *Physical Review Letters*, 75:1260–1263, 1995. doi: 10.1103/PhysRevLett.75.1260.
- Y. Kaspi, E. Galanti, W. B. Hubbard, D. J. Stevenson, S. J. Bolton, L. Iess, T. Guillot, A. Blöcker, J. Levin, P. Lauritzen, W. M. Folkner, R. Helled, A. P. Ingersoll, R. L. Kirk, M. J. Mariani, P. Tortora, S. Wahl, and M. Zannoni. Jupiter’s atmospheric jet streams extend thousands of kilometres deep. *Nature*, 555:223–226, 2018. doi: 10.1038/nature25793.
- Leonid G. Kurakin and Victor I. Yudovich. The nonlinear stability of stationary rotation of a regular vortex polygon. *Chaos*, 12:574–595, 2002.

- Riccardo Rattazzi, Vyacheslav S. Rychkov, Erik Tonni, and Alessandro Vichi. Bounding scalar operator dimensions in 4D CFT. *Journal of High Energy Physics*, 2008(12):031, 2008. doi: 10.1088/1126-6708/2008/12/031.
- A. Sánchez-Lavega, T. del Río-Gaztelurrutia, R. Hueso, et al. The long-term steady motion of Saturn’s hexagon and the stability of its enclosed jet stream under seasonal changes. *Geophysical Research Letters*, 41:1425–1431, 2014.
- Dieter S. Schmidt. The stability of the Thomson heptagon. *Regul. Chaotic Dyn.*, 9: 519–528, 2004.
- Alexander Strohmaier and Ville Uski. An algorithm for the computation of eigenvalues, spectral zeta functions and zeta-determinants on hyperbolic surfaces. *Communications in Mathematical Physics*, 317:827–869, 2013. doi: 10.1007/s00220-012-1557-1. arXiv:1110.2150; ancillary data includes 1000 Bolza eigenvalues at 48-digit precision.
- Gerard ’t Hooft. Canonical quantization of gravitating point particles in 2+1 dimensions. *Classical and Quantum Gravity*, 10:1653–1664, 1993. doi: 10.1088/0264-9381/10/8/022.
- J. J. Thomson. A treatise on the motion of vortex rings. *Macmillan, London*, 1883.