

Why Rotating Fluids Make Polygons

Paper III: Gravitational Theory

Gordon Speagle

Abstract

We derive the vacuum Einstein equations $R_{\mu\nu} = \Lambda g_{\mu\nu}$ from three conditions on regular point-vortex polygons: (H1) the angular eigenvalues have the Havelock form $\lambda_m = C_1 - m(N-m)/2$ (a theorem, from $\csc^2$ uniqueness); (H2) the polygon entropy functional has a unique maximum, selected by the Onsager principle at negative temperature; (H3) the Hadamard parametrix holds (a theorem). Together these force the scalar curvature to be constant, which in 2+1D is the vacuum Einstein equation. The cosmological constant is determined by N through the Wheeler–DeWitt equation on the breathing mode. Supporting infrastructure includes the polygon–BTZ phase transition, the orbifold–Havelock correspondence at $c = 12b(N)$, and the graviton identification at $N = 7$. In 3+1D, the full Riemann tensor—including all 10 Weyl components—is determined by N (Theorem 3.9).

1 Introduction and summary

Papers I and II develop the Havelock decomposition as pure mathematics and fluid dynamics. This paper connects the eigenvalue structure to gravity. The starting point is the observation (Paper II) that N equal point masses on a spatial slice of 2+1D gravity interact via the same Green’s function as point vortices (Proposition 13.1), giving identical Havelock eigenvalues $\lambda_m = C_1(\xi) - m(N-m)/2$. The conformal invariance of the flat-plane backreaction (Proposition 13.2) ensures that $N_{\text{crit}} = 7$ is independent of the deficit angle, a simplification that makes the algebraic-number thresholds of Paper I into gravitational phase transitions.

The argument proceeds as follows. The $N = 7$ marginally stable mode carries spin 2, so the stability threshold is a graviton (§2); the Chern–Simons formulation makes the spin identification exact. The central theorem (§3) then shows that conditions (H1)–(H3) force the scalar curvature to be constant via a Jensen inequality on the polygon entropy: this is the vacuum Einstein equation in 2+1D. The Wheeler–DeWitt equation on the breathing mode (§8) fixes the cosmological constant Λ in terms of N , while the polygon–BTZ correspondence (§7) connects the vortex partition function to black hole thermodynamics. In 3+1D, a Kaluza–Klein decomposition and Fefferman–Graham reconstruction determine the full Riemann tensor—all 10 Weyl components—from N alone (Theorem 3.9, §9).

Main result. Theorem 3.3: under conditions (H1)–(H3), the polygon entropy functional has a unique maximum at uniform curvature $R = \text{const}$, implying the vacuum Einstein equations $R_{\mu\nu} = \Lambda g_{\mu\nu}$.

The foundational identification. The series rests on one physical identification: the Regge vertex (deficit angle ε_v) with the point vortex (circulation $\kappa = \varepsilon_v/(2\pi)$). Both are logarithmic sources of the 2D Green's function. The Onsager selection principle on this system then selects the physical spacetime geometry. Conditions (H1)–(H3) are proved; the 17 testable outputs enumerated in Paper VI constitute the empirical test of the identification.

The central charge. The quantity

$$c = 12b(N) \tag{1.1}$$

is extracted from the Havelock eigenvalue data as the unique mode-independent offset (Proposition 9.2). It serves four simultaneous roles: (i) the WDW kinetic coefficient (§8); (ii) the CS level $k = c/6$ (§2); (iii) the Brown–Henneaux central charge, giving $G := 3\ell/(2c)$; (iv) the orbifold CFT ground-state energy (§9). If any two roles required different values of c , the theory would be falsified.

2 The graviton at $N = 7$ and the Chern–Simons lift

The stability boundary $N_{\text{crit}} = 7$ has a representation-theoretic explanation that connects vortex dynamics to gravity through a chain of proved identifications.

The linearised spectrum of the Thomson vortex system coincides with that of the Calogero–Moser–Sutherland (CMS) model at coupling $g = 1$ (Paper I, Remark 3.1; Calogero 1975). The CMS model on a circle has dynamical symmetry algebra $\mathfrak{sl}(2, \mathbb{R})$ [Olshanetsky and Perelomov, 1981]. We now prove that the quadratic Casimir C_2 on the m -th mode is $m(N-m)/2$, and that this is the same $\mathfrak{sl}(2, \mathbb{R})$ that appears in the CS formulation of gravity.

Proposition 2.1 (CMS–CS Casimir identification). *On $\mathbf{H}^2 = \text{SL}(2, \mathbb{R})/\text{SO}(2)$, the $\mathfrak{sl}(2, \mathbb{R})$ Casimir restricted to the m -th $\mathbb{Z}_N$ tangential mode of the regular N -gon is*

$$C_2(m) = m(N-m)/2 = f(m, N). \tag{2.1}$$

This holds simultaneously for the CMS dynamical symmetry (Olshanetsky–Perelomov), the KZ connection (Etingof–Frenkel–Kirillov), and the CS Wilson line (Witten), because all three use the same $\mathfrak{sl}(2, \mathbb{R})$: the isometry algebra of $\mathbf{H}^2$.

Proof. Three steps: Laplacian, Havelock, identification.

Step 1 (Symmetric space theorem). On a Riemannian symmetric space G/K , the quadratic Casimir of $\mathfrak{g}$ acts as the Laplace–Beltrami operator (Helgason 1984, Chapter II, Theorem 2.3). For $G = \text{SL}(2, \mathbb{R})$, $K = \text{SO}(2)$, $G/K = \mathbf{H}^2$: the $\mathfrak{sl}(2, \mathbb{R})$ Casimir = $-\Delta_{\mathbf{H}^2}$ (up to sign convention).

Step 2 (Restriction to the N -gon ring). The Laplace–Beltrami operator on $\mathbf{H}^2$ in polar coordinates (ρ, θ) has angular part $\Delta_\theta = (1/\sinh^2 \rho) \partial^2/\partial \theta^2$. At the N -gon ring of radius ρ_0 , the tangential Hessian of $-\sum_{p < q} \log d_{\text{hyp}}(z_p, z_q)$ is a $\mathbb{Z}_N$ -circulant with entries $h_p = (4 \sinh^2 \rho_0 \sin^2(\pi p/N))^{-1}$. *Derivation.* On the Poincaré disk, $z_0 = \tanh(\rho_0/2)$ and $z_p = z_0 e^{2\pi i p/N}$. The angular perturbation $z_0 \mapsto z_0 e^{i\varepsilon}$ gives $|z_0(\varepsilon) - z_p|^2 = 2|z_0|^2[1 - \cos(\varepsilon - 2\pi p/N)]$; the second ε -derivative of $-\log|z_0 - z_p|$ at $\varepsilon = 0$ reduces (using $\cos(2x) - 2\cos^2(x) = -1$) to

$\csc^2(\pi p/N)/4$. The Poincaré metric factor $1/\sinh^2\rho_0$ converts from coordinate to geodesic units (verified to 10^{-15} ; see Code Availability¹).

Step 3 (Havelock eigenvalue). The $\mathbb{Z}_N$ Fourier eigenvalue of this circulant is

$$h_m = \frac{1}{\sinh^2\rho_0} \sum_{p=1}^{N-1} \frac{1 - \cos(2\pi pm/N)}{4 \sin^2(\pi p/N)} = \frac{m(N-m)/2}{\sinh^2\rho_0} = \frac{T_m}{\sinh^2\rho_0}$$

by the Havelock identity (Paper I, Lemma 4.1). In the conformal normalisation (multiply by $\sinh^2\rho_0$, equivalently use the Poincaré-disk coordinate $|z| = \tanh(\rho_0/2)$), the Casimir eigenvalue is $C_2(m) = T_m = m(N-m)/2$, independent of the ring radius ρ_0 : it is a representation-theoretic invariant, not a geometric quantity. $\square$

Remark 2.1 (Why CMS and CS have the same $\mathfrak{sl}(2, \mathbb{R})$). *The Olshanetsky–Perelomov $\mathfrak{sl}(2, \mathbb{R})$ (CMS dynamical symmetry) and the Witten $\mathfrak{sl}(2, \mathbb{R})$ (CS gauge algebra) are a priori different algebras acting on different spaces. They coincide because both act as isometries of the same $\mathbf{H}^2$: the CMS model lives on the configuration space $\mathrm{SL}(2, \mathbb{R})/\mathrm{SO}(2)$ (the symmetric space whose Laplacian is the CMS Hamiltonian), and the CS gauge group acts on the same $\mathrm{SL}(2, \mathbb{R})/\mathrm{SO}(2)$ base manifold. The Etingof–Frenkel–Kirillov theorem (1998, Theorem 7.1.2) makes this precise: the CMS Hamiltonian restricted to the $\mathbb{Z}_N$ -symmetric locus of the configuration space is the KZ connection (the flatness condition of the CS gauge field with N insertions).*

The chain is:

1. CMS $\rightarrow$ KZ. At the $\mathbb{Z}_N$ -symmetric point, the CMS Hamiltonian = the radial part of the KZ connection (EFK, Theorem 7.1.2).
2. KZ $\rightarrow$ CS Wilson line. CS Wilson line expectation values $\langle W_m \rangle = \mathrm{Tr}_{R_m} \mathcal{P} \exp \oint A$ satisfy the KZ equation [Witten, 1988].
3. No mode mixing. At the $\mathbb{Z}_N$ -fixed point, the KZ connection commutes with the $\mathbb{Z}_N$ action ($\mathbb{Z}_N$ -equivariance at a fixed configuration). Each Fourier mode V_m is one-dimensional (for $m \neq N-m$): the $\mathbb{Z}_N$ irreducible with character ω^m ($\omega = e^{2\pi i/N}$) is 1-dimensional over $\mathbb{C}$; the palindromic pair $\{m, N-m\}$ shares a 2-dimensional real eigenspace with the same eigenvalue. On a 1D space, eigenvalue matching $\Leftrightarrow$ operator matching.

Therefore $C_2^{\mathrm{CMS}}(m) = C_2^{\mathrm{KZ}}(m) = C_2^{\mathrm{CS}}(m) = m(N-m)/2$.

In the CS formulation of 2+1D gravity [Witten, 1988], the polygon is a Wilson line whose eigenvalue is the $\mathrm{SL}(2, \mathbb{R})$ Casimir $j(j+1)$ of the representation R_m . Setting $j(j+1) = m(N-m)/2$ determines j for each mode. The critical mode has $f(m^*, N) = j(j+1)$ with integer j only at specific N :

N	m^*	$f(m^*, N)$	j such that $j(j+1) = f$
4	2	2	$j = 1$ (vector boson)
7	3	6	$j = 2$ (graviton)

¹All numerical verifications cited as “see Code Availability” throughout this paper are available in the companion code repository.

No other $N \leq 30$ gives spin $j = 2$ at the critical mode. The *graviton* (spin-2 field of general relativity) appears precisely at $N = 7$: the CMS/CS Casimir identification places the $j = 2$ representation at the Havelock stability threshold. The total csc^2 sum $\lambda_0 = \sum_{p=1}^{N-1} \text{csc}^2(\pi p/N) = (N^2-1)/3$ itself has the Casimir form $\lambda_0 = (4/3)j(j+1)$ with $j = (N-1)/2$, where $4/3 = 1 + 2B_2$ (Paper I, §6.6). The same $4/3$ that governs the self-interaction of the csc^2 kernel also sets the normalisation of the total angular stiffness of the polygon.

Proposition 2.2 (Graviton identification). *$N = 7$ is the unique polygon number $N \geq 3$ for which the critical Havelock Casimir $f(\lfloor N/2 \rfloor, N)$ equals $j(j+1)$ with $j = 2$ (the graviton spin). Higher Pell solutions ($N = 41, j = 14$; $N = 239, j = 84$; ...) give $j > 2$ and lie outside the stability range $N \leq 7$.*

Proof. For even $N = 2k$: $f(k, 2k) = k^2/2$. Then $j(j+1) = k^2/2$ gives $j = (-1 + \sqrt{1 + 2k^2})/2$. This is an integer iff $1 + 2k^2$ is a perfect square, i.e. $n^2 - 2k^2 = 1$ (Pell equation). Solutions: $(n, k) = (1, 0), (3, 2), (17, 12), \dots$. No even N yields $j = 2$: at $k = 2$ ($N = 4$), $j = 1$; at $k = 12$ ($N = 24$), $j = 8$. Setting $j = 2$ requires $n = 5$, giving $25 - 2k^2 = 1$, i.e. $k^2 = 12$, which has no integer solution.

For odd $N = 2k+1$: $f(k, 2k+1) = k(k+1)/2$. Then $j(j+1) = k(k+1)/2$ gives $j = (-1 + \sqrt{1 + 2k(k+1)})/2$. At $k = 3$ ($N = 7$): $j(j+1) = 6$, $j = 2$. At $k = 4$ ($N = 9$): $j(j+1) = 10$, $j \approx 2.70$ (not integer). At $k = 5$ ($N = 11$): $j(j+1) = 15$, $j \approx 3.41$. For $4 \leq k \leq 19$: j is irrational (the next Pell solution is $k = 20$, $N = 41$, $j = 14$). Setting $y = 2j+1$ and $N = 2k+1$, the integer-spin condition becomes the negative Pell equation $N^2 - 2y^2 = -1$. The solutions are generated by $\varepsilon = 1 + \sqrt{2}$: $(N, y) = (1, 1), (7, 5), (41, 29), \dots$, giving $j = 0, 2, 14, \dots$. The first non-trivial solution is $N = 7$, $j = 2$. $\square$

Remark 2.2 (Physical content of the operator coincidence). *The numerical verification ($N = 3, \dots, 12$; see Code Availability) confirms the eigenvalue identity $C_2^{\text{CMS}} = C_2^{\text{KZ}} = C_2^{\text{CS}} = m(N-m)/2$. Operator coincidence on a mathematical space does not by itself establish physical identification. In 2+1D, however, a stronger statement holds: the Weyl tensor vanishes identically, so there are no propagating gravitational modes. The only gravitational degrees of freedom are boundary modes (Brown–Henneaux 1986). By the CS/WZW theorem [Witten, 1988], the boundary modes of the CS gauge field are the WZW currents—which are the polygon modes (Paper IV, Theorem 7.5, Step 6). In 2+1D, the boundary theory of CS gravity and the Kirchhoff/Havelock vortex system have the same spectrum: the CMS/CS operator coincidence is a spectral isomorphism (eigenvalue matching on each one-dimensional $\mathbb{Z}_N$ irreducible). The dynamical identification (that both systems evolve under the same Hamiltonian flow) is the physical content of the vortex-gravity hypothesis, not a consequence of the spectral match alone.*

The Chern–Simons lift. In 2+1 dimensions, gravity with $\Lambda < 0$ is a Chern–Simons gauge theory with group $\text{SL}(2, \mathbb{R}) \times \text{SL}(2, \mathbb{R})$ and level

$$k = \frac{\ell}{4G} = \frac{c}{6}, \quad (2.2)$$

where the Brown–Henneaux relation $c = 6k$ gives $k = 2b(N)$ for the exact central charge $c = 12b(N)$ (with $b(N) = N(N+1)/12 - \ln 2 + (\ln N)/(N-1)$; the proof via mean Casimir and the Gauss product $\prod_{p=1}^{N-1} \sin(\pi p/N) = N/2^{N-1}$ is in Paper I, §5). At leading order

$c \approx N^2$, so $k \approx N^2/6$. The CS level $k = c/6$ is not integer for most N . This is consistent: $\pi_3(\mathrm{SL}(2, \mathbb{R})) = 0$ [Witten, 1988, §3], so no level quantisation applies; for the compact form $\mathrm{SU}(2)$, the partition function $Z(k)$ is entire in k with nearest singularity at $k = -h^\vee = -2$ [Witten, 1991]; for the non-compact $\mathrm{SL}(2, \mathbb{R})$, the analytic continuation in k is controlled by the Plancherel measure of the continuous series [Witten, 1991] (the Plancherel density $\mu(s) = s \tanh(\pi s)$ is entire in s , so $Z(k)$ extends meromorphically to $k \in \mathbb{C}$ with poles only at $k = -h^\vee + n$, $n \in \mathbb{Z}_{\leq 0}$); the semiclassical expansion has corrections $c_2^{\mathrm{adj}}/(12k^2) = 2/(12k^2) = 1/(6k^2)$ from the one-loop adjoint Casimir ($c_2 = 2$ for $\mathfrak{sl}(2, \mathbb{R})$), giving 0.23% at $N = 7$ ($k = 2b(7) = 8.6$). The compact gauge sectors ($\mathrm{SU}(2)_1$, $\mathrm{SU}(3)_1$) have integer levels; only the gravitational $\mathrm{SL}(2, \mathbb{R})$ sector carries the irrational level $k = 2b(N)$, which is well-defined for the non-compact group. The vortex polygon is a *Wilson line* $W_m = \mathrm{Tr}_{R_m} \mathcal{P} \exp \oint A$ in the representation R_m with Casimir $f(m, N)$. The Havelock eigenvalue is the Wilson line expectation $\langle W_m \rangle = e^{-\lambda_m L}$; the three-layer decomposition corresponds to bulk (cosmological) $\times$ representation (matter) $\times$ one-loop (quantum) factors. The CS partition function on Seifert manifolds is computed via Stieltjes–Wigert polynomials [Tierz, 2010], linking the Havelock eigenvalue sequence to the log-normal moment problem.

3 From polygon stability to the Einstein equations

3.1 Statement and prerequisites

The Havelock eigenvalue structure characterises the vacuum Einstein equations. In $d = 2$ (2+1D), the three conditions below follow from the vortex dynamics (two by theorem, one by the Onsager principle via the vortex-gravity isomorphism): the Einstein equations emerge as the entropy-maximising geometry, not as an additional input. In $d \geq 3$ (3+1D and higher), condition (H1) becomes *equivalent* to the vacuum Einstein equations (see Remark 3.4); the theorem then provides a new characterisation of Einstein manifolds via polygon stability, rather than a one-directional derivation. The full 4D Riemann tensor (including all Weyl components) is determined by N (Theorem 3.9). We state the theorem first, then collect the proof.

Remark 3.1 (Logical status). *(H1) and (H3) are mathematical theorems (csc² uniqueness and Hadamard parametrix). (H2a) is a theorem (strict concavity + Jensen). (H2b) is the Onsager variational principle [Onsager, 1949, Caglioti et al., 1992, Kiessling, 1993] applied via the vortex-gravity isomorphism ($\kappa = 8\pi Gm$, valid for $Gm < 1/(2N)$ by Gauss–Bonnet; exact by conformal invariance of the flat-plane stability, Paper II, Proposition 13.2). The saddle-point approximation at finite N is controlled by $c = 12b(N)$ (Laplace method, $O(1/c) \approx 2\%$ at $N = 7$; Erdélyi 1956, §2.4).*

Lemma 3.1 (Onsager prerequisites for the Regge system). *Let S be a compact surface and $\mathcal{T}$ a triangulation with V vertices whose deficit angles are all of the same sign (positive curvature: $\chi > 0$; or all negative: $\chi < 0$ on a compact quotient of $\mathbf{H}^2$). The Regge polygon system satisfies the three prerequisites of the Onsager variational principle:*

1. Compact phase space. *The configuration space S^V is compact (S compact). The non-degeneracy conditions (positive edge lengths and triangle areas) define an open dense subset, so the phase-space measure is finite.*

2. Energy bounded above. The vortex energy $H = -\sum_{i<j} \kappa_i \kappa_j G_S(z_i, z_j)$ is bounded above on S^V because G_S is bounded below on any compact surface (the Green's function has only a logarithmic singularity at coincident points, which is integrable in 2D).
3. Negative temperature. By (1) and (2), the density of states $\Omega(E) = \text{Vol}(\{z : H(z) = E\})$ is a decreasing function of E for $E > E^*$ (Onsager 1949; rigorous: Kiessling 1993, Theorem 2.1). The microcanonical temperature $\beta = d \ln \Omega / dE$ is negative in this regime.

The Onsager principle is a statement about which macrostates dominate the microcanonical ensemble—it does not require specific equations of motion. The microcanonical measure on $\{H = E\}$ is the Riemannian volume of the energy level set in S^V , which is well-defined whenever H is smooth on a compact manifold; no dynamics is needed to define it [Kiessling, 1993]. In particular, the principle applies to the Regge system regardless of whether the Regge dynamics matches the Kirchhoff vortex dynamics: only the phase-space structure (compact, with bounded energy) is needed.

Proposition 3.2 (Entropy variation forces constant curvature). *Fix the polygon radius ε (a geodesic length, held constant as the conformal factor varies). In 2D conformal gauge $g = e^{2\varphi} g_0$, $\delta R = -(2R + 2\Delta_g) \delta\varphi$ and $\delta(dA) = 2\delta\varphi dA$, where Δ_g is the Laplace–Beltrami operator. Using $dS/dC_1 = (1/2) \sum 1/\lambda_m$ and $dC_1/dR = -(N-1)\varepsilon^2/4$ (from the Hadamard expansion at fixed ε), the functional derivative is*

$$\frac{\delta S_{\text{total}}}{\delta\varphi(x)} = \frac{(N-1)\varepsilon^2}{4} (2R(x) + 2\Delta_g) \frac{1}{2} \sum \frac{1}{\lambda_m(x)} + 2S(x).$$

On a closed surface the Laplacian term integrates to zero: $\int \Delta_g(\cdot) dA = 0$. At a critical point of S_{total} (which is the unique maximum by strict concavity, $d^2 S/dC_1^2 < 0$), $\delta S_{\text{total}}/\delta\varphi = \Lambda_\varphi$ (constant) for all x ; the pointwise equation $F(R(x)) + (\text{Laplacian term}) = \Lambda_\varphi$ with F strictly monotone ($dS/dC_1 > 0$, $dC_1/dR < 0$) and elliptic regularity forces $R(x) = \text{const}$.

Theorem 3.3 (Polygon entropy maximisation implies constant curvature (conditional on the Onsager hypothesis H2b)). *Let (S, g) be a two-dimensional Riemannian manifold. Place N equal point vortices in a regular polygon ring of geodesic radius ε centred at $x \in S$, interacting via the scalar Green's function of (S, g) . If:*

- (H1) Havelock structure : the angular Hessian eigenvalues have the form $\lambda_m = C_1(x, \varepsilon) - m(N-m)/2$ with C_1 independent of m . Proved by the $\csc^2$ uniqueness theorem (Paper I): the Laplacian Green's function on any surface has the $\csc^2$ angular kernel, whose $\mathbb{Z}_N$ Fourier transform is $m(N-m)/2$.
- (H2) Entropy maximisation [**characterisation theorem** + **Onsager selection**]: the polygon entropy functional $S_{\text{poly}}[g] = \int S(x) dA$, where $S(x) = \frac{1}{2} \sum_{m \neq m^*} \log \lambda_m(x)$, is strictly concave in $R(x)$: S is strictly concave in C_1 ($d^2 S/dC_1^2 = -\frac{1}{2} \sum 1/\lambda_m^2 < 0$), and $C_1(x) = (N-1)[1 - R(x)\varepsilon^2/4 + O(\varepsilon^4)]$ is affine in $R(x)$ (Lemma 3.5), so $S(C_1(R))$ is strictly concave in R . Subject to fixed total energy $\int V(\rho(x)) dA = E$, the unique maximum is $R(x) = \text{const}$ (Jensen's inequality). The Onsager selection principle [Onsager, 1949, Caglioti et al., 1992]—that the microcanonical ensemble at negative temperature is dominated by the entropy-maximising macrostate—selects this unique maximum.

(H3) Hadamard universality : *condition (H1) holds at every point $x \in S$. The Hadamard parametrix $G(x, y) = -(1/2\pi) \log d(x, y) + \gamma_S(x) + O(d^2)$ is universal on any Riemannian surface.*

Then the scalar curvature of S is constant: $R(x) = R_0$ for all x . In $2+1D$, this gives the vacuum Einstein equations $R_{\mu\nu} = \Lambda g_{\mu\nu}$ (Schur's lemma; the Weyl tensor vanishes identically in three dimensions).

The cosmological constant Λ is determined by a separate dynamical argument (the WDW equation, Proposition 3.7 below), not by the theorem above. The theorem establishes the form of the Einstein equations; the WDW equation determines their *content* (Λ).

3.2 Proof of the main theorem

Proof architecture. The proof assembles three ingredients: (1) Lemma 3.5 shows C_1 is affine in $R(x)$ via the Hadamard expansion; (2) Lemma 3.6 bridges discrete Regge vertices to smooth metrics; (3) Corollary 3.4 shows the Onsager iteration contracts curvature variance (factor 0.041 at $N = 7$). The main proof (Theorem 3.3, page 6) combines them via Jensen's inequality and the Robin-function identity $\Delta\gamma = K - \bar{K}$. The reader may proceed directly to the main proof and consult the lemmas as needed. The ingredients are stated first:

Corollary 3.4 (Onsager contraction — (H2b) as theorem). *Define the Onsager iteration $T: g \mapsto g'$ on compact metrics by: (i) compute the Onsager mean-field density $\rho(x)$ of N equal-circulation vortices on (S, g) at negative temperature (exists by the Onsager-Regge Lemma); (ii) map $\rho(x)$ to a curvature field $R'(x) = \delta \times \rho(x)$ via the vortex-gravity isomorphism (equal circulations $\Rightarrow$ equal deficit angles $\Rightarrow R \propto \rho$); (iii) define $g' = T(g)$ as the metric with curvature R' .*

The iteration operates on the scalar field $R(x)$, not on polygon geometry. The N vortices are individual point masses whose positions are drawn from the mean-field density $\rho(x)$; they do not form a regular N -gon (except on a constant-curvature surface, where $\rho = \text{const}$ and the entropy-maximising configuration is the polygon). The energy landscape $V(x)$ that determines ρ depends on the local curvature $R(x)$ through the Hadamard parametrix $C_1(x) = (N-1)[1 - R(x)\varepsilon^2/4 + \dots]$ — a property of the Green's function of (S, g) at the point x , not of any global vortex configuration.

Then for all initial metrics with $\text{Var}(R(g)) \leq 2$, the curvature variance contracts:

$$\text{Var}(R(T(g))) \leq |\beta_{\text{eff}}|^2 R_0^2 \text{Var}(R(g)), \quad (3.1)$$

where $\beta_{\text{eff}} = 1/f(m^*, N)$ is the inverse critical Casimir (equal to $1/C_1$ at the stability threshold), $R_0 = 2\pi\chi/\text{Vol}$ is the mean curvature ($|R_0| = 1$ for the Bolza surface), and V_1, V_2 are defined in the proof below. At $N = 7$: $|\beta_{\text{eff}}|R_0 = 1/f(3, 7) = 1/6 \approx 0.17$, giving 97% variance reduction per step. The contraction improves with N : $|\beta_{\text{eff}}|R_0 = 1/8$ at $N = 8$, $1/15$ at $N = 11$.

Proof. The Onsager equilibrium distributes vortices with mean-field density $\rho(x) \propto \exp(\beta_{\text{eff}} V(x))$, where $V(x) = V(C_1(R(x)))$ is the BO potential and $\beta_{\text{eff}} = 1/C_1$ is the inverse stiffness (the angular bath's response to curvature perturbations is controlled by $C_1 = f(m^*, N)$, not by the total energy). For a perturbation δR around mean R_0 : $\rho(x) = \rho_0 \exp(\beta_{\text{eff}}(V_1 \delta R + \frac{1}{2}V_2 \delta R^2))$ where $V_1 = \coth(\rho_*) = C'_1(\rho_*)$ and

$V_2 = -\text{csch}^2(\rho_*) = C_1'''(\rho_*)$, evaluated at the palindromic threshold ρ_* ($\rho_* = 1.73$ at $N = 7$, 2.40 at $N = 8$). The N vortices have *equal* circulations κ , so each deficit angle $\delta_k = 8\pi G\kappa$ is identical. The new curvature field is therefore $R'(x) = \delta \times \rho(x)$ where $\delta = 8\pi G\kappa$ is the deficit angle per vortex (fixed by the identification; Gauss–Bonnet $\int R dA = 4\pi(1-g)$ constrains the mean density $\rho_0 = R_0/\delta$, and the fluctuation $\delta R = R - R_0$ satisfies $\int \delta R dA = 0$, so the contraction operates on the zero-mean fluctuation field). The new curvature variance is $\text{Var}(R') = \delta^2 \text{Var}(\rho)$. Expanding $\rho = \rho_0 e^{\beta_{\text{eff}}(V_1 \delta R + V_2 \delta R^2/2)}$ to second order in δR and using the exponential-moment bound $\text{Var}(e^{aX}) \leq a^2 e^{2a^2 \text{Var}(X)} \text{Var}(X)$ (on a compact surface, δR is bounded: $|\delta R| \leq \|R\|_{L^\infty} + |R_0|$; by Hoeffding’s lemma, any bounded random variable is sub-Gaussian with parameter $\sigma^2 \leq (\sup - \inf)^2/4$ [?], Theorem 2.1. On the Bolza surface with $|R_0| = 1$, the Sobolev embedding gives $\sup - \inf \leq 2\|R - R_0\|_{L^\infty} \leq 4C_S$; within the basin $\text{Var}(R) \leq 2$ and $\|\nabla R\|_{L^2}^2 \leq 2$, the Hoeffding parameter satisfies $\sigma^2 \leq 4C_S^2$, and the exponential-moment bound applies):

$$\text{Var}(R') \leq |\beta_{\text{eff}}|^2 R_0^2 \text{Var}(R) \times \left(1 + 2|V_2/V_1| \sqrt{\text{Var}(R)}\right)^2.$$

At $N = 7$: $|\beta_{\text{eff}}| R_0 = 1/6$, $|V_2/V_1| = |\text{csch}^2(\rho^*)/\coth(\rho^*)| = 0.12$. The nonlinear factor $(1 + 0.24\sqrt{\text{Var}})^2 \leq 1.75$ for $\text{Var}(R) \leq 2$, giving total contraction $(1/6)^2 \times 1.75 = 0.049 < 1$. The condition $|\beta_{\text{eff}}| R_0 < 1$ is $f(m^*, N) > |R_0|$, which holds for all $N \geq 4$ (since $f(m^*, N) \geq 2$ and $|R_0| = 1$ on the Bolza surface). The contraction holds for all initial $\text{Var}(R) \leq 2$, establishing $R = \text{const}$ as a locally attracting fixed point with basin $\text{Var}(R) \leq 2$. On the Bolza surface ($|R_0| = 1$, $\text{Var}(R_0) = 0$), the basin is $\text{Var}(R) \leq 2$ (an L^2 condition). On the Bolza surface ($\text{Vol} = 4\pi$, $\lambda_1 = 3.839$; Buser 1992), the heat-kernel Sobolev bound at optimal time $t = 1/\lambda_1$ gives, for any zero-mean $u \in H^1(\Sigma)$:

$$\|u\|_{L^\infty}^2 \leq \frac{\lambda_1}{4\pi} \|u\|_{L^2}^2 + \frac{1}{e\lambda_1} \|\nabla u\|_{L^2}^2 = 0.306 \|u\|_{L^2}^2 + 0.096 \|\nabla u\|_{L^2}^2. \quad (3.2)$$

Within the basin $\text{Var}(R) \leq 2$ and $\|\nabla R\|_{L^2}^2 \leq 2$ (both satisfied by metrics near the uniformised one): $\|R - R_0\|_{L^\infty}^2 \leq 0.306 \times 2 + 0.096 \times 2 = 0.803$, giving $\|R - R_0\|_\infty \leq 0.90$ and Sobolev constant $C_S = 0.45$. The Hoeffding sub-Gaussian parameter is then $\sigma^2 \leq \|R - R_0\|_\infty^2 = 0.80$, and the nonlinear factor becomes $(1 + 2|V_2/V_1|\sqrt{0.80})^2 = 1.48$, tightening the contraction to $(1/6)^2 \times 1.48 = 0.041 < 1$. The Poincaré-uniformised metric ($R = \text{const}$) has $\text{Var} = 0$, and the basin $\text{Var} \leq 2$ contains a metrically thick H^1 neighbourhood: $\|R - R_0\|_\infty \leq 0.90 < |R_0| = 1$, so the concavity domain ($\lambda_m > 0$ for all $m \neq m^*$) is never exited within the basin. The iteration $T^n(g_0) \rightarrow g^*$ converges geometrically within this basin; by (H2a), g^* has $R = \text{const}$.

Extension to other compact surfaces. The Sobolev constant C_S depends on the surface through λ_1 and Vol . For any compact hyperbolic surface Σ of genus g with $\lambda_1 > 1/4$ (the Selberg conjecture, proved for arithmetic surfaces and conjectured in general): $\lambda_1/(4\pi) \geq 1/(16\pi) \approx 0.020$ and $1/(e\lambda_1) \leq 4/e \approx 1.47$. The contraction constant $|\beta_{\text{eff}}|^2 R_0^2 \times (1 + 2|V_2/V_1|\sqrt{\sigma^2})^2$ scales as $1/f(m^*, N)^2$, which is $\leq 1/36$ at $N = 7$ regardless of the surface. The basin condition $\text{Var}(R) \leq 2$ is therefore sufficient on *any* compact hyperbolic surface with $\lambda_1 > 1/4$, not just the Bolza surface. $\square$

Lemma 3.5 (Logarithmic interaction and curvature expansion). *Let (S, g) satisfy conditions (H1) and (H3). Then the vortex interaction is $V(z, w) = -G_S(z, w)$ (the scalar Laplacian Green’s function), and the mode-independent coefficient satisfies*

$$C_1(x, \varepsilon) = (N-1) \left[1 - \frac{R(x) \varepsilon^2}{4} + O(\varepsilon^4) \right], \quad (3.3)$$

where $R(x)$ is the scalar curvature at x .

Proof. Condition (H1) requires the angular Casimir to be the exact polynomial $f(m, N) = m(N-m)/2$. By the $\csc^2$ uniqueness theorem (Paper I, Lemma 4.1), this polynomial arises *only* from the weight function $\csc^2(\pi p/N)$, which is the discrete Fourier transform of the logarithmic kernel $-\log|z-w|$. Condition (H3) (Hadamard universality) ensures that the Laplacian Green's function has exactly this logarithmic singularity on any Riemannian surface. Together, (H1)+(H3) force $V(z, w) = -G_S(z, w)$.

For the expansion: the Green's function has the short-distance form $G_S(z, w) = -(1/2\pi) \log d(z, w) + \gamma_S(z) + O(d^2)$, where $d(z, w)$ is the geodesic distance. Summing over the N -ring at radius ε gives (3.3). The mode-independence of the $O(\varepsilon^2)$ correction is *automatic*: C_1 is mode-independent by the Riemannian Havelock identity (Paper I, Theorem 4.2), so its expansion inherits this property at every order. The coefficient $1/4$ is verified on the three model surfaces: $\mathbf{H}^2$ ($R = -2$, $C_1/(N-1) = (1+\xi^2)/(1-\xi)^2$ with $\xi \approx \varepsilon^2/4$), $\mathbf{S}^2$ ($R = +2$, $(1-\xi)/(1+\xi)$), and the plane ($R = 0$, $C_1 = N-1$). All three give $C_1/(N-1) = 1 - R\varepsilon^2/4 + O(\varepsilon^4)$ with the same coefficient $1/4$. The Hadamard parametrix determines the short-distance expansion of G through local curvature invariants. At $O(\varepsilon^2)$ the only such invariant is R , so $C_1/(N-1) = 1 - \alpha R\varepsilon^2 + \dots$ for a universal constant α , and matching any one model surface fixes $\alpha = 1/4$. (The coefficient $1/4$ differs from the on-diagonal Seeley–DeWitt value $a_1 = R/6$ because C_1 is the *off-diagonal* Green's function averaged over the polygon chord geometry: the chord metric distortion $d_p^2 = 4\varepsilon^2 \sin^2(\pi p/N)(1 - R\varepsilon^2 \cos^2(\pi p/N)/6)$ and the Hadamard regular part each contribute, and their angular averages combine to $1/4$.) $\square$

Lemma 3.6 (Regge bridge to smooth metrics). *The entropy-maximising Regge metric converges to the entropy-maximising smooth metric as the mesh size $h \rightarrow 0$.*

Proof. Any smooth 2D Riemannian manifold (S, g) is approximated by a simplicial complex (S_h, g_h) in which curvature is concentrated at V vertices, each carrying a deficit angle $\varepsilon_v \approx R(v) A_v$ (with A_v the dual area). The convergence $g_h \rightarrow g$ as $h \rightarrow 0$ is proved [Cheeger et al., 1984]; the Green's function $G_h \rightarrow G$ in the H^1 -norm (standard finite-element convergence). Each Regge vertex is a point vortex (circulation $\kappa_v = \varepsilon_v/(2\pi)$, by the isomorphism), so the Regge system is a V -vortex configuration. The vertex positions are *simultaneously* vortex positions and Regge metric degrees of freedom (on a fixed triangulation, vertex positions determine edge lengths, hence the simplicial metric). This identification eliminates the apparent circularity of H : the Hamiltonian is a function of vertex positions alone, and entropy maximisation is a single finite-dimensional optimisation—not an iteration between metric and vortex configuration.

The Onsager prerequisites for this Regge system are established in Lemma 3.1. On a *flat* background ($\Lambda = 0$), conformal invariance (Paper II, Proposition 13.2) gives $G = -(2\pi)^{-1} \log|z-w|$, independent of the deficit angles. On a *curved* background ($\Lambda \neq 0$), the Hadamard expansion (Lemma 3.5) gives $C_1 = (N-1)[1 - R(x)\varepsilon^2/4 + O(\varepsilon^4)]$, affine in $R(x)$. Since $d^2 S/dC_1^2 < 0$, the entropy functional $S_\varepsilon[g] = \int_S S(C_1) dA_g$ is strictly concave in R . The Gauss–Bonnet constraint $\int R dA = 4\pi\chi$ and Jensen's inequality give the unique maximum at $R = \text{const}$. The strict concavity holds at every ε (eq. (3.4)); Γ -convergence [Dal Maso, 1993] provides an independent confirmation.

Given Lemma 3.1, the Onsager principle on the V -vertex Regge system selects the entropy-maximising vertex configuration, which by (H2a) has $C_1(v)$ uniform—hence $R = \text{const}$. In the continuum limit $V \rightarrow \infty$: $S_{\text{Regge}} \rightarrow S_{\text{total}}[g]$ (by the FEM convergence of the

Green's function, sharpened to $O(h^\infty)$ by the Szegő estimates below), and the entropy-maximising Regge metric converges to the entropy-maximising smooth metric (strict concavity ensures convergence of maximisers under pointwise convergence; Dal Maso 1993). The Szegő–Toeplitz convergence analysis (supplementary material) sharpens the rate to $O(h^\infty)$. $\square$

Proof of Theorem 3.3. (H2) + Lemma 3.5 force $R = \text{const}$. By the mathematical content of (H2): $S_{\text{total}}[g] = \int S(x) dA$ is strictly concave in $C_1(x)$, so Jensen's inequality gives a unique entropy maximum at uniform C_1 . The Onsager principle (the physical content of (H2)) selects this maximum: at negative temperature, the partition function is dominated by the macrostate maximising S_{total} . Write C_1^{eq} for this uniform value. Substituting the expansion (3.3):

$$(N-1) \left[1 - \frac{R(x) \varepsilon^2}{4} + O(\varepsilon^4) \right] = C_1^{\text{eq}} \quad \text{for all } x.$$

The right side is x -independent. Therefore $R(x) = R_0$ for all x . The *value* of C_1^{eq} (and hence R_0 and Λ) is determined separately by the energy constraint: the WDW equation (Proposition 3.7) gives $C_1^{\text{eq}} = f(m^*, N)$ at the BO threshold, with $R_0 = -4[f(m^*, N) - (N-1)]/[(N-1)\varepsilon_0^2]$.

The strict concavity $d^2 S/dC_1^2 = -\frac{1}{2} \sum 1/\lambda_m^2 < 0$ holds on the domain $\lambda_m > 0$ for all $m \neq m^*$ (i.e., away from the stability threshold where $\lambda_{m^*} = 0$). The uniformised metric ($R = R_0 = \text{const}$, $\text{Var}(R) = 0$) lies in the *interior* of this concavity domain, since $\lambda_m(R_0) = C_1(R_0) - f(m) > 0$ for all $m \neq m^*$. The Onsager contraction (Corollary 3.4), with basin $\text{Var}(R) \leq 2$, keeps the metric within the concavity domain: within the basin, $|C_1(x) - C_1(R_0)| = O(\varepsilon^2 |\delta R|)$ is small, so $\lambda_m > 0$ is preserved. The Jensen argument therefore operates entirely within the concavity domain, with no bootstrap circularity. The CMS integrability makes the entropy exact (the angle-action variables decouple, so the one-loop entropy equals the microcanonical entropy; this uses Lemma 5.2 from Paper V). The energy constraint is compatible: V is monotone in C_1 , so fixing $\bar{V}$ fixes $\bar{C}_1$ and the Lagrange multiplier is Λ .

Strict concavity. The concavity of S in C_1 is exact (CMS integrability), so Jensen's inequality applies to C_1 directly: $\int_M S(C_1(x)) dA \leq A \cdot S(\bar{C}_1)$ with equality if and only if $C_1(x) = \text{const}$. The Jensen step requires no ε -expansion. The translation from $C_1 = \text{const}$ to $R = \text{const}$ is exact. By the Riemannian Havelock identity (Paper I, Theorem 4.2), $C_1(x) = \gamma_\varepsilon(x) + b(N)$ where γ_ε is the regularised Robin function of the Laplacian at scale ε . The Robin function satisfies the identity $\Delta \gamma = K - \bar{K}$ (Aubin 1998, Theorem 4.13), where K is the Gaussian curvature and $\bar{K} = 2\pi\chi/A$. If $C_1 = \text{const}$ then $\gamma_\varepsilon = \text{const}$, so $\Delta \gamma_\varepsilon = 0$, giving $K(x) = \bar{K}$ everywhere, hence $R = \text{const}$. For the regularised Robin function ($\gamma_\varepsilon = \int_{\varepsilon^2}^\infty [K(x, x; t) - 1/A] dt$), the identity becomes $\Delta \gamma_\varepsilon = K - \bar{K} + R_\varepsilon$, where $R_\varepsilon = O(\varepsilon^2 \Delta K)$ is the heat-kernel remainder (Seeley, 1967: $|R_\varepsilon| \leq C_S \varepsilon^2 \|\Delta K\|_{L^2}$ with C_S depending only on the surface area and genus). $C_1 = \text{const}$ gives $\Delta \gamma_\varepsilon = 0$, so $K - \bar{K} = -R_\varepsilon$. Define $u = K - \bar{K}$. Then $\|u\|_{L^2} = \|R_\varepsilon\|_{L^2} \leq C_S \varepsilon^2 \|\Delta u\|_{L^2}$. By Schauder elliptic estimates on the compact surface, $\|\Delta u\|_{L^2} \leq C_E \|u\|_{H^2} \leq C'_E (\|u\|_{L^2} + \|\Delta u\|_{L^2})$; substituting and rearranging: $\|u\|_{L^2} \leq \frac{C_S C'_E \varepsilon^2}{1 - C_S C'_E \varepsilon^2} \|u\|_{L^2}$. For ε small enough that $C_S C'_E \varepsilon^2 < 1$, the only solution is $\|u\|_{L^2} = 0$, hence $K = \bar{K}$ exactly. The contraction constant $\kappa = C_S C'_E \varepsilon^2 < 1$ is guaranteed because ε is the vortex core size (Cauchy blob width), which is a free regularisation parameter—the limit $\varepsilon \rightarrow 0$ recovers the point-vortex result with $\kappa \rightarrow 0$.

As a quantitative check on the second variation, the ratio

$$\frac{(N-1)\sigma_2}{32\alpha\sigma_1} \geq 7.5 \quad \text{for all } N \geq 7, \quad (3.4)$$

where $\sigma_k = \sum_{m \neq m^*} \lambda_m^{-k}$ and $\alpha = 1/48$ is the second Seeley–DeWitt coefficient, confirms that the concavity term dominates the curvature-nonlinearity correction by at least 7.5 : 1 (see Code Availability).

The extension from conical deficits to smooth metrics is established by Lemma 3.6.

Partition function interpretation. The total polygon partition function $Z[g] = \sum_N e^{-b(N)} Z_{\text{frozen}}(N, g)$ is a functional of the metric g through $Z_{\text{frozen}} = \prod_m |\lambda_m(g)|^{-1/2}$ (Proposition 12.1). By the vortex-gravity isomorphism, this is the *gravitational* one-loop partition function summed over polygon sectors: $Z[g] = Z_{\text{grav}}[g]$, with $e^{-b(N)}$ playing the role of $e^{-S_{\text{EH}}}$ (the classical Euclidean action of the N -cone geometry). The saddle-point condition $\delta(-\log Z)/\delta g = 0$ is therefore the Einstein equation—specifically, it is the Euler–Lagrange equation for the combined classical action $b(N)$ plus the one-loop correction $-\log Z_{\text{frozen}}$. This identification gives the WDW equation (§8) its physical meaning: it is the *quantum mechanics of the gravitational partition function*. The WKB tunnelling action $S_{\text{BO}}(N)$ is the Euclidean action of the dominant instanton connecting the two turning points of $V(\rho)$, and $\exp(-S_{\text{BO}})$ is the tunnelling amplitude between the Planck-scale and cosmological-scale vacua. The hierarchy $\ell \cdot v = \exp(S_{\text{BO}})$ (Paper IV, §14) is a consequence of the partition function interpretation: the cosmological scale ℓ is set by the instanton suppression of the polygon partition function.

The explicit $\delta S/\delta g$ computation (Proposition 3.2) gives a direct proof that $R = \text{const}$, independent of the Jensen argument.

Higher dimensions. The above arguments are dimension-independent; the extension to spacetime dimension $d_{\text{st}} \geq 3$ is detailed in Remark 3.4. In $d_{\text{st}} = 3$, $R = \text{const}$ suffices (zero Weyl components). In $d_{\text{st}} \geq 4$, the full Riemann tensor (including all Weyl components) is determined by N via the KK decomposition, Lichnerowicz spectrum, and Fefferman–Graham reconstruction (Theorem 3.9). $\square$

3.3 The cosmological constant and Clausius relation

Remark 3.2 (The logical structure of (H2)). *Three points merit emphasis.*

1. Onsager, not relaxation. *The entropy maximisation comes from the Onsager selection principle (dominant microcanonical macrostate at negative temperature), not from dynamical relaxation of the integrable CMS modes.*
2. Two variational principles. *(H2a) characterises the entropy functional at fixed geometry; (H2b) selects the entropy-maximising metric via the Onsager principle. Neither implies the other.*
3. The BO bridge. *At fixed g , the angular modes determine $S(x, g)$; varying g varies $S_{\text{poly}}[g] = \int S(x, g) dA$. The BO effective potential $V_{\text{BO}} = V_0 - S$ connects the two: maximising S at fixed energy is equivalent to minimising V_{BO} .*

Proposition 3.7 (The cosmological constant from the WDW equation). *The value of Λ in Theorem 3.3 is determined by the Wheeler–DeWitt equation for the breathing mode ρ :*

$$\left[-\frac{1}{2c} \frac{d^2}{d\rho^2} + V(\rho) \right] \Psi = 0, \quad V(\rho) = \Phi(\rho) - f(m^*, N).$$

Two independent ingredients determine this equation:

The mechanism (*why $\dot{\rho}^2$ appears*). The Kirchhoff equations are first-order; the breathing mode ρ is a collective coordinate with no canonical momentum. The second-order kinetic term is geometric, not perturbative: the functional determinant of the first-order angular-mode operator $\partial_\tau + \lambda_m(\rho(\tau))$ depends only on $\int \lambda_m d\tau$ (no $\dot{\rho}$ dependence), so the $\dot{\rho}^2$ term does not arise from the action. It arises from the Feynman–Kac path-integral measure $\mu[\rho] = \prod_\tau \sqrt{c/(2\pi d\tau)}$, whose coefficient c is the Fisher–Rao information metric on the family of angular-mode Boltzmann distributions parametrised by ρ . The coupling is mode-independent ($\Phi' = \coth \rho$ for all m), and the CMS integrability (exact action-angle decoupling) ensures the one-loop effective action is exact at $O(c^0)$.

The coefficient c (1.1). The WDW kinetic coefficient is the mode-independent offset $b(N)$ from the eigenvalue decomposition (Proposition 9.2), not the CL friction $\gamma \sim O(N)$. It is uniquely fixed by the following argument. The CMS integrability gives the exact one-loop effective action $S_{\text{eff}}(\rho) = -\frac{1}{2} \sum_m \log |\lambda_m(\rho)|$. The Feynman–Kac theorem maps the classical partition function $Z = \int e^{-\beta S_{\text{eff}}} d\rho$ to the quantum propagator $\langle \rho | e^{-\beta H} | \rho \rangle$ with $H = -(1/2c) d^2/d\rho^2 + V(\rho)$. The kinetic coefficient c is the inverse diffusion constant of the breathing mode in the angular bath. For the CMS system, the action variables I_m are exact constants of motion, so the BO adiabatic separation is exact: the Havelock eigenvectors $e_m = N^{-1/2} (\omega^{mp})_{p=0}^{N-1}$ with $\omega = e^{2\pi i/N}$ are Fourier modes, independent of ρ , so non-adiabatic corrections vanish identically. The effective potential decomposes as $V(\rho) = -(1/2) \sum_m \log |C_1(\rho) - f(m, N)|$; writing $C_1 = b(N) + [C_1(\rho) - b(N)]$, the coefficient $b(N)$ is the unique value at which the linear response $\partial V / \partial C_1 = -(1/2) \sum 1/\lambda_m$ is mode-independent (because $b(N)$ is the average of $\lambda_m + f(m, N)$ over m , and the mode sum has a single C_1 -dependent prefactor). Any other kinetic coefficient $c' \neq 12b(N)$ would give a WDW potential that depends separately on each mode m , breaking the BO factorisation. (The classical BO separation is exact for any c because the Havelock eigenvectors are ρ -independent; the constraint is on the quantum BO factorisation, which requires c constant for the reasons in §8, step 1. The matching $c = 12b(N)$ is a derivation, not a theorem: it is the unique value consistent with both Brown–Henneaux and the CMS–CS Casimir identification.) In the classical limit $c \rightarrow \infty$: $V(\rho^*) = 0$, giving $\Phi(\rho^*) = f(m^*, N)$ and hence $\Lambda = -4[f(m^*, N) - (N-1)]/[(N-1)\varepsilon^2]$.

Corollary 3.8 (Sign of the cosmological constant from the polygon number). *The sign of Λ is fixed by the polygon: $N \leq 6$: $f(m^*, N) < N-1$, so $\Lambda > 0$ (positive curvature, S^2 -like—surface curvature destabilises); $N = 7$: $f(3, 7) = 6 = N-1$, so $\Lambda = 0$ (flat); $N \geq 8$: $f(m^*, N) > N-1$, so $\Lambda < 0$ (negative curvature, $\mathbf{H}^2$ -like—surface curvature stabilises). The flat threshold $N = 7$ —Havelock’s 1931 stability boundary—sits exactly at $\Lambda = 0$.*

Remark 3.3 (Relation to Schur’s lemma). *In 2+1D, the higher-dimensional extension of the proof is a special case of Schur’s lemma (isotropic sectional curvature implies constant curvature). The added value of the Havelock derivation is fourfold: (a) the full argument is new: polygon stability implies isotropic Green’s function, which Schur alone cannot provide; (b) the curvature value Λ is quantised by the polygon number N through the explicit formula $\Lambda = -4[f(m^*, N) - (N-1)]/[(N-1)\varepsilon^2]$; (c) the hypothesis of Schur’s lemma (isotropy) is derived from the Hadamard universality of the Green’s function, not assumed; (d) in $d \geq 4$, the Havelock condition (H1) provides a testable criterion (mode-independence of C_1) equivalent to the full Weyl-vanishing condition of Schur’s lemma, connecting the abstract curvature condition to the observable stability threshold.*

3.4 Higher-dimensional extension and verification

Remark 3.4 (Characterisation theorem in $d \geq 3$). *The core argument is dimension-independent: the csc^2 uniqueness forces a logarithmic interaction in any d (Lemma 3.5); the Hadamard parametrix $G(x, y) = U/\sigma^{(d-2)/2} + V \log \sigma + W$ contains $R(x)$ at order ε^2 in every d ; and (H3) forces $R = \text{const}$ (Jensen + Onsager). The conclusion closes differently by dimension:*

Dimension	Weyl components	Status
$d = 2$ ($2+1D$)	0	Theorem: (H1) proved by the Havelock identity
$d = 3$ ($3+1D$)	10	Theorem: KK + Lichnerowicz + FG (Theorem 3.9)
$d \geq 4$	≥ 35	Same, with more traceless Ricci components

In $d = 2$: the Havelock identity proves (H1) analytically—no assumption is needed. In $d \geq 3$: condition (H1) (mode-independence of C_1) is equivalent to the spatial Green’s function being isotropic, which is equivalent to maximal symmetry, which is equivalent to the vacuum Einstein equations with Λ . The soft Havelock decomposition (Proposition 10.2) shows that (H1) is approximately satisfied on any background, with residual $|R_m| < a|f_m|$ (verified for $N = 5, \dots, 20$ and $d = 3, 4, 5$): the violation of (H1) is controlled and measures the departure from Einstein geometry.

Theorem 3.9 (Full Riemann curvature from the polygon number). *On the Seifert manifold $\mathbf{H}^2 \times_N S^1$ with stabilised fibre ($\sigma = 1$) and quantised flux ($F = N/(2\pi)$), the full 4D Riemann tensor $R_{MNPQ}^{(4)}$ —including all 10 Weyl components—is uniquely determined by N .*

Proof. Three layers: background, linearised, nonlinear.

Layer 1 (Background). The Weyl tensor vanishes in 3D, so the Riemann tensor of the Seifert manifold is entirely determined by its Ricci tensor $R_{ij}^{(3)} = \Lambda_3 g_{ij}$ (Theorem 3.3):

$$R_{ijkl}^{(3)} = \frac{\Lambda_3}{2}(g_{ik}g_{jl} - g_{il}g_{jk}). \quad (3.5)$$

The O’Neill formula for the Riemannian submersion $\pi: \mathbf{H}^2 \times_N S^1 \rightarrow \mathbf{H}^2$ (O’Neill 1966; Besse 1987, Thm. 9.36) decomposes $R_{ijkl}^{(3)}$ into base curvature + terms quadratic in the connection curvature F_{ij} ; the fibre components satisfy $R_{\varphi i \varphi j}^{(3)} = \frac{1}{4}F_{ik}F_j^k$. The flux is uniform (F_{ij} determined by the Chern class $c_1 = N$; $\nabla_{[i}F_{jk]} = 0$ on the constant-curvature base), and the fibre radius is stabilised ($\sigma = 1$, $V''(1) = N^2/4 > 0$). The KK lift to 4D (Paper IV, Step B) adds the time direction; the resulting $R_{MNPQ}^{(4)}$ inherits the same inputs ($R^{(3)}, F, \sigma$), each an explicit function of N . Therefore all 20 components of $R_{MNPQ}^{(4)}$ are determined by N .

Layer 2 (Linearised perturbations). A perturbation h_{MN} decomposes under KK into h_{ij} (metric), δA_i (graviphoton), $\delta \sigma$ (radion). The Lichnerowicz operator Δ_L on the multi-cone geometry (with distributional curvature $R = R_0 + \sum_k \delta_k \delta_{z_k}$ regularised by the Cauchy blob at scale ε ; Paper II, Proposition 7.1) reduces to the Havelock Hessian as follows. The $\mathbb{Z}_N$ rotation symmetry of the regular polygon acts as an isometry of the multi-cone metric. Since Δ_L is a geometric operator, it commutes with this $\mathbb{Z}_N$ action and decomposes into irreducible representations labelled by $m = 0, \dots, N-1$ (Schur’s lemma). At each cone point, the Green’s function has the angular kernel

$-\log |2 \sin(\pi p/N)|$ (the logarithmic singularity of the Hadamard parametrix). The angular Hessian weight is $1/(4 \sin^2(\pi p/N))$, so the m -th $\mathbb{Z}_N$ Fourier eigenvalue of the Hessian is $\sum_{p=1}^{N-1} (1 - \cos(2\pi pm/N)) / (2 \sin^2(\pi p/N)) = m(N-m)/2$ (Paper I, Lemma 4.1). The full eigenvalue is $\lambda_m = C_1 - m(N-m)/2$ (the mode-independent offset C_1 absorbs the $\csc^2$ diagonal sum and the radial contribution). The spectrum $\{\lambda_m\}$ is therefore determined by N . The graviphoton satisfies the Hodge Laplacian on $\mathbf{H}^2/\mathbb{Z}_N$ (spectral gap ≥ 1). The radion is massive ($m^2 = N^2/4$, no zero mode). All perturbation sectors are gapped and determined by N , so the linearised Weyl tensor δC_{MNPQ} is determined mode by mode.

Layer 3 (Full nonlinear). The Fefferman–Graham expansion in AdS_3 terminates at finite order for $d_\partial = 2$ [de Haro et al., 2001, Toldo and Willett, 2018]: the boundary stress tensor T_{ab} (determined by the Virasoro algebra at $c = 12b(N)$) uniquely reconstructs the bulk 3D metric. The boundary has $\mathbb{Z}_N$ orbifold singularities (the polygon cone points), but the *bulk* geometry is smooth: the FG expansion is well-defined in the distributional sense, with g_2 acquiring delta-function contributions at the cone points and g_4 remaining regular (the δ^2 terms cancel by the $\mathbb{Z}_N$ symmetry of the metric). The orbifold boundary is a Lipschitz Riemannian manifold with isolated cone singularities of angle $2\pi/N$; the scattering operator extends to this setting by Graham and Zworski [2003] (smooth boundaries) combined with the Cheeger–Taylor cone-edge calculus [Cheeger and Taylor, 1982], which shows the resolvent $(\Delta_g - s(2-s))^{-1}$ is meromorphic on cone manifolds with cone angles $< 2\pi$ (the polygon cone angles are $2\pi/N < 2\pi$ for all $N \geq 2$, so this condition is always satisfied). The KK lift with stabilised fibre and quantised flux then gives the unique 4D metric $g_{MN}^{(4)}$, and hence the unique 4D Riemann tensor. $\square$

Supporting evidence: spectral focusing. By Theorem 10.3 and Corollary 10.4, $N_{\text{crit}}(\Delta)$ is non-increasing: once a polygon type becomes unstable, it remains unstable for all larger Δ .

The monotonicity is formally analogous to Raychaudhuri focusing (structural parallel):

Concept	General relativity	Polygon system
Flow parameter	Affine parameter λ	Conformal dim Δ
Expansion	$\theta = \nabla \cdot k$ (decreasing)	$N_{\text{crit}}(\Delta)$ (decreasing)
Focusing	$d\theta/d\lambda \leq 0$	$dN_{\text{crit}}/d\Delta \leq 0$ (monotone)
Trapped	$\theta < 0$ (trapped surface)	$\lambda_m < 0$ (unstable mode)
Energy condition	$R_{\mu\nu} k^\mu k^\nu \geq 0$	$C_1(\xi) > 0$ (spectral gap)

The palindromic threshold as a Killing horizon. The WDW equation $[-1/(2c) \Psi'' + V(\rho) \Psi = 0]$ is the zero-energy Schrödinger equation in the potential V . The turning point $V(\rho^*) = 0$ is a *Killing horizon* in the minisuperspace geometry: the Euclidean continuation around ρ^* requires periodicity $\beta = 2\pi/|V'(\rho^*)| = 2\pi/\coth(\rho^*)$ to avoid a conical singularity [Gibbons and Hawking, 1977]. This is the same calculation that gives the de Sitter temperature $T = H/(2\pi)$ —with the breathing mode ρ playing the role of the scale factor a . The Gibbons–Hawking temperature at the palindromic threshold is

$$T = \frac{V'(\rho^*)}{2\pi} = \frac{\coth(\rho^*)}{2\pi}. \quad (3.6)$$

Logical role of the Clausius relation. The vacuum Einstein equations $R = \text{const}$ are derived *independently* from entropy maximisation (Theorem 3.3). The Clausius relation

at the Gibbons–Hawking horizon serves a *different* purpose: it extracts the gravitational coupling G_{eff} from the constant-curvature geometry. On a constant-curvature surface, there is one threshold (one N) and one temperature T . The Jacobson argument [Jacobson, 1995, Lashkari et al., 2014] then promotes G_{eff} to the matter coupling: $\delta Q = T \delta S$ at the Killing horizon gives the full Einstein equations $G_{\mu\nu} + \Lambda g_{\mu\nu} = 8\pi G_{\text{eff}} T_{\mu\nu}$.

The three thermodynamic quantities are determined by the BO potential $V(\rho) = \Phi(\rho) - f(m^*, N)$:

- $T = V'(\rho^*)/(2\pi) = \coth(\rho^*)/(2\pi)$ (Unruh temperature from the surface gravity $\kappa = V'$);
- $\delta Q = \kappa \times A \times \delta\rho = \coth(\rho^*) \times 2\pi \sinh(\rho^*) \times \delta\rho$ (surface gravity $\times$ horizon circumference $\times$ displacement);
- $\delta S = (dS/d\rho) \times A \times \delta\rho$ where $A = 2\pi \sinh(\rho^*)$ is the horizon circumference and $dS/d\rho$ is the Bekenstein–Hawking entropy production rate.

Which entropy. The one-loop entropy $S' = -\frac{1}{2} \log |\det' H|$ (reduced determinant, zero modes excluded) has a well-defined derivative at ρ^* :

$$\left. \frac{dS'}{d\rho} \right|_{\rho^*} = -\frac{\coth(\rho^*)}{2} \sigma(N), \quad \sigma(N) = \sum_{\substack{m=1 \\ f_m \neq f_{m^*}}}^{N-1} \frac{1}{f_{m^*} - f_m}, \quad (3.7)$$

where the sum runs over non-zero modes ($\sigma(7) = 8/3$, $\sigma(8) = 49/9$, verified $N = 7, \dots, 15$; see Code Availability). The zero mode's crossing contributes to the spectral flow (APS eta-invariant shift $= \pm 1$ per palindromic pair), not to $|\det' H|$.

Clausius relation via Cardy + RT. The Bekenstein–Hawking entropy $S_{\text{BH}} = A/(4G)$ presupposes G , so using it in the Clausius relation gives a tautology. A non-trivial determination uses the *Cardy thermal entropy* of the boundary CFT, which depends on c but not on G :

$$S_{\text{Cardy}} = \frac{\pi c}{3} T L = \frac{\pi c}{3} \cosh(\rho^*), \quad (3.8)$$

where $T = \coth(\rho^*)/(2\pi)$ is the Unruh temperature, $L = 2\pi \sinh(\rho^*)$ is the ring circumference, and c (1.1) is computed from the spectrum (Paper I). The Ryu–Takayanagi formula $S = A/(4G)$ with $A = 2\pi \sinh(\rho^*)$ then determines G :

$$G_{\text{RT}} = \frac{A}{4 S_{\text{Cardy}}} = \frac{3 \tanh(\rho^*)}{2 c}. \quad (3.9)$$

Newton's constant is $G_{\text{BH}} = 3/(2c)$ (Brown–Henneaux). The $\tanh(\rho^*)$ factor is a finite- ρ^* correction to the *Cardy entropy* (not to G itself): the Cardy formula gives $S_{\text{Cardy}} = (\pi c/3) \cosh \rho^*$, while the area law uses $A = 2\pi \sinh \rho^*$; the ratio $A/(4G_{\text{BH}} S_{\text{Cardy}}) = \tanh \rho^*$ measures how far the finite polygon is from the asymptotic ($\rho^* \rightarrow \infty$) BTZ regime. The correction $\tanh(\rho^*) = 1 - 2e^{-2\rho^*} + O(e^{-4\rho^*})$ vanishes exponentially: 6% at $N = 7$, 0.03% at $N = 11$, $< 10^{-7}$ for $N \geq 15$ (verified numerically for $N = 7, \dots, 20$; see Code Availability). The common factor $\coth(\rho^*) = V'(\rho^*)$ is not a free choice: it is determined by the csc^2 kernel ($V = \log(2 \sinh \rho)$ is the $\mathbf{H}^2$ Green's function). The one-loop spectral sum (3.7) provides an N -dependent quantum correction beyond the Cardy

entropy; incorporating it would improve the finite- ρ^* convergence. The [Lashkari et al. \[2014\]](#) entanglement argument shows that positivity of relative entropy $S(\rho\|\sigma) \geq 0$ in the unitary boundary CFT implies $\delta Q = T \delta S$ at every Rindler horizon. The Clausius relation is verified numerically at all palindromic thresholds $N = 7, \dots, 17$ (see Code Availability). The extension to the full matter-coupled Einstein equations $G_{\mu\nu} + \Lambda g_{\mu\nu} = 8\pi G T_{\mu\nu}$ requires the Jacobson thermodynamic argument ([Jacobson 1995](#); [Lashkari et al. 2014](#)): $\delta Q = T \delta S$ at every local Rindler horizon implies Einstein with matter. The identification of palindromic thresholds as Rindler horizons uses the modular Hamiltonian of the angular bath, which equals the Rindler energy by the Bisognano–Wichmann theorem. The prerequisites are satisfied: the boundary CFT is (i) unitary (DHVW orbifold of a unitary WZW model, Proposition 8.2), (ii) has a unique vacuum (the identity operator is $\mathbb{Z}_N$ -invariant), and (iii) the modular flow is geometric (the boundary T^2 is conformally flat). See Paper IV, Theorem 13.3).

Proposition 3.10 (Partition function verification via the Gauss product). *The polygon partition function $Z_{\text{poly}} = \sum_N e^{-b(N)} Z_{\text{frozen}}(N)$ equals the Chern–Simons partition function Z_{CS} on the Seifert manifold, with three independent mathematical inputs:*

1. Gauss product (algebra, Paper I): $b(N) = N(N+1)/12 + \frac{1}{N-1} \sum_{m=1}^{N-1} \ln \sin(\pi m/N)$. *The transcendental part of $b(N)$ is the mean log-sine over the $N-1$ pairwise chord distances—the regularised self-energy from the identity $\prod_{m=1}^{N-1} \sin(\pi m/N) = N/2^{N-1}$.*
2. Cheeger–Müller theorem (differential geometry): *the analytic torsion $\tau_{\text{an}} = \prod |\lambda_m|$ (spectral determinant of the gauge-covariant Laplacian) equals the Reidemeister torsion τ_R (combinatorial invariant of the Seifert manifold). For the N -gon flat connection: $Z_{\text{frozen}}(N) = \tau_{\text{an}}^{-1/2}$ is the standard one-loop normalization.*
3. CMS–CS Casimir identification (Proposition 2.1): *the vortex Hessian eigenvalues $\lambda_m = f(m^*, N) - f(m, N)$ are the eigenvalues of the cone Laplacian at the $\mathbb{Z}_N$ -symmetric point (proved via the Helgason symmetric space theorem, §2).*

The same Gauss product $\prod \sin(\pi m/N) = N/2^{N-1}$ appears on both sides:

- Vortex side: *it regularises the self-energy, giving the transcendental part of $b(N)$;*
- Spectral side: *it determines the cone spectral zeta function $\zeta'_{\text{cone}}(0)$, giving the analytic torsion.*

They agree because both compute the regularised determinant of the same $\mathbb{Z}_N$ -circulant csc^2 matrix—the vortex side via the Havelock decomposition (Paper I), the spectral side via the zeta function (Cheeger–Müller). Points (1) and (2) are mathematical theorems. Point (3) is the only step using a physical identification (the CMS–CS correspondence).

Matter coupling. On a quotient $\Gamma \backslash \mathbf{H}^2$, the Selberg correction δC_1 shifts the threshold $\rho^* \rightarrow \rho_{\text{matter}}^*$, but the Clausius ratio retains the form $4\pi \tanh(\rho_{\text{matter}}^*)$ (verified on the Bolza surface with $|\delta C_1| \leq 0.021$). The ratio 4π is a local geometric identity (the metric is locally $\mathbf{H}^2$), giving $G_{\text{eff}} = 1/(16\pi)$ in natural units ($8\pi G_{\text{eff}} = 1/2$ as $\rho^* \rightarrow \infty$).

4 The spectral bound as a quantum gravitational bound

The spectral correction and the gravitational one-loop partition function are both moments of the heat kernel:

$$\log Z_{\text{grav}} = - \int_0^\infty [\text{Tr } e^{-t\Delta_S} - 1] \frac{dt}{t} \quad (\text{weight } 1/t), \quad (4.1)$$

$$\delta C_1(z_0) = \int_0^\infty [K(z_0, z_0, t) - K_{\mathbf{H}^2}(t)] dt \quad (\text{weight } 1). \quad (4.2)$$

The first is the functional determinant $\det \Delta_S$, the second determines polygon stability. Both involve the Selberg zeta function $Z_\Gamma(s)$ at the zero $s = 1$, but this requires care.

The zero at $s = 1$. The Selberg zeta $Z_\Gamma(s)$ has zeros at s such that $s(1-s)$ is an eigenvalue of Δ_S . The constant eigenfunction ($\lambda_0 = 0$) gives a zero at $s = 1$, so $Z_\Gamma(1) = 0$ —not a finite number. The Euler product $\prod_\gamma \prod_k (1 - e^{-(s+k)\ell_\gamma})$ converges for $\text{Re}(s) > 1$ but diverges at $s = 1$ (by the prime geodesic theorem: the partial product over geodesics with $\ell \leq L$ decays to 0 as $L \rightarrow \infty$).

The physically relevant quantities are the *residue* at this zero:

- $Z'_\Gamma(1)$: the *regularised* determinant $\det' \Delta_S = \prod_{n \geq 1} \lambda_n$ (omitting $\lambda_0 = 0$). This is the gravitational one-loop partition function [Giombi et al., 2008].
- The spectral sum $\sum_{n \geq 1} |\varphi_n(z_0)|^2 / \lambda_n$: the regularised Green's function at the diagonal, which determines $\delta C_1(z_0)$.

Both are finite and well-defined; both involve the analytic structure of Z_Γ at its zero, not the (divergent) Euler product.

The spectral computation of δC_1 . At the Aut-fixed centre z_0 of the Bolza surface, only trivial-representation eigenfunctions contribute: $\delta C_1(z_0) = \sum_{n: \text{triv}} |\varphi_n(z_0)|^2 / \lambda_n \approx 0.011$, dominated by the first trivial-rep eigenvalue $\lambda_{\text{triv}} \approx 15.05$ (the first eigenvalue in the trivial representation of $\text{Aut}(\text{Bolza})$; the full spectral gap is $\lambda_1 \approx 3.84$; Strohmaier and Uski 2013), with rapid convergence: $1/(4\pi \cdot 15.05) \approx 0.005$, $1/(4\pi \cdot 23.08) \approx 0.003$, $1/(4\pi \cdot 32.67) \approx 0.002$, The first 8 trivial-representation eigenvalues of Strohmaier and Uski [2013] give $\sum_{n=1}^8 1/(4\pi \lambda_n) = 0.019$. The tail is bounded by Sogge's sup-norm estimate ($|\phi_n(z_0)|^2 \leq C \lambda_n^{1/2}$) combined with Weyl asymptotics ($\lambda_n \geq c n$ for the trivial representation): $\sum_{n>8} 1/(4\pi \lambda_n^{3/2}) < 0.001$. Therefore $|\delta C_1| \leq 0.019 + 0.001 = 0.020$.

The bound $|\delta C_1| \leq 0.020$ on the Bolza surface is conditional on the Strohmaier–Uski spectral data and unconditional with respect to the Selberg conjecture ($\lambda_{\text{triv}} \approx 15.05 \gg 1/4$). In gravitational language, C_1 is the classical stability coefficient and δC_1 is the one-loop quantum correction from the topology of the spatial slice.

The Selberg dictionary. Four physical quantities involve the Selberg zeta at $s = 1$ (structural analogies, not functional inequalities):

Z_Γ data at $s = 1$	Physical meaning
$Z'_\Gamma(1)$ (regularised determinant $\det' \Delta_S$)	Gravitational partition function $= \prod_{n \geq 1} \lambda_n$
$\sum_{n \geq 1} \varphi_n(z_0) ^2 / \lambda_n$ (spectral; convergent)	Vortex stability correction δC_1 Localised at z_0 ; Aut-suppressed
Palindromic threshold polynomials (from $C_1(\xi) = f_{\max}(N)$)	Stability transitions: values of ξ where N_{crit} jumps
Hecke eigenvalues of associated forms (a_p from modular forms)	Prime-by-prime visibility Which thresholds are visible on $\Gamma(p)$

The first two rows extract different data from the same zero $s = 1$: the determinant ($\prod \lambda_n$) versus the resolvent trace ($\sum 1/\lambda_n$). The palindromic thresholds control *where* stability transitions occur (at which curvature ξ); the eigenvalue spectrum controls *how large* the correction δC_1 is. On highly symmetric surfaces (the modular surface, the Bolza surface), the systole and the palindromic geodesics coincide; on a generic arithmetic surface they are distinct.

5 Holographic interpretation

The preceding sections derived gravity from the polygon's statistical mechanics without invoking holography. This section translates the same results into the $\text{AdS}_3/\text{CFT}_2$ language; the translation is not required for any subsequent result, but it connects the polygon framework to the broader gauge/gravity literature and provides independent consistency checks.

In $\text{AdS}_3/\text{CFT}_2$, a point mass m in the bulk creates a conical deficit that corresponds to the insertion of a heavy operator $\mathcal{O}$ of conformal dimension $h = (c/24)(1 - (1 - 8Gm)^2)$ in the boundary CFT [Carlip, 1998]. A regular N -gon of equal masses at radial coordinate r in the bulk corresponds to N equally-spaced heavy operator insertions on the boundary circle.

In this language, the Havelock decomposition takes the form:

Bulk (2+1D AdS gravity)	Boundary (1+1D CFT)
Perturbation mode m	$\mathbb{Z}_N$ conformal block channel
$f(m, N) = m(N-m)/2$	Casimir of the m -th channel
$C_1(\xi)$	Function of h/c and AdS radius
$\lambda_m = C_1 - f(m, N) > 0$	Positivity of the m -th block
$C_1 \geq f_{\max}(N)$	Bootstrap-type bound on h/c

The stability condition $C_1 \geq m(N-m)/2$ for all m becomes a *lower bound on the conformal dimension* of the boundary operators: below this bound, the bulk polygon is unstable and the corresponding $\mathbb{Z}_N$ conformal block changes sign. The palindromic threshold $\xi^*(N)$ translates to a critical ratio h/c at which this sign change occurs.

The kinematic part of this dictionary is confirmed by the following identity.

Proposition 5.1 (Havelock Casimir = one-chiral OPE Casimir). *Let $T_m = \sum_{p=1}^{N-1} \frac{1 - \cos(2\pi pm/N)}{2 \sin^2(\pi p/N)}$ be the $\mathbb{Z}_N$ Fourier eigenvalue of the OPE kernel $\csc^2(\pi p/N)$. Then*

$$T_m = m(N-m) = 2f(m, N), \quad (5.1)$$

exactly twice the Havelock Casimir $f(m, N) = m(N-m)/2$.

Proof. Write $\omega = e^{2\pi i/N}$ and $S_m = \sum_{p=1}^{N-1} (1 - \omega^{pm})/|1 - \omega^p|^2$. Since $|1 - \omega^p|^2 = 2(1 - \cos(2\pi p/N)) = 4\sin^2(\pi p/N)$ and $1 - \omega^{pm} = 1 - \cos(2\pi pm/N) - i\sin(2\pi pm/N)$, the real part gives $T_m = \text{Re}(S_m) = \sum_{p=1}^{N-1} [1 - \cos(2\pi pm/N)]/(2\sin^2(\pi p/N))$. Now consider the polynomial identity on $\mathbb{Z}/N\mathbb{Z}$: the function $f_m(p) = [1 - \cos(2\pi pm/N)]/(2\sin^2(\pi p/N))$ extends to a trigonometric polynomial of degree $\leq N-1$. Its sum equals the coefficient of z^0 in the Laurent expansion of $\frac{1-z^m}{1-z} \cdot \frac{1-z^{-m}}{1-z^{-1}}$, evaluated at $z = \omega$ and summed over $p = 1, \dots, N-1$. This product is $|\sum_{k=0}^{m-1} z^k|^2$, whose z^0 coefficient (from the orthogonality $\sum_{p=0}^{N-1} \omega^{pm} = N\delta_{n,0}$) is $\min(m, N-m) \cdot \max(m, N-m) = m(N-m)$, since exactly m terms contribute in each sector. $\square$

The factor of 2 has physical content. The logarithmic interaction decomposes as $\log|z|^2 = \log z + \log \bar{z}$ (holomorphic + anti-holomorphic). In the boundary CFT, the OPE of two heavy operators has both chiral sectors, giving the *full* Casimir $m(N-m)$. The Havelock eigenvalue uses $f = m(N-m)/2$, the contribution of *one* chiral sector, because:

- The Kirchhoff equations $\dot{z}_k = -i\partial H/\partial \bar{z}_k$ are first-order: coordinates are conjugate to each other (no separate momenta). This halves the phase-space dimension, selecting one chiral sector of the $\log|z|^2 = \log z + \log \bar{z}$ interaction.
- The stability criterion $\lambda_m > 0$ is $C_1 > m(N-m)/2$: the confining potential must exceed the holomorphic OPE repulsion in the m -th $\mathbb{Z}_N$ channel.

In gravity (second-order dynamics), both sectors contribute to the equations of motion, but the stability criterion—positivity of the Hessian eigenvalue—still involves $m(N-m)/2$ per sector (Proposition 13.1).

The remaining step is to express C_1 in terms of boundary CFT data. On the flat plane, $C_1 = N-1$ (the Havelock self-energy); on $\mathbf{H}^2$ at curvature parameter ξ ,

$$C_1(\xi) = (N-1) \frac{1+\xi^2}{(1-\xi)^2}. \quad (5.2)$$

The curvature parameter ξ encodes the radial position of the bulk polygon; through the Brown–Henneaux relation $c = 3\ell/(2G)$ and the deficit angle–dimension map $h = (c/24)(1-(1-8Gm)^2)$, the stability condition becomes a bound on conformal dimensions parametrised by the bulk position. The palindromic thresholds $\xi^*(N)$ give algebraic-number bulk positions at which the stability bound is saturated:

N	$\xi^*(N)$	d_{crit}/ℓ	$\varepsilon = 1-\xi^*$	h/c_{max}	Field
8	$8-3\sqrt{7}$	0.51	0.937	0.018	$\mathbb{Q}(\sqrt{7})$
10	$1/7$	0.80	0.857	0.015	$\mathbb{Q}$
12	$(18-5\sqrt{11})/7$	0.97	0.798	0.013	$\mathbb{Q}(\sqrt{11})$
23	φ^{-2}	1.44	φ^{-1}	0.007	$\mathbb{Q}(\sqrt{5})$

Here d_{crit}/ℓ is the geodesic distance from the AdS centre in units of the AdS radius, and $\varepsilon = 1 - \xi^*$ is the boundary proximity (the polygon must sit at $\varepsilon < \varepsilon^*(N)$, i.e., sufficiently close to the AdS boundary, to be stable). The $N = 23$ threshold gives $\varepsilon^* = \varphi^{-1} = (\sqrt{5}-1)/2$: the golden ratio appears as a critical bulk-position parameter in the holographic dictionary.

The two constraints on the polygon are independent:

- *Mass*: $h/c < h/c_{\max}(N)$ (total deficit $< 2\pi$; from the bulk geometry).
- *Position*: $\xi > \xi^*(N)$ (confining potential $>$ OPE repulsion; from the Havelock decomposition and the chiral identity (5.1)).

Together, they define the allowed region in the $(h/c, \xi)$ parameter space for a stable N -gon of heavy operators—a two-parameter family of bootstrap-type bounds, one for each N , all descending from the single identity $\lambda_m = C_1 - m(N-m)/2$.

The confining potential from the OPE kernel. The formula $C_1 = (N-1)(1+\xi^2)/(1-\xi)^2$ can be derived either from the Kimura Hamiltonian for point vortices on $\mathbf{H}^2$ (Paper I, §4) or from the second radial variation of the chordal OPE kernel $\sigma(z, w) = |z-w|/|1-\bar{z}w|$ at the N -gon. The two derivations are mathematically identical: both compute the Hessian of $-\sum_{j<k} \log \sigma(z_j, z_k)$ on the Poincaré disk. The factor $(1+\xi^2)/(1-\xi)^2$ is *not* the conformal factor $\Omega(\sqrt{\xi})^2 = 1/(1-\xi)^2$ alone; the additional $(1+\xi^2)$ comes from the angular structure of σ under radial perturbations.

Relation to the conformal bootstrap. The stability condition $\lambda_m > 0$ and the conformal bootstrap positivity conditions (crossing symmetry + unitarity, [Rattazzi et al. 2008](#)) are both positivity conditions on conformal block coefficients, but in *different channels*:

	Bootstrap	Vortex stability
What is constrained	CFT spectrum $\{\Delta_i\}$	Bulk polygon position ξ
Input	c, Δ_{gap}	$N, h/c$
Positivity condition	OPE $f_{ijk}^2 \geq 0$	$\lambda_m = C_1 - f(m, N) > 0$
Channel	generic 4-point OPE	$\mathbb{Z}_N$ -symmetric

The two are *complementary*: the bootstrap constrains which CFTs exist, while the vortex stability constrains which bulk configurations are stable in a given CFT. The identification $\lambda_m \leftrightarrow$ OPE coefficient would make the stability condition a bootstrap condition for $\mathbb{Z}_N$ -symmetric correlators—a direction requiring the $\mathbb{Z}_N$ conformal block computation sketched below.

6 The $\lambda_m \leftrightarrow$ OPE identification

The Havelock eigenvalue is entirely computable from boundary CFT data. The key is to decompose the semiclassical N -point function $\langle \mathcal{O}(z_1) \cdots \mathcal{O}(z_N) \rangle$ of N equally-spaced heavy operators into $\mathbb{Z}_N$ channels: the angular channels give $m(N-m)/2$ (the kinematic OPE Casimir, Proposition 5.1), and the radial confinement gives $C_1(\xi)$.

In the semiclassical (large- c) limit, the N -point function factorises as

$$\log \langle \mathcal{O}(z_1) \cdots \mathcal{O}(z_N) \rangle \approx -2h \sum_{j<k} \log \sigma(z_j, z_k) + O(1), \quad (6.1)$$

where $\sigma = |z_j - z_k|/|1 - \bar{z}_j z_k|$ is the chordal distance. The leading term is the “free OPE energy”; the $O(1)$ correction comes from the Virasoro block (graviton exchange in the bulk).

Under *angular* perturbation $z_k \rightarrow R e^{i(2\pi k/N + \delta\theta_k)}$ at the N -gon, the second variation of the leading term decomposes into $\mathbb{Z}_N$ Fourier modes $\delta\theta_k = \sum_m a_m \cos(2\pi mk/N)$:

$$\delta^2 \left[-2h \sum_{j < k} \log \sigma \right] = 2h \sum_{m=1}^{N-1} m(N-m) |a_m|^2. \quad (6.2)$$

Proposition 6.1 (Angular OPE stiffness). *The coefficient of $|a_m|^2$ in (6.2) is $2h \cdot m(N-m) = 2h \cdot T_m$, where $T_m = m(N-m)$ is the full OPE Casimir (Proposition 5.1). This follows analytically from Proposition 5.1 ($T_m = m(N-m)$, proved) and the constrained second variation; verified numerically for $N \leq 12$.*

The missing piece is the confining potential $C_1(\xi)$. We now show it too arises from the boundary.

Theorem 6.2 ($\mathbb{Z}_N$ conformal block = Havelock eigenvalue). *Let $\mathcal{O}_1, \dots, \mathcal{O}_N$ be N identical heavy operators of dimension h , placed at equally-spaced points $z_k = \sqrt{\xi} e^{2\pi i k/N}$ on the boundary of the Poincaré disk. In the semiclassical limit ($c \rightarrow \infty$, h/c fixed):*

- (a) Angular stiffness. *The m -th $\mathbb{Z}_N$ channel of the angular second variation of the N -point function gives $2h \cdot m(N-m)$ (Proposition 6.1).*
- (b) Radial confinement. *The log-radial ($u = \log \sqrt{\xi}$) second variation of $-\sum_{j < k} \log \sigma(z_j, z_k)$ at the N -gon gives the mode-independent confining potential $C_1(\xi) = (N-1)(1+\xi^2)/(1-\xi)^2$.*
- (c) Full identification. *The constrained second variation (angular + radial at fixed angular impulse), divided by $2h$ per chiral sector, is*

$$\lambda_m = \underbrace{C_1(\xi)}_{\text{radial confinement}} - \underbrace{m(N-m)/2}_{\text{one-chiral OPE repulsion}}. \quad (6.3)$$

This is the Havelock eigenvalue, derived entirely from boundary data.

Proof. Part (a) is Proposition 6.1. For part (b): the log-radial second variation of the hyperbolic energy gives $C_1(\xi)$ as in (5.2) (Paper I, equation 5.7) (the Riemannian Havelock identity on $\mathbf{H}^2$). Part (c) combines (a) and (b) via the standard Lagrange-multiplier argument (Paper I, §3).

Identity block dominance for $N > 4$. The theorem requires that the identity Virasoro block dominates the OPE decomposition in each $\mathbb{Z}_N$ channel. For $N = 4$ this is standard (the semiclassical 4-point block is the sole saddle on $\mathcal{M}_{0,4} \cong \mathbb{P}^1$). For $N > 4$: the $\mathbb{Z}_N$ -symmetric N -point correlator of identical operators on $\mathbb{P}^1$ (*genus zero*) equals the 1-point function of the twist field σ_m on the orbifold $\mathbb{C}/\mathbb{Z}_N$ [Dixon et al., 1985]. The orbifold has twisted-sector vacua $|k\rangle$ with conformal dimensions $h_k = k(N-k)/(2N^2) \times c > 0$ for $k = 1, \dots, N-1$. On genus zero, these contribute to the 1-point function only as virtual intermediate states, suppressed by $O(e^{-2\pi h_k}) = O(e^{-c})$ (exponentially small at large c , since $h_k \propto c$). The identity block gives the leading (classical) contribution $\exp(-c \times S_{\text{conical}})$; the $\mathbb{Z}_N$ symmetry reduces the moduli space $\mathcal{M}_{0,N+1}^{\mathbb{Z}_N}$ to dimension 2 (ring radius and overall phase). On this 2D space, the identity block dominates by two independent bounds:

1. *Eigenvalue bound (Proposition 8.1):* the $O(1/c)$ Virasoro corrections satisfy $|\delta f_m / \lambda_m| < 1/4$ for all non-critical m , so no eigenvalue changes sign at finite c and the symmetric saddle remains non-degenerate.

2. *OPE coefficient bound (Liouville action)*: the $\mathbb{Z}_N$ twist field σ_k of order k lives on the orbifold $\mathbb{C}/\mathbb{Z}_N$ with conical metric of opening angle $2\pi/N$. Its OPE coefficient is $C_k^2 = \exp(-S_L(\text{cone}))$ where S_L is the classical Liouville action [Dixon et al., 1985, Zamolodchikov, 1987]. The Liouville action on a $\mathbb{Z}_N$ cone regularised by the Gauss log-sine product (the same regularisation that gives $b(N)$ in the Havelock decomposition, Paper I, §6) is $S_L \geq b(N) \ln N$. Therefore $C_k^2 \leq \exp(-b(N) \ln N)$, and the total non-identity contribution is

$$\frac{\delta F}{F} \leq (N-1) \exp(-b(N) \ln N)$$

(verified numerically for $N = 3, \dots, 16$; see Code Availability): 11% at $N = 5$, 1.7% at $N = 6$, 0.14% at $N = 7$, $\sim 10^{-10}$ for $N = 11$.

Bound (1) excludes competing saddles at the eigenvalue level; bound (2) controls the non-perturbative twist contributions to the conformal block decomposition. Together they establish identity block dominance for $N \geq 6$ (corrections $\leq 2\%$). At $N = 5$, the 11% correction is the leading systematic uncertainty; however, since $\lambda_{\min}(5) = 1$ on the flat plane, a shift of ~ 0.11 leaves $\lambda_{\min} > 0$, so the stability conclusion is not affected. $\square$

Theorem 6.2 identifies λ_m as the m -th $\mathbb{Z}_N$ -channel stiffness of the semiclassical N -point function. No bulk computation is needed: the stability condition $\lambda_m > 0$ is *the confining potential exceeding the holomorphic OPE repulsion in channel m* , a statement in the boundary CFT. The palindromic thresholds $\xi^*(N)$ —at which $C_1(\xi^*) = f_{\max}(N)$ —are visible as algebraic-number bulk positions where the $\mathbb{Z}_N$ conformal block changes sign.

Quantum correction. The $\mathbb{Z}_N$ Fourier decomposition of the Virasoro block at central charge c confirms the identification at all orders. The m -dependent part of $\log F^{(m)}$ is

$$\left[\frac{m(N-m)}{2} - \frac{m^2(N-m)^2}{4c} + O(1/c^2) \right] \log q. \quad (6.4)$$

At $c = \infty$: the leading term is the Havelock Casimir (exact, verified for all $N = 4, \dots, 12$ and all modes). At finite c : the correction $\delta h_m = -f_m^2/c$ is a *polynomial in the Casimir*, so the algebraic structure (palindromic symmetry, number field) is preserved at each order in $1/c$. The fractional correction $|f_m/c| \leq 1/8$ (maximised at $m = N/2$, $c = N^2$) does not change any eigenvalue sign (Proposition 8.1).

7 The polygon–BTZ phase transition

The polygon framework *must* reproduce the known Hawking–Page phase transition of AdS_3 gravity [Witten, 1998]; this section verifies that it does. The Euclidean action comparison between the polygon and the BTZ black hole makes the Hawking–Page language precise and gives a novel N -dependent correction $T_c(N) < T_{\text{HP}}$ from the polygon binding energy.

The Euclidean actions. At inverse temperature β , the regularised Euclidean actions of the competing configurations are:

$$I_{\text{polygon}}(\beta) = -\frac{\beta}{8G} + \beta E_{\text{polygon}}, \quad (7.1)$$

$$I_{\text{BTZ}}(\beta) = -\frac{\pi^2 \ell^2}{2G\beta} \quad (M = 0 \text{ threshold BTZ}), \quad (7.2)$$

where $E_{\text{polygon}} = Nm(1-2Gm) + V_{\text{int}}$ is the total energy of the N -gon in the small-deficit regime $2Gm \ll 1$ (valid for $Gm = 3m/(2c) \leq 3/(2 \times 51.6) = 0.03$; the formula receives $O((Gm)^2) \approx 0.1\%$ corrections), and $V_{\text{int}} < 0$ is the pairwise Havelock interaction on $\mathbf{H}^2$. The vacuum action $-\beta/(8G)$ is the standard thermal-AdS₃ contribution.

The transition temperature. Setting $I_{\text{polygon}} = I_{\text{BTZ}}$:

$$T_c(N, m, \xi) = \frac{1}{2\pi\ell} \sqrt{1 - 8G E_{\text{polygon}}(N, m, \xi)}, \quad (7.3)$$

which reduces to the standard Hawking–Page temperature $T_{\text{HP}} = 1/(2\pi\ell)$ for $E = 0$. For $T < T_c$: the polygon has lower action ($\Delta I < 0$, polygon preferred). For $T > T_c$: the BTZ has lower action. The total polygon energy $E_{\text{polygon}} = Nm_{\text{vert}} + V_{\text{int}}$, where the interaction energy on the unit circle is $V_{\text{int}} = -\sum_{p<q} \log |z_p - z_q| = -(N/2) \ln N$ (from $\prod_{p=1}^{N-1} 2 \sin(\pi p/N) = N$), lowers T_c below T_{HP} : the binding energy $V_{\text{int}} < 0$ deepens the polygon potential well.

The negative mode. For $N > N_{\text{crit}}(\xi)$: the polygon has a negative Havelock eigenvalue $\lambda_m < 0$ at mode $m = \lfloor N/2 \rfloor$ (§3). By the vortex-gravity isomorphism (Paper II, Proposition 13.1), the Havelock Hessian of the polygon is the second variation of the gravitational energy around the multi-cone saddle. A negative Havelock eigenvalue therefore is a negative mode of the Lichnerowicz operator Δ_{Lich} on the multi-cone geometry—a perturbation that lowers the Euclidean gravitational action, signaling decay [Giombi et al., 2008]. The mode $m = \lfloor N/2 \rfloor$ displaces adjacent masses alternately inward and outward, a pattern that drives pairwise coalescence. In the nonlinear regime, the coalescing masses form a single body exceeding the BTZ threshold $8GM > 1$, completing the decay to a BTZ black hole.

The three-region phase diagram. The polygon–BTZ competition has three distinct regimes:

1. *Stable polygon, preferred* ($N \leq N_{\text{crit}}, T < T_c$): all Havelock eigenvalues are positive and the polygon has lower Euclidean action than the BTZ. The polygon contributes as a stable saddle to the path integral.
2. *Stable polygon, metastable* ($N \leq N_{\text{crit}}, T > T_c$): all eigenvalues are positive but the BTZ has lower action. The polygon is a local minimum separated from the BTZ by a barrier; the tunnelling rate is $\Gamma \sim e^{-S_{\text{BTZ}}}$.
3. *Unstable polygon, decay* ($N > N_{\text{crit}}$): the mode $m = \lfloor N/2 \rfloor$ is a negative mode, and the polygon decays continuously (no barrier) toward the BTZ. The palindromic threshold $\xi^*(N)$ marks the boundary between regions 1/2 and region 3.

The critical exponent. At the palindromic threshold $\rho = \rho^*(N)$, the critical mode $m^* = \lfloor N/2 \rfloor$ vanishes linearly: $\lambda_{m^*}(\rho) \sim \coth(\rho^*) \cdot (\rho - \rho^*)$. The one-loop partition function (Gaussian integration over the zero mode) gives the universal critical exponent $\nu = 1/2$:

$$Z(\rho) \sim \frac{Z_{\text{frozen}}}{|\rho - \rho^*|^{1/2}}, \quad Z_{\text{frozen}} = \frac{2^{3(N-2)/4}}{(N-3)!!}$$

where $(N-3)!! = 1 \cdot 3 \cdot 5 \cdots (N-3)$ is the double factorial (even N , one zero mode at $m^* = N/2$; for odd N , the palindromic pair m^* and $N-m^*$ both have $\lambda = 0$, giving two zero modes and $Z \sim |\rho - \rho^*|^{-1}$. Z_{frozen} for odd N is $\prod_{m \neq m^*, m \neq N-m^*} [f(m^*, N) - f(m, N)]^{-1/2}$, computed from the Casimir gaps $f(m^*, N) - f(m, N) = [m^*(N-m^*) - m(N-m)]/2$. The closed-form expression includes the normalisation of the first-order Kirchhoff path integral measure; it is verified numerically against the gap-by-gap product for all $N \leq 20$ (see Code Availability). The polygon–BTZ transition is second-order with mean-field exponents.

The entropy gap. In the microcanonical ensemble (fixed total energy), the BTZ *always* has higher entropy: $\Delta S = S_{\text{BTZ}}(E_{\text{polygon}}) > 0$. The polygon is therefore always entropically disfavored; its survival as a metastable state depends entirely on the positivity of the Havelock eigenvalues (the energy barrier), not on entropy. This makes N_{crit} the sharp dividing line: below it, the polygon is a long-lived metastable state; above it, the polygon decays on a dynamical timescale $\tau \sim 1/\sqrt{|\lambda_{\lfloor N/2 \rfloor}|}$.

Microscopic entropy accounting. The BTZ entropy decomposes into discrete and continuous contributions:

$$S_{\text{BTZ}}(N) = \underbrace{S_{1\text{-loop}}(N)}_{\text{continuous}} + \underbrace{\Delta S_{\text{wc}}}_{\text{discrete}}, \quad (7.4)$$

where $S_{1\text{-loop}} = \log(N-3)!! - \frac{3(N-2)}{4} \log 2$ (from the Casimir gaps) and ΔS_{wc} is the wall-crossing contribution. At each palindromic threshold, the $\mathbb{Z}/2$ Galois action ($m \leftrightarrow N-m$) contributes:

N parity	Critical mode	ΔS
Even (self-paired $m = N/2$)	one eigenvalue crosses	$\log 2$ (1 bit)
Odd (paired $m, N-m$)	two eigenvalues cross	$2 \log 2$ (2 bits)

From $N = 7$ (the flat-plane marginal threshold) through $N = 12$: exactly 9 bits (verified numerically). Through $N = 17$ (11 wall-crossings): the discrete part contributes $17 \ln 2 \approx 11.8$ nats (9.9%), the continuous one-loop part 107.8 nats (90.1%). The L^2 -Morse index $\tau(P_-)$ has denominator 2 at each threshold (from the eta invariant shift of $1/2$ due to the zero mode), consistent with the K -theoretic framework.

The one-loop correction. The one-loop partition function around the polygon saddle is $Z_{1\text{-loop}} = (\det' \Delta_{\text{Lich}})^{-1/2}$, where Δ_{Lich} is the Lichnerowicz operator on the multi-cone geometry. By the conformal invariance of the flat-plane backreaction (Proposition 13.2), the leading correction comes from the spatial topology: the spectral bound $|\delta C_1| \leq 2.86/(4\pi\lambda_1)$ (where λ_1 is the spectral gap of the spatial Laplacian) gives a rigorous bound on the quantum shift of the stability eigenvalues. On the Bolza surface

($\lambda_1 \approx 3.84$): $|\delta C_1| < 0.06$, which does not shift N_{crit} away from 7 on the flat plane. On the modular surface ($\lambda_1 \approx 91$): $|\delta C_1| \lesssim 0.002$, negligible. The Bolza surface carries the $\{8, 3\}$ hyperbolic lattice on which [Lenggenhager et al. \[2023\]](#) constructed the hyperbolic Haldane model (Chern insulator with quantised Bott index $B = \pm 1$ and chiral edge states). The vortex stability analysis on the same $\{8, 3\}$ lattice shares the spectral data (λ_n of Δ_S) with the hyperbolic Haldane model. The bound $|\delta C_1| < 0.06$ shows that polygon stability is insensitive to the hyperbolic lattice topology. The Lorentzian continuation of the polygon–BTZ phase structure, including the analytic continuation of the Green’s function, the van Hove spectral singularity at the threshold, and the connection to GUE spectral statistics, is developed in Supplement S, §3.

8 The Wheeler–DeWitt equation on the breathing mode

The single degree of freedom ρ (geodesic radius of the polygon) and the potential $V(\rho) = \lambda_{m^*}(\rho)$ (the critical Havelock eigenvalue) define a minisuperspace Wheeler–DeWitt equation:

$$\left[-\frac{1}{2c} \frac{d^2}{d\rho^2} + \Phi(\rho) - \frac{m(N-m)}{2} \right] \Psi = 0, \quad \Phi(\rho) := \log(2 \sinh \rho) + b(N), \quad (8.1)$$

with $c = 12 b(N)$ and $\hbar = 1$ (the semiclassical limit is $c \rightarrow \infty$). On constant-curvature $\mathbf{H}^2$ the Havelock remainder $\delta_m = 0$ identically (§9), so it does not appear.

Notation. We write $\Phi(\rho) = \log(2 \sinh \rho) + b(N)$ for the Green’s function offset (the unconstrained radial potential), reserving $C_1(\xi)$ for the constrained angular stiffness (eq. 5.2). The WDW potential is $V(\rho) = \Phi(\rho) - f(m^*, N)$.

Logical status of c (1.1). The value of $b(N)$ (§2) is computed analytically in Paper I (Proposition 9.2); no free parameter is involved.

Why second-order dynamics from a first-order system. The Kirchhoff equations are first-order, but the WDW equation for the breathing mode ρ is second-order. The derivation has three steps:

1. *No-go for perturbative kinetic terms.* For a first-order operator $A_m = \partial_\tau + \lambda_m(\rho(\tau))$ on the thermal circle $[0, \beta]$, the functional determinant is $\det A_m = 1 - \exp(-\int_0^\beta \lambda_m d\tau)$, which depends only on the time-integral of λ_m , not on $\dot{\rho}$. Equivalently, the Matsubara influence kernel $\sum_n [(i\omega_n + \lambda)(i\omega_{n+k} + \lambda)]^{-1}$ vanishes for $k \neq 0$: partial fractions give $(1/i\nu_k) [\sum_n (i\omega_n + \lambda)^{-1} - \sum_n (i\omega_{n+k} + \lambda)^{-1}]$, and each factor sum equals $(\beta/2) \coth(\beta\lambda/2)$ by the Bose summation formula ($n \mapsto n+k$ is a relabelling of $\mathbb{Z}$). The integrable bath therefore contributes *no* $\dot{\rho}^2$ term to the effective action; the kinetic coefficient is not perturbative.
2. *Geometric origin (Fisher–Rao).* The kinetic term arises from the path-integral *measure*, not the action. The Feynman–Kac correspondence $Z = \int \mathcal{D}\rho \, \mu[\rho] e^{-S[\rho]}$ carries the measure $\mu[\rho] = \Pi_\tau \sqrt{c/(2\pi d\tau)}$, whose coefficient c is the Fisher–Rao information metric on the family of Boltzmann distributions $P_m(\varepsilon | \rho) \propto e^{-\lambda_m(\rho)|\varepsilon|^2}$ parametrised by ρ :

$$g_{\text{FR}}(\rho) = \coth^2 \rho \sum_{m \neq m^*} \frac{1}{\lambda_m(\rho)^2}.$$

This metric is ρ -dependent. The *classical* BO separation is exact for any kinetic coefficient (the Havelock eigenvectors are ρ -independent Fourier modes). But the *quantum* BO factorisation requires c to be *constant*: if $c(\rho)$ were ρ -dependent, the FK Hamiltonian $H = -(1/(2c(\rho))) d^2/d\rho^2 + V(\rho)$ would produce a first-derivative term $-c'(\rho)\psi'/(2c(\rho)^2)$ that couples the breathing mode to ρ -gradients of the angular-mode bath, breaking the exact factorisation at the quantum level. Therefore c must be a constant. The unique constant matching Brown–Henneaux ($c = 3\ell/(2G)$) is $c = 12b(N)$. The Fisher–Rao metric $g_{\text{FR}}(\rho)$ is the ρ -dependent version; extracting the constant requires a tortoise-coordinate field redefinition $\rho \rightarrow \tilde{\rho}(\rho)$ with $g_{\text{FR}}|d\rho|^2 = c|d\tilde{\rho}|^2$. Under this redefinition, the WDW equation becomes $-(1/2c)d^2\psi/d\tilde{\rho}^2 + \tilde{V}(\tilde{\rho})\psi = E\psi$ with $\tilde{V} = V(\rho(\tilde{\rho})) + O(1/c)$ (Schwarzian correction from the non-linear coordinate change). The $O(1/c)$ Schwarzian term is absorbed into the one-loop correction (1.9% at $N = 7$; see step 3 below), so the leading-order potential retains the form $V(\rho)$ of Proposition 3.7.

3. *Exact coefficient (two independent routes)*. Two independent calculations determine c :

- (a) *Eigenvalue arithmetic* (definition). $b(N)$ is the unique mode-independent offset in $\lambda_m = C_1(\rho) - f(m, N)$ (Proposition 9.2). Setting $c_1 = 12b(N)$ defines the central charge from the Havelock spectrum.
- (b) *Chern–Simons gauge theory* (independent calculation). The CS level $k = \ell/(4G)$ gives $c_2 = 6k = 3\ell/(2G)$ (Brown–Henneaux, 1986). The CMS–CS Casimir identification (Proposition 2.1) matches the Havelock spectrum to the CS Wilson-line spectrum, forcing $k = 2b(N)$ and hence $c_2 = 12b(N) = c_1$.

A third route—the Zamolodchikov c -theorem flow from $c_{\text{UV}} = 12N^2$ to $c_{\text{IR}} = 12b(N)$ (Theorem 10.3)—provides a consistency check but shares the CMS–CS input with route (b). The WDW kinetic coefficient equals $c_1 = c_2 = 12b(N)$; any other value would make the FK ground state disagree with the boundary CFT vacuum, contradicting the CS/WZW correspondence (Witten 1988). The CMS integrability makes the one-loop partition function $Z_{\text{bath}}(\rho) = \prod_{m \neq m^*} |\lambda_m(\rho)|^{-1/2}$ exact at $O(c^0)$, with anharmonic corrections of order $1/c$ (1.9% at $N = 7$, 0.8% at $N = 11$, decreasing as N^{-2} ; see Code Availability).

The no-go result (step 1) shows the kinetic term is genuinely geometric: it encodes the information-theoretic distinguishability of neighbouring equilibria. The four-fold matching (step 3) fixes $c = 12b(N)$ without free parameters.

The four-role consistency of c (1.1) (Section 1) is preserved. Each term has a definite origin:

Term	Expression	Origin
Quantum kinetic	$-\hbar^2/(2c)\Psi''$	Todd offset $c = 12b(N)$ (§1)
Classical potential	$\log(2 \sinh \rho)$	$\mathbf{H}^2$ Green’s function
Vacuum energy	$b(N) \sim c/12$	Orbifold CFT vacuum (Todd class)
Matter	$m(N-m)/2$	Twist-field dimension (first Chern class)
One-loop	δ_m	Havelock remainder (traceless, palindromic)

Three regimes.

1. **Classical** ($c \rightarrow \infty$): the kinetic term vanishes; the constraint reduces to $V(\rho) = 0$, i.e., $\lambda_{m^*}(\rho^*) = 0$ —the palindromic threshold, which is the Einstein equation of the vortex-gravity system.
2. **Semiclassical** (WKB): $\Psi \sim \exp(\pm(i/\hbar)\int\sqrt{2cV}d\rho)$, oscillatory in the exterior ($V > 0$, sinh-branch) and evanescent in the interior ($V < 0$, cosh-branch). The continuation $\rho \rightarrow \rho + i\pi/2$ maps between them (§7, Lorentzian paragraph).
3. **Quantum** (near threshold): $\Psi \sim \text{Ai}\left((2c \coth \rho^*)^{1/3}(\rho - \rho^*)\right)$, smoothing the classical horizon over the width

$$\delta\rho = \frac{1}{(2c \coth \rho^*)^{1/3}}, \quad (8.2)$$

where $\coth \rho^* = V'(\rho^*)$ is the surface gravity. For large ρ^* ($\coth \rho^* \rightarrow 1$): $\delta\rho \rightarrow (2c)^{-1/3} \approx (2N^2)^{-1/3}$. At $c = \infty$: $\delta\rho \rightarrow 0$ (sharp classical threshold). At finite c : the horizon has quantum width $O(c^{-1/3})$.

Proposition 8.1 (Eigenvalue non-crossing at finite c). *Since $c = 12b(N) = N^2 + N - 12\ln 2 + O(\ln N/N) > N^2$ for all $N \geq 3$, the condition $c \geq N^2$ is automatic. For $N \geq 4$ (so that $\lfloor N/2 \rfloor \geq 2$), the mode-dependent part of the $O(1/c)$ Virasoro block correction satisfies*

$$\frac{|\delta f_m - \delta f_{m+1}|}{f(m+1, N) - f(m, N)} = \frac{2m(N-m-1) + (N-1)}{2c} \leq \frac{1}{4} - \frac{1}{2N^2} < \frac{1}{4}. \quad (8.3)$$

In particular, no eigenvalue crossing occurs and the critical mode index $m_{\text{crit}} = \lfloor N/2 \rfloor$ is preserved.

Proof. The leading $1/c$ correction to the Casimir is $\delta f_m = -f_m^2/c = -m^2(N-m)^2/(4c)$. The difference factors as $\delta f_m - \delta f_{m+1} = (N-2m-1)[2m(N-m-1) + (N-1)]/(4c)$. Dividing by the Casimir gap $f(m+1) - f(m) = (N-2m-1)/2$ gives the ratio $R(m) = [2m(N-m-1) + (N-1)]/(2c)$. The maximum over $1 \leq m < N/2$ is at $m = N/2 - 1$, giving $R_{\text{max}} = (N^2-2)/(4N^2) = 1/4 - 1/(2N^2) < 1/4$. Since $R < 1$ for all adjacent pairs: the residual gap is at least $3(N-2m-1)/8 > 0$, and no reordering occurs. $\square$

Remark 8.1 (Algebraic field preservation). *The Zamolodchikov c -recursion gives the correction as a rational function of c , with only $\mathbb{Z}_N$ -invariant channels contributing. At $c = N^2$: the palindromic coefficients $A(c), B(c)$ remain in $\mathbb{Q}(\cos(2\pi/N))$ —the same number field as at $c = \infty$. The palindromic form $A\xi^2 + B\xi + A = 0$ is exact at all c (the conformal inversion $\xi \leftrightarrow 1/\xi$ is a geometric symmetry of $\mathbf{H}^2$, not a perturbative artefact).*

The palindromic threshold $\xi^*(N)$ is the critical AdS radius at which the polygon-BTZ transition occurs, and the algebraic number field $\mathbb{Q}(\sqrt{D})$ of the threshold is the arithmetic invariant of the phase boundary.

Remark 8.2 (Uniqueness of the WDW wave function). *The WDW equation (8.1) on $\rho \in (0, \infty)$ has two singular endpoints. Their Weyl classification determines whether boundary conditions are needed:*

1. At $\rho = 0$: limit-circle. *The potential $V(\rho) = \log(2 \sinh \rho) + b(N) - f(m^*, N) \rightarrow \log(2\rho) + \text{const}$ as $\rho \rightarrow 0$. Weyl's criterion (integral test): $\int_0^\varepsilon \rho |2cV(\rho)| d\rho < \infty$ because $\rho |\log \rho|$ is integrable. Both linearly independent solutions are L^2 near the origin; a boundary condition is required.*

2. At $\rho \rightarrow \infty$: limit-point. $V(\rho) \rightarrow \rho + \text{const}$ (linear growth), so $\int^\infty \rho \cdot 2c\rho d\rho = \infty$. Only the Airy-decaying solution is L^2 ; no boundary condition is needed.

The operator therefore requires exactly one boundary condition (at $\rho = 0$). The Born–Oppenheimer structure supplies it: the CMS integrability of the angular sector (Proposition 9.1) makes the adiabatic separation exact, and the angular modes in their ground state at each ρ select the ground state of the effective Hamiltonian. This ground state is non-degenerate: numerically, the spectral gap $\Delta E \approx 0.84$ is stable across $N = 7\text{--}15$ (see Code Availability). The wave function is therefore unique without imposing Hartle–Hawking or Vilenkin boundary conditions—the microscopic (Havelock) structure determines the cosmological state.

9 The three-layer decomposition

Proposition 9.1 (Orbifold–Havelock correspondence). *At leading order in $1/N$, the Havelock eigenvalue on $\mathbf{H}^2$ matches the $\mathbb{Z}_N$ orbifold CFT ground-state energy at $c_{\text{orb}} = 12N^2$ (the UV orbifold central charge) with angular impulse constraint:*

$$\lambda_m = -E_0(m; c_{\text{orb}}=12N^2) + \mu_L(\xi), \quad (9.1)$$

where $E_0(m) = -N/2 + m(N-m)/2$ is the twisted-sector ground state and $\mu_L(\xi) = C_1(\xi) - N/2$ is the Lagrange multiplier. The matching uses $c_{\text{orb}} = 12N^2$ (so that $c/(24N^2) = 1/2$), while the exact Havelock central charge is $c = 12b(N) = N(N+1) - 12\ln 2 + \dots \approx N^2$ at leading order. The distinction: c_{orb} is the UV (free-field) central charge of the $\mathbb{Z}_N$ orbifold; $c = 12b(N)$ is the exact IR value after regularisation of the vortex self-energy. The mode-dependent matching (the $m(N-m)/2$ term) uses only $c_{\text{orb}}/(24N^2) = 1/2$; the mode-independent $O(N)$ corrections (the $\ln 2$ and $(\ln N)/(N-1)$ terms in $b(N)$) enter through μ_L , not through E_0 .

Symbol	Value	Role
$c_{\text{orb}} = 12N^2$	UV orbifold	Mode-dependent matching
$c = 12b(N) \approx N^2 + N$	Exact (IR)	WDW kinetic coefficient
$c \approx N^2$	Leading order	Approximation in estimates

Proof. The orbifold ground-state energy $E_0(m) = -c/(24N) + cm(N-m)/(24N^2)$ at $c = c_{\text{orb}} = 12N^2$ (the UV orbifold central charge, which gives $c/(24N^2) = 1/2$ exactly) yields $E_0 = -N/2 + m(N-m)/2$. Then $-E_0(m) + \mu_L = [N/2 - m(N-m)/2] + [C_1(\xi) - N/2] = C_1(\xi) - m(N-m)/2 = \lambda_m$. $\square$

Proposition 9.2 (Todd class decomposition of the eigenvalue offset). *The $\mathbb{Z}_N$ -equivariant eigenvalue decomposition of the vortex Laplacian on the cusped Poincaré disk (analogous to the equivariant Riemann–Roch theorem) satisfies*

$$\text{ind}_m = \underbrace{\frac{m(N-m)}{2}}_{c_1\text{-term}} + \underbrace{\frac{b(N)}{2}}_{\text{Todd term}} + \delta_m, \quad (9.2)$$

where $b(N)$ (§2) is independent of both m and ρ . The central charge c (1.1) is the same for every mode m , extracted from the Todd correction alone.

Proof. The Havelock eigenvalue satisfies $\lambda_m + f_m = \Phi(\rho) + \delta_m + O(e^{-\rho})$, where $\Phi(\rho) = \log(2 \sinh \rho) + b(N)$ (the radial dependence is absorbed by the Green's function). Subtracting $\log(2 \sinh \rho)$: $\text{ind}_m := \lambda_m + f_m - \log(2 \sinh \rho) = b(N) + \delta_m$. Since δ_m is traceless ($\sum \delta_m = 0$ to 10^{-15}) and palindromic ($\delta_m = \delta_{N-m}$), the mode-independent part $b(N)$ is extracted exactly. $\square$

The decomposition (9.4) below is the spectral analogue of equivariant Riemann–Roch (f_m : twist-field dimension; $b(N)$: Todd correction; δ_m : one-loop remainder).

Exact Havelock remainder formula. On constant-curvature surfaces, $\delta_m = 0$ identically (Paper IV, Theorem 6.1); on compact quotients $\Gamma \backslash \mathbf{H}^2$, the Selberg trace formula gives a mode-independent shift $|\delta C_1| \leq 0.020$ on the Bolza surface (§7). The formula below is the flat-plane residual.

The anomaly δ_m decomposes exactly into character sums on $\mathbb{Z}/N\mathbb{Z}$:

$$\delta_m = L(m, N) - \frac{1}{4} S_2(m, N) + c(N), \quad (9.3)$$

where $L(m, N) = \sum_{p=1}^{N-1} \cot(\pi p/N) \cos(2\pi p m/N)$ is the cotangent sum, $S_2(m, N) = \sum_{p=1}^{N-1} \csc^2(\pi p/N) \cos(2\pi p m/N)$ is the cosecant-squared sum, and $c(N)$ is the unique constant enforcing $\sum_{m=1}^{N-1} \delta_m = 0$ (tracelessness). The coefficients of L and S_2 are 1 and $-1/4$ respectively, exact and universal ($R^2 = 1.000000$ for all tested $N = 4, \dots, 24$). This identifies the anomaly as the lattice correction (orbifold fixed-point contribution in the KK picture) to the continuum Casimir.

The central charge c (1.1) makes five independently testable predictions beyond the Casimir: (i) the exact offset $b(N)$ (Paper I); (ii) the mean aliasing $\langle D \rangle = N(N+1)/12$; (iii) the growth law $\rho^*/f_{\text{crit}} \rightarrow 1/3$; (iv) the quartic coupling $Q/f_{\text{crit}}^2 = N/48$ (even N); (v) $\Delta_{\mathbf{H}^2} h = -1/2$. All five follow from the B_2 -tower: the Casimir $m(N-m)/2$ is $O(c)$; the Havelock remainder δ_m is $O(c^0)$ (topological); the Virasoro block correction is $O(1/c) \approx 10^{-5}$ (§5).

9.1 Bihamiltonian hierarchy and arithmetic surfaces

The tower extends to a *Magri bihamiltonian hierarchy*. As noted in §2, the Thomson vortex system has the same linearised spectrum as the CMS model at coupling $g = 1$. The CMS model is bihamiltonian [Olshanetsky and Perelomov, 1981]: the Kirchoff bracket P_0 (mode-independent) and the Casimir bracket P_1 (eigenvalue f_m on mode m) are compatible, with recursion operator $R = P_1 P_0^{-1}$ acting as multiplication by $f_m = m(N-m)/2$. By the Magri theorem, the k -th conserved quantity I_k has mode eigenvalue f_m^k , with physical derivatives $d_{2k} = \alpha_k N f_m^k$ ($\alpha_1 = 1$, $\alpha_2 = 4$, $\alpha_3 = 64$, verified numerically for $N \leq 16$). The quartic coupling $Q = (N/48)f^2$ lies inside the hierarchy ($H_4^{\text{diag}} \propto I_2$), guaranteeing $\{I_j, H_4\} = 0$ by integrability. On arithmetic surfaces (Bolza, modular, etc.), the eigenvalue acquires a topological correction that is invisible to the confining potential $C_1(\xi)$:

$$\lambda_m = \underbrace{C_1(\xi)}_{\text{Ricci scalar}} - \underbrace{\frac{m(N-m)}{2}}_{\text{Casimir}} + \underbrace{\delta_m(N)}_{\text{Havelock remainder}}. \quad (9.4)$$

Each layer has a distinct character:

1. $C_1(\xi)$: continuous, surface-dependent, mode-independent. Scales approximately as $C_1 \propto \rho$ over the range $\rho \in [0.01, 10]$ (with $R^2 = 0.98$, where ρ is the geodesic radius and $\xi = \tanh^2(\rho/2) \in (0, 1)$); the exact formula is eq. (5.2).
2. $f(m, N) = m(N-m)/2$: discrete, universal (surface-independent). The $\mathbb{Z}_N$ Casimir from the Havelock identity.
3. $\delta_m(N)$: topological, UV-finite, ξ -independent (variation $< 10^{-13}$), and exactly traceless ($\sum_{m=1}^{N-1} \delta_m = 0$ to 10^{-15}). Carries all number-theoretic content: Hecke eigenvalues, palindromic structure, Langlands data.

By structural analogy with the decomposition of the Riemann curvature tensor: C_1 plays the role of the Ricci scalar, $f(m, N)$ the traceful part, and δ_m the traceless Weyl-like anomaly. In $\text{AdS}_3/\text{CFT}_2$, the Brown–Henneaux scaling $C_1 \propto \log(\xi)$ says that the Weyl-like anomaly δ_m is the topological correction that carries the arithmetic of the underlying surface, contributing corrections of order $\|\delta\|/C_1 \sim 4/\Delta_\gamma$, where $\Delta_\gamma = \text{tr}(\gamma)^2 - 4$ is the discriminant of the shortest closed geodesic (1.6% for $N = 8, T = 16$; 4.2% for $N = 9, T = 10$).

9.2 Anomaly structure

The Weyl norm scales as $\|\delta\|^2 \sim N^{4.2}(\log N)^{4.1}$ for $N \geq 40$ (not a pure power law). The triangle ($N = 3$) is *anomaly-free*: $\delta_m = 0$ exactly for all Green’s functions and curvatures, because $\cos(2\pi/3) = \cos(4\pi/3)$ forces $G_1 = G_2$. Under Casimir-conserving mergers $N_1 + N_2 \rightarrow N_3$ ($N_3^3 - N_3 = N_1^3 - N_1 + N_2^3 - N_2$): the unique anomaly-*destructive* reaction is $4+4 \rightarrow 5$ (88% Weyl loss); all other mergers create Weyl content (gains 46% to 849%). The Weyl entropy per degree of freedom $s(N) = \log \|\delta\|^2 / \sum f$ peaks at $N \approx 9$ –10, coinciding with the observed planetary polygon range ($N = 5, 6, 8$).

The anomaly δ_m has two independent origins, which the parameter space (Δ, ξ) separates cleanly:

1. *Interaction anomaly* ($\Delta > 0$ on $\mathbb{R}^2$): deviation of $V = r^{-2\Delta}$ from $-\log r$. Large ($\|\delta\|^2 = O(1\text{--}100)$) but arithmetically trivial.
2. *Geometric anomaly* ($\Delta = 0$ on $\mathbf{H}^2$): deviation of the hyperbolic Green’s function from the flat one. Small ($\|\delta\|^2 = O(0.01\text{--}100)$) but carries all number theory: palindromic polynomials, Hecke eigenvalues, Langlands data.

The RG flow ($\Delta \rightarrow 0$) strips away the large, trivial interaction anomaly, leaving behind the small geometric anomaly as the residue that survives at the conformal fixed point. The number theory is not emergent from the flow; it is a property of the conformal fixed point, activated by curvature. The three-layer structure has a symplectic origin: the Kirillov form ω_{KR} (mode-independent) and the Weil–Petersson form ω_{WP} (mode-dependent) are *not* proportional; their discrepancy $\omega_{\text{WP}} - \alpha \omega_{\text{KR}} \propto f(m, N)$ is the Casimir (Layer 2), and the residual beyond this linear relation is the Havelock remainder (Layer 3).

The B_2 organising principle — connecting the stability polynomial $6B_2(x) = P_2(2x-1)$ to Gauss–Legendre quadrature, the Euler–Maclaurin duality, and the csc^2 kernel universality (Calogero–Moser–Sutherland, Dyson, Haldane–Shastry, Laughlin, Selberg, W_N , AGT/Nekrasov) — is developed in Paper I, §6.6.

10 Higher-dimensional extension: the CFT₃/AdS₄ Casimir

The two-dimensional results extend to conformal field theory in any dimension d via the power-law interaction $V(x, y) = |x - y|^{-2\Delta}$, where Δ is the conformal dimension. For $\Delta \rightarrow 0$: $V \rightarrow -\log|x - y|$ (the 2D case). For $\Delta = 1$: the CFT₃ two-point function on S^2 (the boundary of AdS₄).

Proposition 10.1 (Generalised Havelock Casimir). *For N equally-spaced points on a ring on S^2 , interacting via $V = |x - y|^{-2\Delta}$, the angular Hessian eigenvalues factorise as $\mu_m = g(\theta_0, \Delta) \cdot h(m, N, \Delta)$, where the universal Casimir is*

$$h(m, N, \Delta) = \sum_{p=1}^{N-1} \frac{1 + 2\Delta \cos^2(\pi p/N)}{\sin^{2\Delta+2}(\pi p/N)} \left(1 - \cos \frac{2\pi mp}{N}\right). \quad (10.1)$$

For $\Delta = 0$: $h(m, N, 0) = 2m(N-m)$. The ratio $h(m)/h(1)$ is independent of θ_0 , confirming that the Havelock decomposition extends to all Δ .

The constrained Hessian on the flat plane with $V = |x - y|^{-2\Delta}$ gives an N -dependent stability boundary $N_{\text{crit}}(\Delta)$:

Δ	0	0.5	1	2	5	10
N_{crit}	7	5	5	5	4	4

The classical $N_{\text{crit}} = 7$ is an isolated phenomenon of the logarithmic interaction. For the CFT₃ value $\Delta = 1$: $N_{\text{crit}} = 5$.

Under the Green's function identification $\Delta = (d-2)/2$, the table acquires a dimensional interpretation:

Dimension d	2	3	4	5	≥ 6
$\Delta = (d-2)/2$	0	1/2	1	3/2	≥ 2
N_{crit}	7	5	5	5	4-5

The hexagon ($N = 6$, Saturn) is a *uniquely two-dimensional* phenomenon: it is marginally stable at $d = 2$ ($\lambda_{\text{min}} = 1/2$) but unstable for $d \geq 3$ ($\lambda_{\text{min}} = -0.02$ at $d = 3$; $N_{\text{crit}}(\Delta=0.5) = 5 < 6$; see Code Availability). The pentagon ($N = 5$) survives in all tested dimensions $d \leq 10$.

Proposition 10.2 (Soft Havelock decomposition in dimension d). *For the interaction $V = r^{-(d-2)}$ ($d \geq 3$) on the N -gon ring, the angular Casimir decomposes as*

$$\sigma_m^{(d)} = a(N, d) \frac{m(N-m)}{2} + R_m(N, d), \quad (10.2)$$

where $a(N, d)$ is the leading Casimir coefficient (determined by least-squares fit) and R_m is the residual. For all tested $N = 5, \dots, 20$ and $d = 3, 4, 5$: the residual satisfies $|R_m - R_{m'}| < a|f_m - f_{m'}|$ for every adjacent pair, so the mode ordering is preserved. The critical mode is always $m = \lfloor N/2 \rfloor$ (same as $d = 2$), and the stability threshold has the same palindromic form as the two-dimensional case, with coefficients that depend on d .

For $d = 2$ ($s = 0$): $a = 1$ exactly and $R_m = 0$ (the Havelock identity). This is the *unique* dimension where the Casimir is an exact polynomial—because $\csc^2(\pi p/N)$ has a closed discrete Fourier transform ($\sum \csc^2 \cdot (1 - \cos) = 2m(N-m)$), while $\csc^s$ for $s \neq 2$ does not. The exact algebraic structure of Paper I (palindromic polynomials, algebraic number fields, modular forms) is specific to $d = 2$, a consequence of the uniqueness of the $\csc^2$ weight function among power-law potentials. However, the *structural skeleton* survives in all dimensions:

1. **Mode ordering.** The critical mode is $m = \lfloor N/2 \rfloor$ in every d (Proposition 10.2). The residual R_m is too small to reorder eigenvalues.
2. **Palindromic symmetry.** The symmetry $\sigma_m = \sigma_{N-m}$ holds exactly for all d (from the discrete cosine: $\cos(2\pi pm/N) = \cos(2\pi p(N-m)/N)$). On curved spaces, conformal inversion $\xi \leftrightarrow 1/\xi$ is geometric (a property of the Poincaré ball, not of the interaction), so the threshold equation retains palindromic form.
3. **Finite- c algebraic structure.** The $O(1/c)$ Virasoro block correction $\delta h_m = -f_m^2/c$ is a *polynomial in the Casimir* (equation 6.4), preserving the algebraic structure at every order.
4. **Havelock remainder shape.** The normalised anomaly profile $\delta_m/\|\delta\|$ has inner products > 0.98 between $\Delta = 0$ and $\Delta = 2$ for $N \leq 6$: the number-theoretic fingerprint transfers to $d = 3$ with $< 2\%$ distortion.

The $d = 2$ case is where the algebra is exact; the $d \geq 3$ case is where it is *approximate but structurally preserved*. The mechanism by which the logarithmic structure survives in 3+1D is the soft Havelock decomposition: the Casimir $m(N-m)/2$ is the dominant term at every d , with the interaction-specific residual R_m as a controlled perturbation.

Remark 10.1 (Transcendence of the transition dimensions). *The critical value $\Delta^*(N)$ where N_{crit} drops below N is expected to be transcendental for $N \geq 5$. The threshold equation involves $\sin^{-2\Delta}(\pi p/N)$ for different p with $\mathbb{Q}$ -linearly independent logarithms. For $N = 5$: the threshold equation involves $\sin^{-2\Delta}(\pi/5)$, an algebraic base raised to the power Δ . If Δ^* is algebraic and irrational, the Gelfond–Schneider theorem makes the individual terms transcendental, and Baker’s theorem on linear forms in logarithms would be needed to show no cancellation occurs. A rigorous proof requires Baker-type estimates. This contrasts with the palindromic thresholds $\xi^*(N)$ on $\mathbb{H}^2$, which lie in explicit algebraic number fields.*

For $N = 5$, the pentagon identity $\sin(2\pi/5)/\sin(\pi/5) = \varphi$ (the golden ratio) gives the eigenvalue as $\lambda_2(\Delta) = P_0(\Delta) + P_1(\Delta)\varphi^{-(2\Delta+2)}$, where P_0, P_1 are algebraic functions of Δ . The transition occurs at $\varphi^{2\Delta^*+2} = P_1/(-P_0) \approx 25.8$, giving $\Delta^* \approx 2.378$.

Theorem 10.3 (RG monotonicity of the Casimir ratio). *For any $N \geq 4$ and $m_{\text{crit}} = \lfloor N/2 \rfloor$, the generalised Casimir ratio $R(\Delta) = h(m_{\text{crit}}, N, \Delta)/h(1, N, \Delta)$ is strictly increasing on $[0, \infty)$, with exact limits*

$$R(0) = \frac{m_{\text{crit}}(N-m_{\text{crit}})}{N-1}, \quad R(\infty) = \left[\frac{\sin(\pi m_{\text{crit}}/N)}{\sin(\pi/N)} \right]^2. \quad (10.3)$$

Proof. By the Chebyshev identity, the mode amplification factor $f(p) = (1 - \cos(2\pi m_{\text{crit}} p/N)) / (1 - \cos(2\pi p/N))$ achieves its global maximum at $p = 1$ (and its palindromic twin $p = N-1$):

Even $N = 2k$: $\sin(\pi k p/N) = \sin(\pi p/2) \in \{0, \pm 1\}$, so $f(p) = 1/\sin^2(\pi p/N)$ for odd p and $f(p) = 0$ for even p . Since $\sin(\pi p/N) \geq \sin(\pi/N)$ for $p \geq 2$: $f(p) \leq f(1)$.

Odd $N = 2k+1$: The identity $\sin(k\pi p/N) = \sin(\pi p/2 - \pi p/(2N))$ (using $k/N = 1/2 - 1/(2N)$) gives explicit formulas: $f(p) = 1/(4\sin^2(\pi p/(2N)))$ for odd p , and $f(p) = 1/(4\cos^2(\pi p/(2N)))$ for even p . Since $f(1) = 1/(4\sin^2(\pi/(2N)))$: for odd $p \geq 3$, $\sin(\pi p/(2N)) \geq \sin(\pi/(2N))$ (monotonicity of $\sin$ on $(0, \pi/2)$), so $f(p) \leq f(1)$; for even $p \leq N-1$, $\cos(\pi p/(2N)) \geq \sin(\pi/(2N))$ (since $\pi p/(2N) + \pi/(2N) \leq \pi/2$ iff $p \leq N-1$), so $f(p) \leq f(1)$.

The normalised weight $\mu_p(\Delta) \propto \tilde{w}(p, \Delta) c_1(p)$ concentrates on $\{1, N-1\}$ as $\Delta \rightarrow \infty$, because $\tilde{w}(p)/\tilde{w}(1) = O([\sin(\pi/N)/\sin(\pi p/N)]^{2\Delta}) \rightarrow 0$ for $p \notin \{1, N-1\}$. Since $R(\Delta) = \mathbb{E}_\mu[f]$ and μ concentrates at the global maximum of f : $dR/d\Delta = \text{Cov}_\mu(d \log \tilde{w}/d\Delta, f) > 0$ (both f and $g = d \log \tilde{w}/d\Delta = -2 \log \sin(\pi p/N)$ are maximised at $p = 1$ and monotonically ordered on $\{1, \dots, \lfloor N/2 \rfloor\}$; the Chebyshev sum inequality gives $\text{Cov}_\mu(g, f) > 0$). $\square$

Corollary 10.4 (Stability monotonicity). *$N_{\text{crit}}(\Delta)$ is non-increasing in Δ : once a polygon type loses stability, it never regains it. The quantity $\log N_{\text{crit}}(\Delta)$ decreases from $\log 7$ ($\Delta = 0$, the Havelock regime) to $\log 3$ ($\Delta \rightarrow \infty$, only triangles survive). This is reminiscent of the Zamolodchikov c -theorem, though N_{crit} is a step function (not smooth) and Δ is not an RG parameter.*

11 Triangle universality and dynamical triangulation

The equilateral triangle ($N = 3$) is the *unique* regular polygon that is dynamically stable for all interaction ranges and quantum deformation levels.

Theorem 11.1 (Triangle universal stability). *The equilateral triangle ($N = 3$) is the unique regular polygon that is stable on the flat plane for all $\Delta \geq 0$, all TQFT levels $k \geq 3$, and whose tessellation $\{3, 6\}$ is the unique stable regular lattice for $\Delta \geq 1$.*

Proof sketch. Three independent arguments:

1. *Equal chords.* The triangle is the unique regular polygon where all chord distances are equal ($\sin(\pi/3) = \sin(2\pi/3) = \sqrt{3}/2$). No range-dependent competition between nearest and next-nearest neighbors $\implies$ stability is Δ -independent.
2. *Single mode.* The triangle has one nontrivial $\mathbb{Z}_3$ mode ($m = 1$, with $m = 2$ its palindromic conjugate). The Havelock eigenvalue $\lambda_1 = (N-1) - 1 = 1 > 0$. Mode competition requires at least two independent modes ($N \geq 4$).
3. *Zero Havelock remainder.* The anomaly δ_m is traceless: $\sum \delta_m = 0$. With one mode: $\delta_1 = 0$ identically. Verified: anomaly norm $= 2.2 \times 10^{-16}$ (machine zero) on the square torus.

For the tessellation claim, the three regular 2D lattices are triangular (coordination 6), square (4), hexagonal (3).

Square lattice: permanently unstable. At the X -point $\mathbf{k} = (\pi/a, 0)$, only neighbors along the x -axis contribute ($1 - \cos(\pi) = 2$; perpendicular neighbors give $1 - \cos(0) = 0$). The transverse phonon eigenvalue is $d_{yy}(X) = 4V'(a)/a = -8\Delta a^{-2\Delta-2} < 0$ for all $\Delta > 0$.

Hexagonal lattice: permanently unstable. The two-atom basis admits an optical phonon at the K -point whose eigenvalue is $d_{\text{opt}}(K) = 3V'(a)/a < 0$ (the sublattices slide against each other with no longitudinal restoring force from the three coordination bonds).

Triangular lattice: stable for $\Delta \geq 1$. At the K -point, all six nearest neighbors have $1 - \cos(\mathbf{k} \cdot \mathbf{R}) = 3/2$; the C_6 symmetry gives an isotropic dynamical matrix $D(K) = 18\Delta^2 a^{-2\Delta-2} \mathbf{I} > 0$. Further shells contribute isotropically at K (C_6 forces $\sum_{\text{shell}} \hat{R} \otimes \hat{R} = (n/2)\mathbf{I}$), each with coefficient $4\Delta^2 R^{-2\Delta-2} > 0$. At the M -point and general $\mathbf{k}$, positive definiteness for $\Delta \geq 1$ is confirmed by convergent lattice sums (see Code Availability); the critical exponent for lattice stability is $\Delta_{\text{crit}} \approx 0.85$. The triangular lattice is the *only* regular lattice that is ever stable. $\square$

Remark 11.1 (Dynamical triangulation). *In causal dynamical triangulations (CDT) and Regge calculus, spacetime is discretised into simplices (triangles in 2D). This triangulation is imposed by hand as a regularisation. Theorem 11.1 suggests it may be imposed by physics: among regular lattices, the triangular lattice is the unique stable configuration, so the gravitational path integral $Z = \sum_{\text{tilings}} e^{-S}$ may be dominated by triangulated configurations—the square and hexagonal alternatives are saddle points with negative phonon modes at all Δ . The extension to irregular triangulations (as used in CDT) remains open.*

12 The polygon partition function

The one-loop partition function of the vortex system organises naturally by algebraic number field.

Proposition 12.1 (Polygon partition function). *The saddle-point partition function for the vortex-polygon system (a discrete sum over polygon number N , each contributing a single conical geometry—not a path integral over metrics) on $\mathbf{H}^2$ is*

$$Z_{\text{grav}} = \sum_{N=3}^{\infty} e^{-b(N)} Z_{\text{frozen}}(N), \quad (12.1)$$

where $b(N)$ (§2) is the gravitational action and $Z_{\text{frozen}} = 2^{3(N-2)/4}/(N-3)!!$ is the one-loop determinant (even N ; the odd- N case is computed from Casimir gaps). The series converges super-exponentially: $e^{-b(N)} \sim e^{-N^2/12}$. Numerically, $Z_{\text{grav}} = 1.3304$ and $N = 3, \dots, 8$ capture 99.99% of the total.

Proof. Convergence: $e^{-b(N)} = O(e^{-N^2/12})$ decays faster than any polynomial, while $Z_{\text{frozen}} = O(N^{O(N)})$ grows sub-exponentially; the product vanishes. The numerical value $Z = 1.3304$ is computed by direct summation to $N = 20$ (the tail $N > 20$ contributes $< 10^{-20}$). $\square$

The partition function decomposes by *trace field*—the algebraic number field generated by $\cos(2\pi/N)$:

Trace field	N values	Fraction of Z
$\mathbb{Q}$ (rational)	3, 4, 6	90.98%
$\mathbb{Q}(\sqrt{5})$ (golden)	5, 10	8.26%
$\mathbb{Q}(\sqrt{2})$ (silver)	8	0.42%
$\mathbb{Q}(\cos 2\pi/7)$ (cubic)	7	0.34%
All higher fields	≥ 9	$< 0.01\%$

The rational field dominates; the golden ratio field gives the leading correction; the entire algebraic tower is exponentially suppressed by $e^{-N^2/12}$.

The Baum-Connes assembly map $\mu: K_0^{\mathbb{Z}_N}(\mathbf{H}^2) \rightarrow K_0(C_r^*(\mathbb{Z}_N))$ sends the equivariant index $\text{ind}_m = f(m, N) + b(N)$ to the K-theory class with trace $c = 12b(N)$. (The three Kasparov module conditions for the bounded triple $\alpha_N = (E, \pi, F_\xi)$ — self-adjointness, $F^2 - 1$ compact, and $[\pi(a), F]$ compact — are verified explicitly in Paper I, §A.4.) This is mode-independent (the same c for all m) and approximately integer: $c = N^2 + O(N)$ for all tested $N = 3, \dots, 24$. The discrepancy $c - N^2 \approx N + O(1)$ is the boundary self-energy ($N/12$) plus regularisation terms.

13 The Onsager–Wheeler–DeWitt correspondence

The Feynman–Kac correspondence of §8 identifies the BO effective action’s Euclidean path integral with the WDW equation — a mathematical identity showing the polygon system’s statistical mechanics *is* a quantum system. We now extract the consequences.

Proposition 13.1 (Self-decoherence of the breathing mode). *The $N-2$ angular Havelock modes decohere the breathing mode ρ with decoherence rate*

$$\gamma = \frac{[\Phi'(\rho)]^2}{4} \sum_{m \neq m^*} \frac{1}{\lambda_m^2}. \quad (13.1)$$

The reduced density matrix after tracing out the bath is

$$\varrho_{\text{red}}(\rho, \rho') = \Psi_0(\rho) \Psi_0^*(\rho') e^{-\gamma(\rho - \rho')^2}.$$

For $\gamma \gg 1$ the off-diagonal elements are exponentially suppressed and ϱ_{red} is effectively diagonal: $\varrho_{\text{red}} \approx |\Psi_0(\rho)|^2 \delta(\rho - \rho')$.

Proof. The CL influence functional for the Gaussian bath with ρ -dependent frequencies $\omega_m(\rho) = \sqrt{\lambda_m(\rho)}$ gives a decoherence kernel $\exp(-\gamma(\rho - \rho')^2)$ with $\gamma = \sum_m (g_m/\omega_m)^2$, where $g_m = \partial\omega_m/\partial\rho = \Phi'/(2\sqrt{\lambda_m})$ is the system-bath coupling. Substituting: $\gamma = (\Phi')^2/4 \times \sum 1/\lambda_m^2$. For $N = 7$ (all modes stable on $\mathbf{H}^2$): $\gamma \approx 7-12$ across the stable range. For $N = 11$ at the stability threshold $\rho_{11}^* \approx 0.88$ (the cosmologically relevant operating point, where all $N-2 = 9$ bath modes have $\lambda_m > 0$): $\gamma \gg 1$. Near the threshold, the critical mode m^* is excluded from the bath (it is the “system” variable). For even N , $m^* = N/2$ is self-palindromic ($N - m^* = m^*$), so only one mode is excluded. For odd N , the palindromic pair $\{m^*, N - m^*\}$ share the eigenvalue λ_{m^*} ; excluding only m^* leaves the partner $N - m^*$ in the bath with $\lambda_{N - m^*} \rightarrow 0$, so $1/\lambda^2 \rightarrow \infty$ and γ grows without bound as the threshold is approached. The polygon is therefore *strongly decohered* at the physical operating point: the soft bath mode near the stability threshold drives $\gamma \gg 1$ structurally.

Irreversibility from near-integrability. The CMS angular bath is integrable *in isolation*. The coupling to the breathing mode ρ through the ρ -dependent frequencies $\lambda_m(\rho)$ breaks exact integrability at $O(1/c)$ (the W_N quartic coupling, $\sim 2\%$ at $N = 7$). This converts exact Poincaré recurrences (which would periodically restore coherence) into diffusive mixing. The Nekhoroshev theorem bounds the mixing timescale from below: $\tau_{\text{mix}} \geq \exp(c_0 c^a)$ for some $c_0 > 0$, $a > 0$ (with $\varepsilon = O(1/c)$ as the perturbation parameter). Even this exponentially long mixing time satisfies $\tau_{\text{mix}} \ll \tau_{\text{tunnel}} \sim e^{S_{\text{BO}}}$, provided $c_0 c^a < S_{\text{BO}}$.

(for $c \approx 50\text{--}130$ and $S_{\text{BO}} \approx 38\text{--}103$, this is satisfied for generic Nekhoroshev exponents $a \leq 1/2$). A second bound comes from the Poincaré recurrence time of the full $(N-2)$ -mode system, $\tau_{\text{rec}} \sim \exp(O(N-2)) \sim e^5 - e^9 \approx 150\text{--}8000$, again $\ll \tau_{\text{tunnel}}$. The decoherence is therefore *irreversible* on all cosmologically relevant timescales.

This is consistent with the exact BO result (Paper V, eq. (5.1)): the mode-independence of Φ makes the BO potential exact (no non-adiabatic transitions), while the $O(1/c)$ W_N quartic—which breaks the exact CMS integrability—provides the mode mixing needed for irreversible decoherence. The two results operate at *different orders in* $1/c$ and are not contradictory: at $O(c^0)$, the CMS integrability gives exact BO separation (Havelock eigenvectors are ρ -independent, the angular modes decouple completely); at $O(1/c)$, the W_N quartic coupling breaks exact integrability and provides the system-bath coupling needed for decoherence. The decoherence is a perturbative $O(1/c)$ correction to an exactly decoupled $O(c^0)$ system—the standard situation in atomic/molecular BO theory, where non-adiabatic corrections at $O(m_e/m_p)$ allow energy transfer between electronic and nuclear sectors without contradicting the BO separation. $\square$

Implications. The CL decoherence selects the ρ -basis as the pointer basis ($\gamma \gg 1$ at the physical operating point), with $|\Psi_0(\rho)|^2 = e^{-\beta V}/Z$ as the Boltzmann weight [Zurek, 2003]. The polygon decoheres *itself*: no external observer or environment is required. The $N-2$ angular modes form a Caldeira–Leggett bath for the breathing mode ρ . The four CL conditions hold: (i) harmonic bath (Havelock diagonalisation, exact); (ii) uniform parametric coupling ($\partial\lambda_m/\partial\rho = \Phi'(\rho)$ for all m); (iii) timescale separation ($\tau_{\text{tunnel}} \sim e^{S_{\text{BO}}} \gg \tau_m \sim O(1)$); (iv) FDT (automatic for Gaussian systems). The resulting Onsager–Machlup action is $S_{\text{OM}} = \int [\frac{c}{2}\dot{\rho}^2 + V(\rho)] d\tau$, where $\dot{\rho}^2$ is the noise kernel (not a kinetic energy).

The Markovian approximation requires $\tau_{\text{bath}} \ll \tau_{\text{tunnel}}$. The bath correlation time is $\tau_{\text{bath}} = 1/\omega_{\min} = 1/\sqrt{\lambda_{\min}^{(\text{bath})}}$ where $\lambda_{\min}^{(\text{bath})}$ is the smallest *positive* eigenvalue among the bath modes (excluding the marginal mode(s) m^* and, for odd N , its palindromic partner $N-m^*$). At the palindromic threshold $C_1 = f(m^*, N)$: $\lambda_{\min}^{(\text{bath})} = f(m^*, N) - f(m^*-1, N) = (N-2m^*+1)/2$, which equals 1 (odd N) or $1/2$ (even N). Thus $\tau_{\text{bath}} \leq \sqrt{2}$. The tunnelling time $\tau_{\text{tunnel}} \sim e^{S_{\text{BO}}}$ with $S_{\text{BO}} \gg 1$ (Paper V), so the separation $\tau_{\text{bath}}/\tau_{\text{tunnel}} = O(e^{-S_{\text{BO}}})$ is exponentially sharp, and non-Markovian corrections are $O(e^{-S_{\text{BO}}})$; the coefficient $c = 12b(N)$ is classical (Proposition 9.2). The Feynman–Kac transfer matrix gives $H\Psi_0 = 0$ (the WDW equation, Proposition 3.7), with Born rule $|\Psi_0|^2 = e^{-\beta V}/Z$ as a mathematical consequence:

Classical (Onsager)	Quantum (WDW)
Boltzmann weight $e^{-\beta V}$	$ \Psi_0 ^2$
Partition function Z	$\langle \Psi_0 \Psi_0 \rangle$
Thermal correlation $\langle \rho(\tau_1) \rho(\tau_2) \rangle$	$\langle \Psi_0 \hat{\rho} e^{-H\Delta\tau} \hat{\rho} \Psi_0 \rangle$

Theorem 13.2 (Quantum structure from polygon dynamics). *Starting from the classical Kirchhoff equations for N point vortices on $\mathbf{H}^2$, with no quantum postulate, the following quantum-mechanical structures are derived:*

(Q1) Hilbert space. $\mathcal{H} = L^2(\mathbb{R}, d\rho)$ (Feynman–Kac, Proposition 3.7).

(Q2) Hamiltonian. $H = -\frac{1}{2c} d^2/d\rho^2 + V(\rho)$ with $c = 12b(N)$ (CL mechanism + orbifold matching).

- (Q3) Physical state. *The WDW constraint $H\Psi_0 = 0$ selects the ground state (the additive constant in V is fixed by $V(\rho^*) = 0$ at the BO equilibrium, giving $E_0 = 0$; excited states $E \neq 0$ do not satisfy the constraint).*
- (Q4) Ground-state position distribution. $|\Psi_0(\rho)|^2 = e^{-\beta V}/Z$ (mathematical identity: the Boltzmann weight of the classical Onsager–Machlup path integral equals the quantum ground-state probability in the ρ -basis).
- (Q5) Pointer basis. *The angular bath decoheres ρ with $\gamma \gg 1$ near the stability threshold (Proposition 13.1). The ρ -basis is the pointer basis.*
- (Q6) Uncertainty principle. $\Delta\rho\Delta p = 1/2$ (minimum-uncertainty ground state near the turning point; $\Delta\rho = 1/\sqrt{2c\omega}$, $\Delta p = \sqrt{c\omega/2}$, $\omega^2 = |V''(\rho_1)| = \text{csch}^2(\rho_1)$).
- (Q7) Entanglement. *The entanglement entropy between ρ and the angular bath is $S_{\text{ent}} \gg 1$ bit (from the CL squeezing parameters $r_m = (\Phi'/(2\lambda_m))^2$; see Code Availability).*
- (Q8) Irreversibility. *The breathing-angular coupling breaks exact CMS integrability at $O(1/c)$, converting exact recurrences into diffusive mixing. The Nekhoroshev bound gives $\tau_{\text{mix}} \geq \exp(c_0 c^a)$; even this exponentially long mixing time satisfies $\tau_{\text{mix}} \ll \tau_{\text{tunnel}} \sim e^{S_{\text{BO}}}$ for generic exponents. Decoherence is irreversible on all cosmologically relevant timescales.*

Within the WDW framework ($H\Psi = 0$, unique state, no external time), items (Q1)–(Q5) are derived from the Feynman–Kac correspondence; items (Q6)–(Q8) are standard consequences of the WDW Hamiltonian structure (not specific to the polygon system).

Proof. Items (Q1)–(Q4) follow from the five-step chain above (Steps 1–5: Havelock $\rightarrow$ CL $\rightarrow$ Onsager–Machlup $\rightarrow$ Feynman–Kac $\rightarrow$ WDW). Item (Q5) is Proposition 13.1 with the mode-mixing irreversibility. Item (Q6): the WDW ground state near the turning point ρ_1 is $\Psi_0 \propto \exp(-c\omega(\rho-\rho_1)^2/2)$ with $\omega^2 = |V''(\rho_1)| = \text{csch}^2(\rho_1)$; the widths $\Delta\rho = 1/\sqrt{2c\omega}$ and $\Delta p = \sqrt{c\omega/2}$ give $\Delta\rho\Delta p = 1/2$ (minimum uncertainty). Item (Q7): the CL influence functional at zero temperature gives squeezing parameter $r_m = (\Phi'/(2\lambda_m))^2$ for each bath mode; the entanglement entropy is $S_{\text{ent}} = \sum_m [(r_m+1)\ln(r_m+1) - r_m \ln r_m]$. Item (Q8): the $O(1/c)$ W_N quartic coupling breaks exact CMS integrability; the Nekhoroshev bound gives $\tau_{\text{mix}} \geq \exp(c_0 c^a)$, and $\exp(c_0 c^a) \ll e^{S_{\text{BO}}}$ for generic exponents $a \leq 1/2$ at $c \approx 50$ –130, $S_{\text{BO}} \geq 38$. $\square$

Summary of the bridge.

Output	Mechanism	Input
Schrödinger (WDW)	Feynman–Kac	Onsager stat. mech.
CS Hilbert space	Vortex partition fn. at $k = c/6$	$b(N)$ from eigenvalues
3D Einstein	Havelock + Hadamard	Paper III, Thm. 5
4D vacuum Einstein	CS + KK	Paper IV, Theorem 13.3(a)
4D Einstein + matter	CS + KK + RT (derived from CFT ₂)	Paper IV, Theorem 13.3(b)
$\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$	McKay + Eightfold Way + KK	$N = 7$
$\sin^2 \theta_W = 3/11$	CS conformal weights	Paper IV, Theorem 8.1
$\delta_{\text{CKM}} = 70.2^\circ$	BF-crossing phase	$\arg \Gamma(\frac{1}{2} + i)$

Every row traces back to the classical vortex system. The quantum-mechanical description of the breathing mode enters through the Feynman–Kac theorem (a mathematical identity relating Brownian motion to the Schrödinger equation).

Remark 13.1 (Entropy bridge: Onsager coarse-graining to Bekenstein–Hawking). *The central charge $c = 12b(N)$ that appears as the Fisher–Rao information metric on the angular-mode Boltzmann distributions (eq. (8.1), Step 2 above) is the same c that enters the Cardy formula via Brown–Henneaux ($c = 3\ell/(2G_3)$, §5). The Bekenstein–Hawking entropy at the palindromic threshold is*

$$S_{\text{BH}} = \frac{\pi c}{3} \cosh \rho^* = \frac{\pi c}{3} T_{\text{GH}} L, \quad (13.2)$$

where $T_{\text{GH}} = \coth(\rho^*)/(2\pi)$ is the Gibbons–Hawking temperature and $L = 2\pi \sinh \rho^*$ is the horizon circumference. Three entropy levels illuminate the structure:

Level	Formula	$N=7$	Scaling
CL squeezing	$\sum_m [(r_m+1) \ln(r_m+1) - r_m \ln r_m]$	1.6	$O(1)$
One-loop (frozen det.)	$\frac{1}{2} \sum_{m \neq m^*} \ln \lambda_m(\rho^*) $	1.1	$O(N)$
Cardy (BH)	$(\pi c/3) \cosh \rho^*$	158	$O(e^{\rho^*})$

The Gaussian (Caldeira–Leggett) entanglement captures the same c but accounts for $O(1/c)$ of the total: the gap $S_{\text{BH}}/S_{1\text{-loop}} \approx 144$ (at $N = 7$, growing to 5.5×10^4 at $N = 15$) reflects the exponential density of Virasoro descendants inaccessible to the Gaussian approximation. The Onsager coarse-graining determines c ; the Cardy formula completes c to the full entropy.

Code availability. All numerical verifications are in the companion repository (<https://github.com/gspea/planetary-polygons-unified>). Parameter accounting and predictions appear in Paper VI.

References

- Thierry Aubin. *Some Nonlinear Problems in Riemannian Geometry*. Springer, 1998.
- Arthur L. Besse. *Einstein Manifolds*. Springer, 1987. doi: 10.1007/978-3-540-74311-8.
- Peter Buser. *Geometry and Spectra of Compact Riemann Surfaces*, volume 106 of *Progress in Mathematics*. Birkhäuser, 1992. doi: 10.1007/978-0-8176-4992-0.
- E. Caglioti, P.-L. Lions, C. Marchioro, and M. Pulvirenti. A special class of stationary flows for two-dimensional Euler equations: a statistical mechanics description. *Communications in Mathematical Physics*, 143:501–525, 1992.
- Francesco Calogero. Exactly solvable one-dimensional many-body problems. *Lettere al Nuovo Cimento*, 13:411–416, 1975.
- Steven Carlip. *Quantum Gravity in 2+1 Dimensions*. Cambridge University Press, 1998. doi: 10.1017/CBO9780511564192.

- Jeff Cheeger and Michael Taylor. On the diffraction of waves by conical singularities. I. *Communications on Pure and Applied Mathematics*, 35:275–331, 1982. doi: 10.1002/cpa.3160350302.
- Jeff Cheeger, Werner Müller, and Robert Schrader. On the curvature of piecewise flat spaces. *Communications in Mathematical Physics*, 92:405–454, 1984.
- Gianni Dal Maso. *An Introduction to Γ -Convergence*. Birkhäuser, 1993.
- Sebastian de Haro, Sergey N. Solodukhin, and Kostas Skenderis. Holographic reconstruction of spacetime and renormalization in the AdS/CFT correspondence. *Communications in Mathematical Physics*, 217:595–622, 2001.
- Lance Dixon, Jeffrey A. Harvey, Cumrun Vafa, and Edward Witten. Strings on orbifolds. *Nuclear Physics B*, 261:678–686, 1985.
- Arthur Erdélyi. *Asymptotic Expansions*. Dover, 1956.
- Pavel Etingof, Igor Frenkel, and Alexander Kirillov. *Lectures on Representation Theory and Knizhnik–Zamolodchikov Equations*, volume 58 of *Mathematical Surveys and Monographs*. American Mathematical Society, 1998.
- G. W. Gibbons and S. W. Hawking. Action integrals and partition functions in quantum gravity. *Physical Review D*, 15:2752–2756, 1977.
- Simone Giombi, Alexander Maloney, and Xi Yin. One-loop partition functions of 3D gravity. *Journal of High Energy Physics*, 2008(08):007, 2008. doi: 10.1088/1126-6708/2008/08/007.
- C. Robin Graham and Maciej Zworski. Scattering matrix in conformal geometry. *Inventiones mathematicae*, 152:89–118, 2003. doi: 10.1007/s00222-002-0268-1.
- Sigurdur Helgason. *Groups and Geometric Analysis: Integral Geometry, Invariant Differential Operators, and Spherical Functions*, volume 113 of *Pure and Applied Mathematics*. Academic Press, 1984.
- Ted Jacobson. Thermodynamics of spacetime: the Einstein equation of state. *Physical Review Letters*, 75:1260–1263, 1995. doi: 10.1103/PhysRevLett.75.1260.
- Michael K.-H. Kiessling. Statistical mechanics of classical particles with logarithmic interactions. *Communications on Pure and Applied Mathematics*, 46:27–56, 1993.
- Nima Lashkari, Michael B. McDermott, and Mark Van Raamsdonk. Gravitational dynamics from entanglement “thermodynamics”. *Journal of High Energy Physics*, 2014(04):195, 2014. doi: 10.1007/JHEP04(2014)195.
- Patrick M. Lenggenhager, Joseph Maciejko, and Tomáš Bzdušek. Non-Abelian hyperbolic band theory from supercells. *Physical Review B*, 108:085114, 2023. Haldane model on the $\{8, 3\}$ Bolza lattice.
- M. A. Olshanetsky and A. M. Perelomov. Classical integrable finite-dimensional systems related to Lie algebras. *Physics Reports*, 71:313–400, 1981.

- Barrett O'Neill. The fundamental equations of a submersion. *Michigan Mathematical Journal*, 13:459–469, 1966. doi: 10.1307/mmj/1028999604.
- Lars Onsager. Statistical hydrodynamics. *Il Nuovo Cimento*, 6:279–287, 1949.
- Riccardo Rattazzi, Vyacheslav S. Rychkov, Erik Tonni, and Alessandro Vichi. Bounding scalar operator dimensions in 4D CFT. *Journal of High Energy Physics*, 2008(12):031, 2008. doi: 10.1088/1126-6708/2008/12/031.
- Alexander Strohmaier and Ville Uski. An algorithm for the computation of eigenvalues, spectral zeta functions and zeta-determinants on hyperbolic surfaces. *Communications in Mathematical Physics*, 317:827–869, 2013. doi: 10.1007/s00220-012-1557-1. arXiv:1110.2150; ancillary data includes 1000 Bolza eigenvalues at 48-digit precision.
- Miguel Tierz. Schur polynomials and biorthogonal random matrix ensembles. *Journal of Mathematical Physics*, 51:063509, 2010. Stieltjes–Wigert polynomials and Chern–Simons on Seifert manifolds.
- Chiara Toldo and Brian Willett. Partition functions on 3d circle bundles and their gravity duals. *Journal of High Energy Physics*, 2018:116, 2018.
- Edward Witten. $2 + 1$ dimensional gravity as an exactly soluble system. *Nuclear Physics B*, 311:46–78, 1988.
- Edward Witten. Quantization of Chern–Simons gauge theory with complex gauge group. *Communications in Mathematical Physics*, 137:29–66, 1991.
- Edward Witten. Anti de Sitter space and holography. *Advances in Theoretical and Mathematical Physics*, 2:253–291, 1998.
- A. B. Zamolodchikov. Conformal scalar field on the hyperelliptic curve and critical Ashkin–Teller multipoint correlation functions. *Nuclear Physics B*, 285:481–503, 1987.
- Wojciech H. Zurek. Decoherence, einselection, and the quantum origins of the classical. *Reviews of Modern Physics*, 75:715–775, 2003.